

# Machine learning: Neural networks

Bulat Ibragimov

bulat@di.ku.dk

Department of Computer Science  
University of Copenhagen

UNIVERSITY OF COPENHAGEN



# Learning Objectives


- One-neuron network
- Loss function
- Backpropagation
- Activation function
- Multi-neuron network

# Machine learning (ML) in 1 line

1-line summary of ML:  $y = \text{sign}(f(z))$






- Simple example of  $f(x) = Wz + b$
- ... you just need to have the right  $z$ .

Data representation



# Machine learning in medical image analysis


- How useful hand-crafted features are?
- Very useful, but limited




# Machine learning in medical image analysis

Which hand-crafted features will turn to be useful for a selected classifier?

Decision forests



Support vector machines



# Machine learning in medical image analysis

Hand-crafted features work well when target objects are:


- Oriented and positioned in a predictable way



- Relatively similar to each other in appearance



# Neural networks




Neuron components:

- Feature  $\rightarrow W^T x + b$
- Activation  $\rightarrow \sigma$

# Neural networks

## Deep neural network

- Output of a neuron is a number
- Output of a neuron can be passed to next neurons as an input






# 1-neuron network

Patient survival depends on cancer size:

- Bigger tumors are worse
- Can we model this dependency?


	Tumor Size	Survival
Case1	0.5	3.2
Case2	2.3	1.9
Case3	2.9	1.0



# 1-neuron network

Network design:

- Input  $x$  is 1-dimensional (tumor size)
- Output  $y$  is 1-dimensional (survival time)
- Activation function  $\sigma$  is identity function



# 1-neuron network: initialization


Let's initialize our network:

- $W = 0.01$ ;  $b = 0.01$

	Tumor Size	Survival
Case1	0.5	3.2
Case2	2.3	1.9
Case3	2.9	1.0

The performance of the network is low:

- $y_1 = Wx_1 + b = 0.01 * 0.5 + 0.01 = 0.015$
- $y_2 = Wx_2 + b = 0.01 * 2.3 + 0.01 = 0.033$
- $y_3 = Wx_3 + b = 0.01 * 2.9 + 0.01 = 0.039$



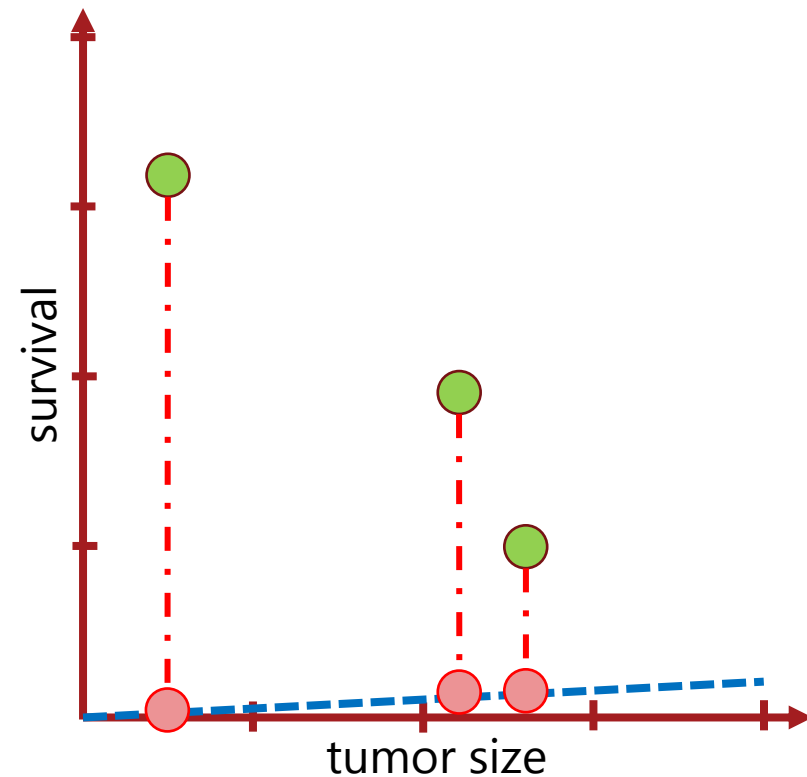
# 1-neuron network: loss function

The performance of the network is low, but how low?

	Tumor Size	Survival
Case1	0.5	3.2
Case2	2.3	1.9
Case3	2.9	1.0

Sum of absolute differences:

- $Loss = \sum_i |y'_i - y_i| = \sum_i |(Wx_i + b_i) - y_i|$
- $Loss =$   
 $|0.015 - 3.2| +$   
 $|0.033 - 1.9| +$   
 $|0.039 - 1.0| = 6.013$




# 1-neuron network: loss function

The performance of the network is low, but how low?

	Tumor Size	Survival
Case1	0.5	3.2
Case2	2.3	1.9
Case3	2.9	1.0

Mean squared error:

- $Loss = \sum_i (y'_i - y_i)^2 = \sum_i ((Wx + b) - y)^2$
- $Loss =$   
 $(0.015 - 3.2)^2 +$   
 $(0.033 - 1.9)^2 +$   
 $(0.039 - 1.0)^2 = 14.55$



# 1-neuron network: derivatives


We want the loss to be as small as possible, i.e. find its minimum.

We use derivatives to find minima/maxima of a function:

- How fast function changes
- Will it increase or decrease

- Chain rule:

$$\frac{\partial}{\partial x} f(g(x)) = \frac{\partial}{\partial g} f(g(x)) \cdot \frac{\partial}{\partial x} g(x)$$



derivatives are:

← negative      positive →


# 1-neuron network : backpropagation

Initialization:

- $W = 0.01, b = 0.01$

	Tumor Size	Survival
Case1	0.5	3.2
Case2	2.3	1.9
Case3	2.9	1.0

Forward pass:



# 1-neuron network: backpropagation

Compute the derivatives:

$$Loss = \sum_i (y' - y)^2$$


- How changes in  $\sigma(Wx + b)$  affect the loss

- How changes in  $W$  and  $b$  affect  $g$

- $\frac{\partial Loss}{\partial W}$  and  $\frac{\partial Loss}{\partial b}$  will be computed using the chain rule:

- $\frac{\partial Loss}{\partial W} = 2 \sum_i ((g(x_i) - y_i) \cdot x_i)$

- $\frac{\partial Loss}{\partial b} = 2 \sum_i ((g(x_i) - y_i) \cdot 1)$





# 1-neuron network: optimization step

Update  $W$  and  $b$  according to the derivatives:

- $W \leftarrow W - \lambda \frac{\partial Loss}{\partial W}; \quad b \leftarrow b - \lambda \frac{\partial Loss}{\partial b}$

Learning rate

The results of the optimization step:

- $W \leftarrow W - 0.1 \cdot 2 \sum_i ((g(x_i) - y_i) \cdot x_i) = 0.01 - 0.1 \cdot 2 \cdot ((0.015 - 3.2)0.5 + (0.033 - 1.9)2.3 + (0.039 - 1.0)2.9) = 0.01 - 0.1 \cdot -8.67 = 0.88$

W	b
0.01	0.01

- $b \leftarrow b - 0.1 \cdot 2 \sum_i ((g(x_i) - y_i)) = 0.61$

	Tumor Size	Survival	First solution
Case1	0.5	3.2	0.015
Case2	2.3	1.9	0.033
Case3	2.9	1.0	0.039

# 1-neuron network: optimization step

After parameter update:


- $W = 0.88$ ;  $b = 0.61$

The performance of the network is better:

- $y_1 = Wx_1 + b = 0.70 * 0.5 + 0.61 = 1.05$
- $y_2 = Wx_2 + b = 0.70 * 2.3 + 0.61 = 2.63$
- $y_3 = Wx_3 + b = 0.70 * 2.9 + 0.61 = 3.16$

Loss improves  $Loss = 9.8$

	Tumor Size	Survival
Case1	0.5	3.2
Case2	2.3	1.9
Case3	2.9	1.0



# 1-neuron network: optimization step

After next parameter update:


- $W = 0.19$ ;  $b = 0.54$

The performance of the network is better:

- $y_1 = Wx_1 + b = 0.45 * 0.5 + 0.38 = 0.63$
- $y_2 = Wx_2 + b = 0.45 * 2.3 + 0.38 = 0.98$
- $y_3 = Wx_3 + b = 0.45 * 2.9 + 0.38 = 1.10$

Loss improves  $Loss = 7.46$

	Tumor Size	Survival
Case1	0.5	3.2
Case2	2.3	1.9
Case3	2.9	1.0




# 1-neuron network: optimization step

After next parameter update:

- $W = 0.50; b = 0.88, Loss = 6.07$
- $W = 0.19; b = 0.50, Loss = 7.75$
- $W = 0.53; b = 0.85, Loss = 6.29$
- $W = 0.19; b = 0.90, Loss = 5.37$
- $W = 0.29; b = 1.13, Loss = 4.67$
- $W = 0.12; b = 1.23, Loss = 4.14$
- ...
- $W = 0.04; b = 1.51, Loss = 3.27$
- $W = -0.15; b = 1.98, Loss = 2.06$
- $W = -0.22; b = 2.18, Loss = 1.63$
- $W = -0.29; b = 2.35, Loss = 1.31$
- $W = -0.36; b = 2.50, Loss = 1.04$
- $W = -0.43; b = 2.63, Loss = 0.82$

	Tumor Size	Survival
Case1	0.5	3.2
Case2	2.3	1.9
Case3	2.9	1.0



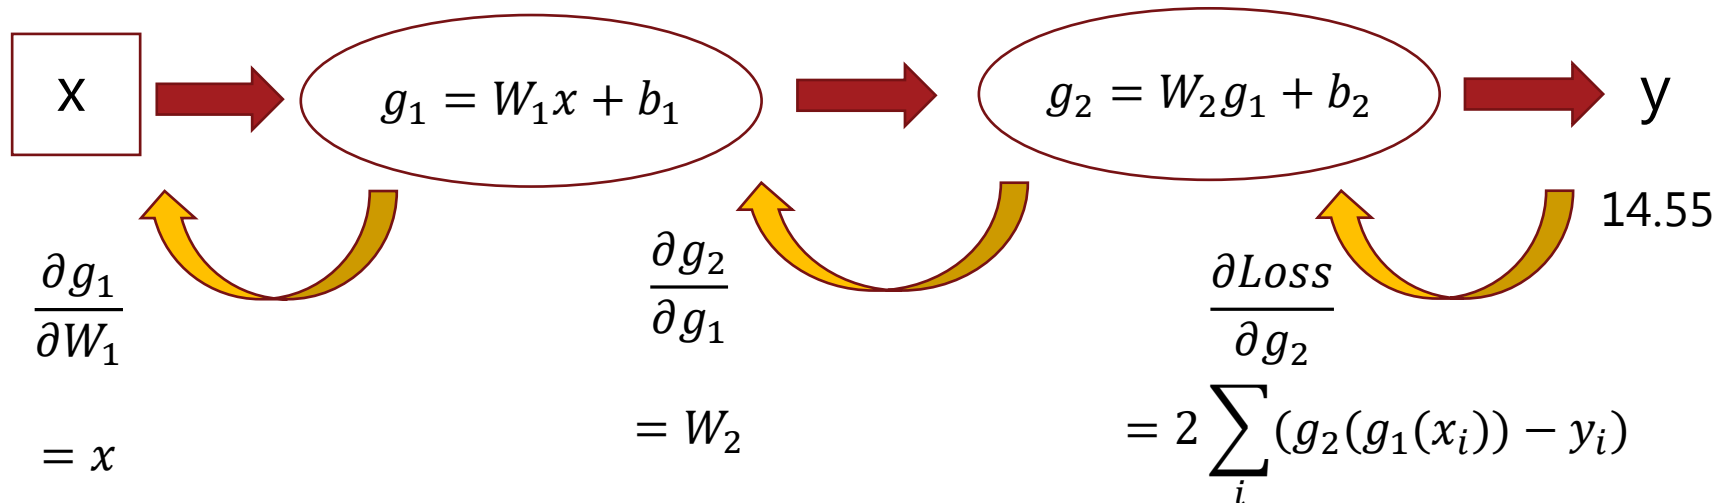
# More neurons

The chain rule can be applied for longer chains of neurons:

Let's initialize our network:

- $W_1 =$  ;  $b_1 =$   $W_2 = 0.01; b_2 = 0.01$

$$Loss = \sum_i (g_2(g_1(x_i)) - y_i)^2$$



$$\frac{\partial Loss}{\partial W_1} = 2 \sum_i (g_2(g_1(x_i)) - y_i) W_2 x$$

# More neurons: parameter update


We need to update all network parameters  $W_1, b_1, W_2, b_2$ :

- First we update the parameters on the lowest level:  $W_1, b_1$

$$\frac{\partial Loss}{\partial W_1}; \frac{\partial Loss}{\partial b_1}$$

- Then we move to the next level:  $W_2, b_2$

$$\frac{\partial Loss}{\partial W_2}; \frac{\partial Loss}{\partial b_2}$$




# Multidimensional input

$W$  will be in agreement with input  $x$  dimensions :


	Tumor Size	Temperature	Survival
Case1	0.5	37.6	3.2
Case2	2.3	39.1	1.9
Case3	2.9	36.2	1.0

- $g = \sigma(W_1(tumor_{size}) + W_2(temperature) + b)$




# Activation function

Let's compute this:



$$y' = g_2(g_1(x)) = W_2(W_1x + b_1) + b_2 = W_2W_1x + (W_2b_1 + b_2)$$

Can we simplify the network?

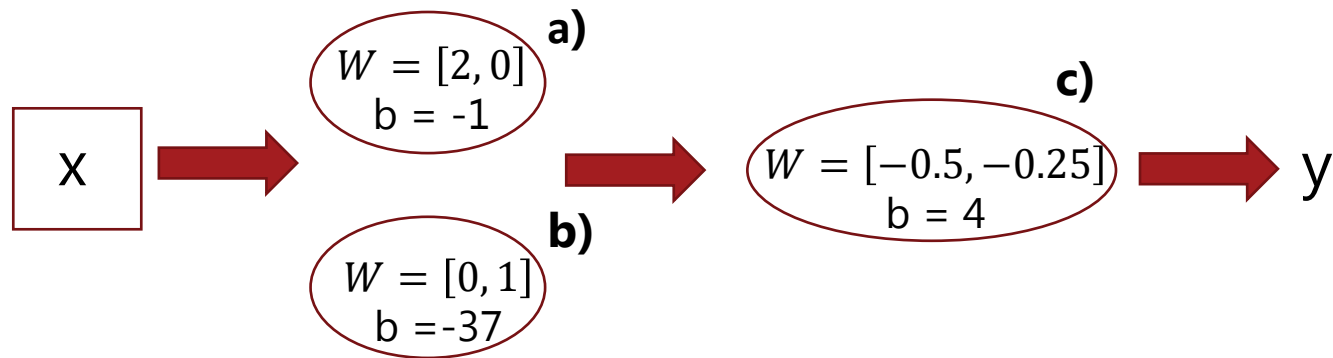


The network is not deep at all



# Activation function

On high level, what the following neurons try to capture?



Neuron a) gets excited about

Neuron b) gets excited about

Neuron c) gets predicts


	Tumor Size	Temperature	Survival
Case1	0.5	37.6	3.2
Case2	2.3	39.1	1.9
Case3	2.9	36.2	1.0

# Activation function

Let's add some more data into the database


What will happen with the network behavior?

	Tumor Size	Temperature	Survival
Case1	0.5	37.6	3.2
Case2	2.3	39.1	1.9
Case3	2.9	36.2	1.0
Case4	2.8	22	0.0
Case5	2.1	23	0.0



# Activation function

Maybe we can add a neuron to catch dead patients



Will this new neuron fix the network behavior on dead patients?

# Activation function

Neuron

**b)**  
 $W = [0, 1]$   
 $b = -37$

**d)**  
 $W = [0, -1]$   
 $b = 20$

	Tumor Size	Temperature	Survival
Case2	2.3	39.1	1.9
Case5	2.1	23	0.0

2.1

-19.1

-14


3

The idea of neuron b) is to capture fever and not to care about low temperature


The idea of neuron d) is to capture dead body temperature and not to care about fever

# Activation function

Sigmoid activation:



ReLU:



# Regularization

Let's augment our loss functions:

$$Loss = \sum_i (y' - y)^2 + \mu \sum_j |W_j|$$

$$Loss = \sum_i (y' - y)^2 + \mu \sum_j (W_j)^2$$


What will these additional terms do?

# Regularization

L1 regularization tries to minimize the number of features used in the analysis

More features you use – the more likely it is to find some dependencies.


One person with a rare name X is carpenter, dose this mean that X can predict the profession?



# Regularization

L2 regularization tries to use all features equally

Do not rely on one feature too much, you can get outliers.






# Brief CNNs: L2 regularization

- Regularization makes your network more robust

$\mu = 0.001$



$\mu = 0.01$




$\mu = 0.1$



# Regression/classification

L2 regularization tries to use all features equally

Do not rely on one feature too much, you can get outliers.



# Questions?