

# Machine learning: Neural networks

Bulat Ibragimov

[bulat@di.ku.dk](mailto:bulat@di.ku.dk)

Department of Computer Science  
University of Copenhagen

UNIVERSITY OF COPENHAGEN



# Learning Objectives

- Convolution operation
- Backpropagation for convolution
- Convolutional neural network in keras
- Loss functions useful in medical imaging

# Machine learning in medical image analysis

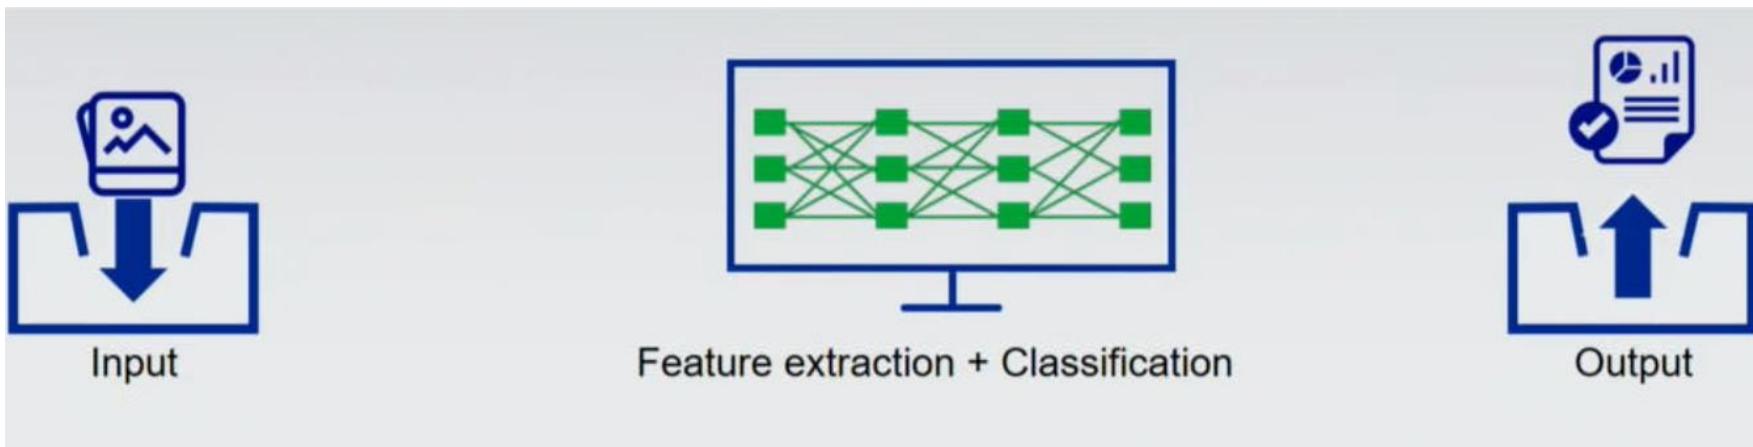
Which hand-crafted features will turn to be useful for a selected classifier?



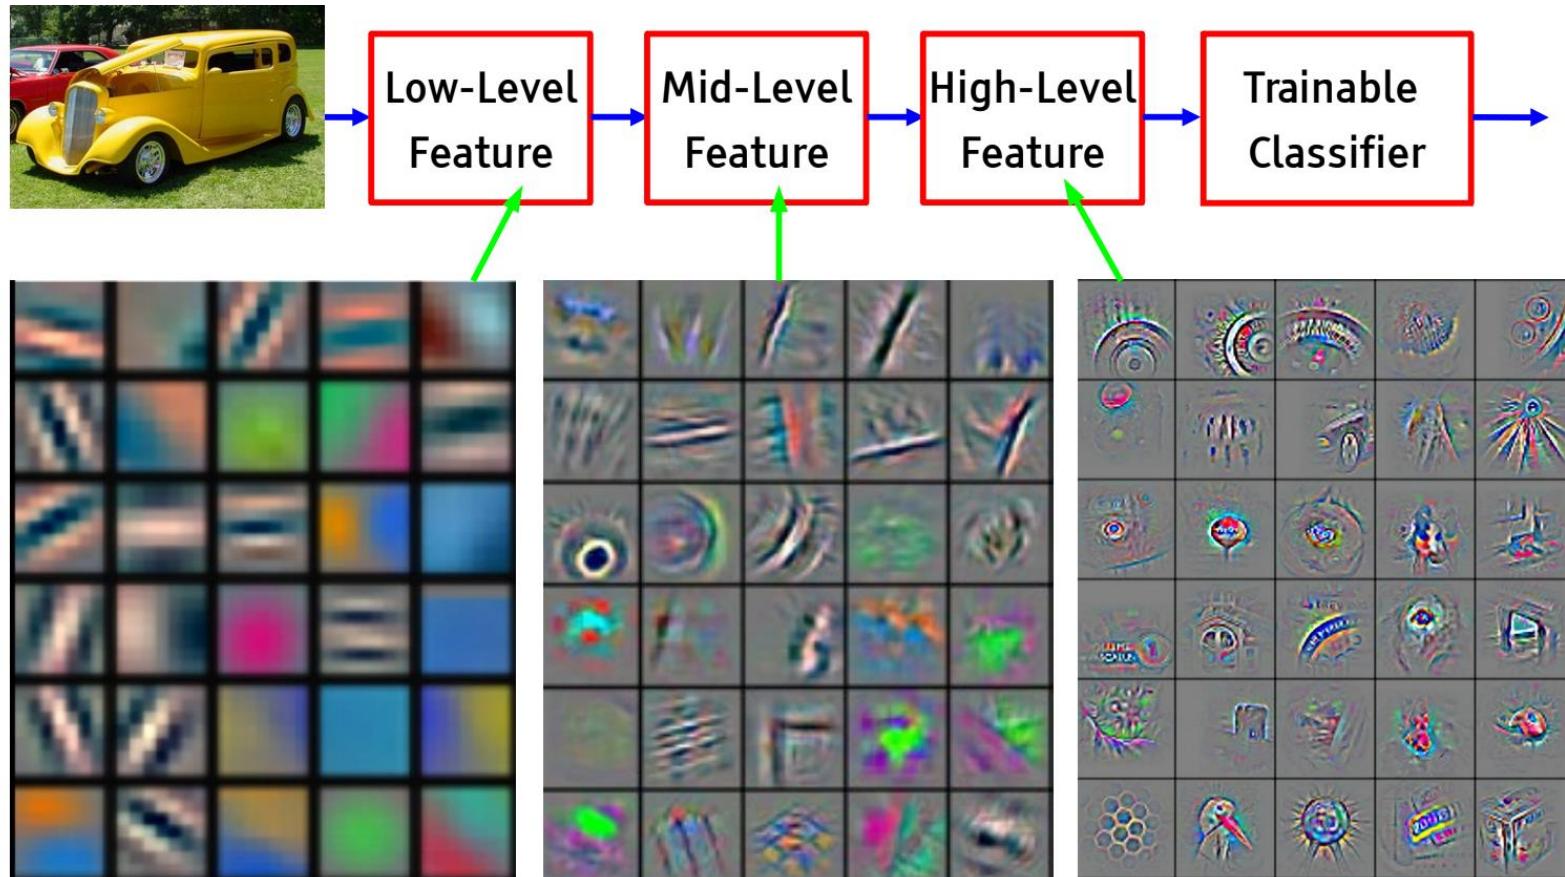
# Deep learning in image analysis

Complex features are often needed

Deep learning



# Deep learning in image analysis



Feature visualization of convolutional net trained on ImageNet from [Zeiler & Fergus 2013]

# DL: convolution operation

Image **convolution** is an **element-wise multiplication** of two matrices **followed by a sum**

The diagram illustrates the convolution operation between two 3x3 input matrices and one 3x3 kernel matrix.

**Input Matrices:**

1	0	-1
1	0	-1
1	0	-1

3	0	1
1	5	8
2	7	2

**Kernel Matrix:**

3	0	-1
1	0	-8
2	0	-2

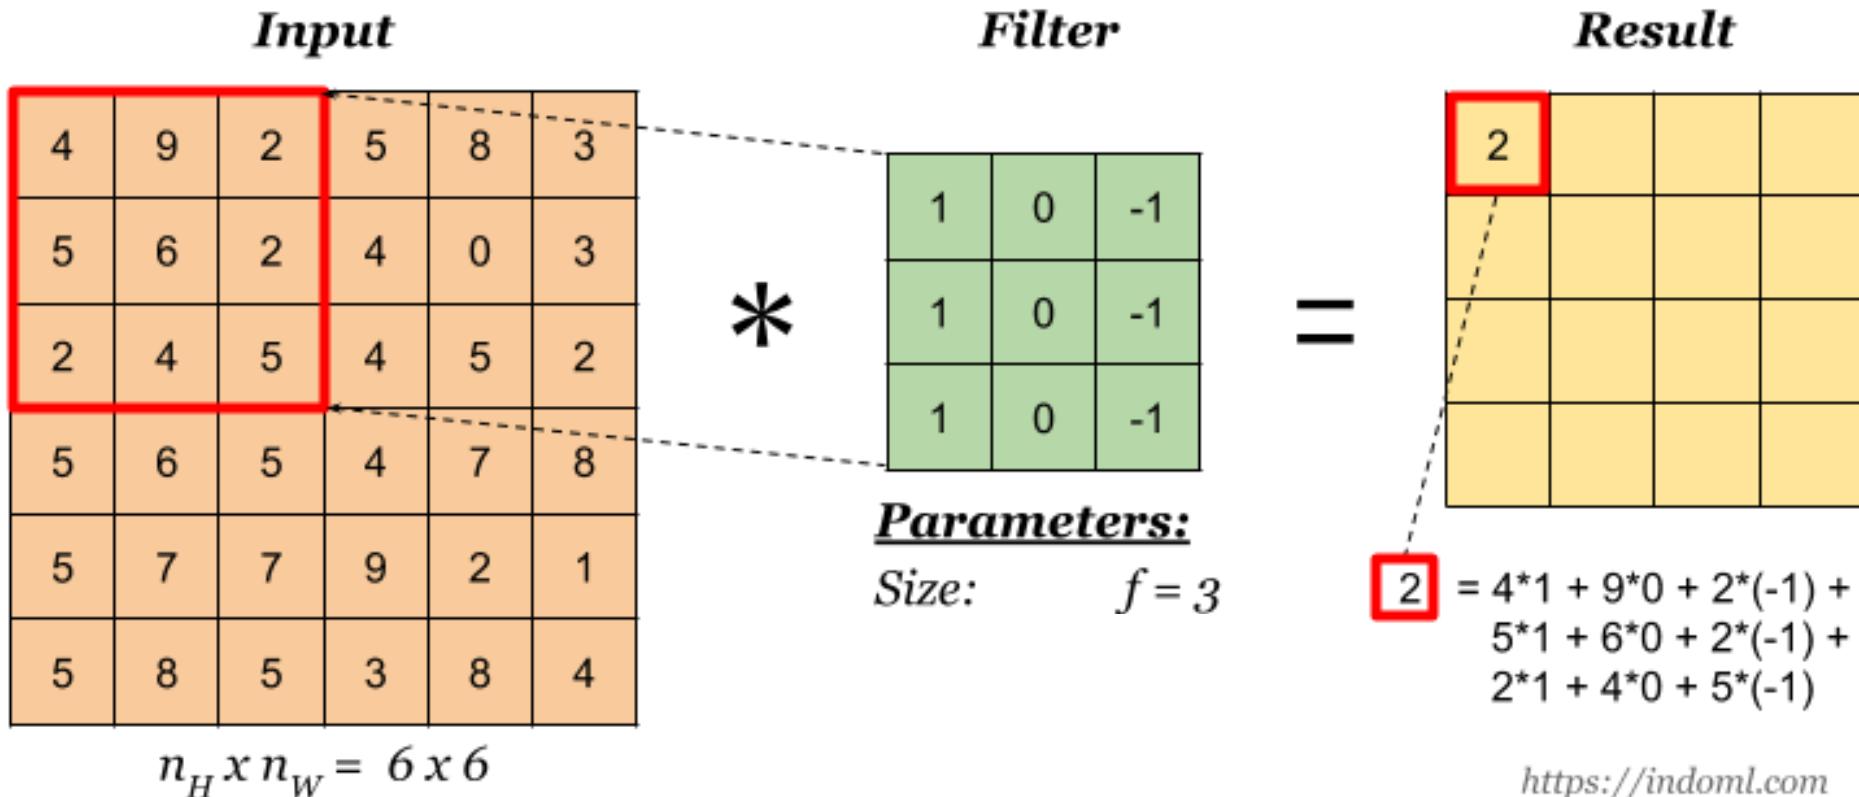
**Operation:** The diagram shows the element-wise multiplication of the input matrices with the kernel matrix. A red arrow points from the first element of the kernel to the first element of the top-left input matrix. The text "element-to-element multiplication" is written in red between the two input matrices.

**Result:** The result of the element-wise multiplication is shown below the inputs. An equals sign follows the multiplication result, and a red arrow points down to the sum of all elements, labeled "Sum all Elements".

**Final Result:** The final result is  $-5$ .

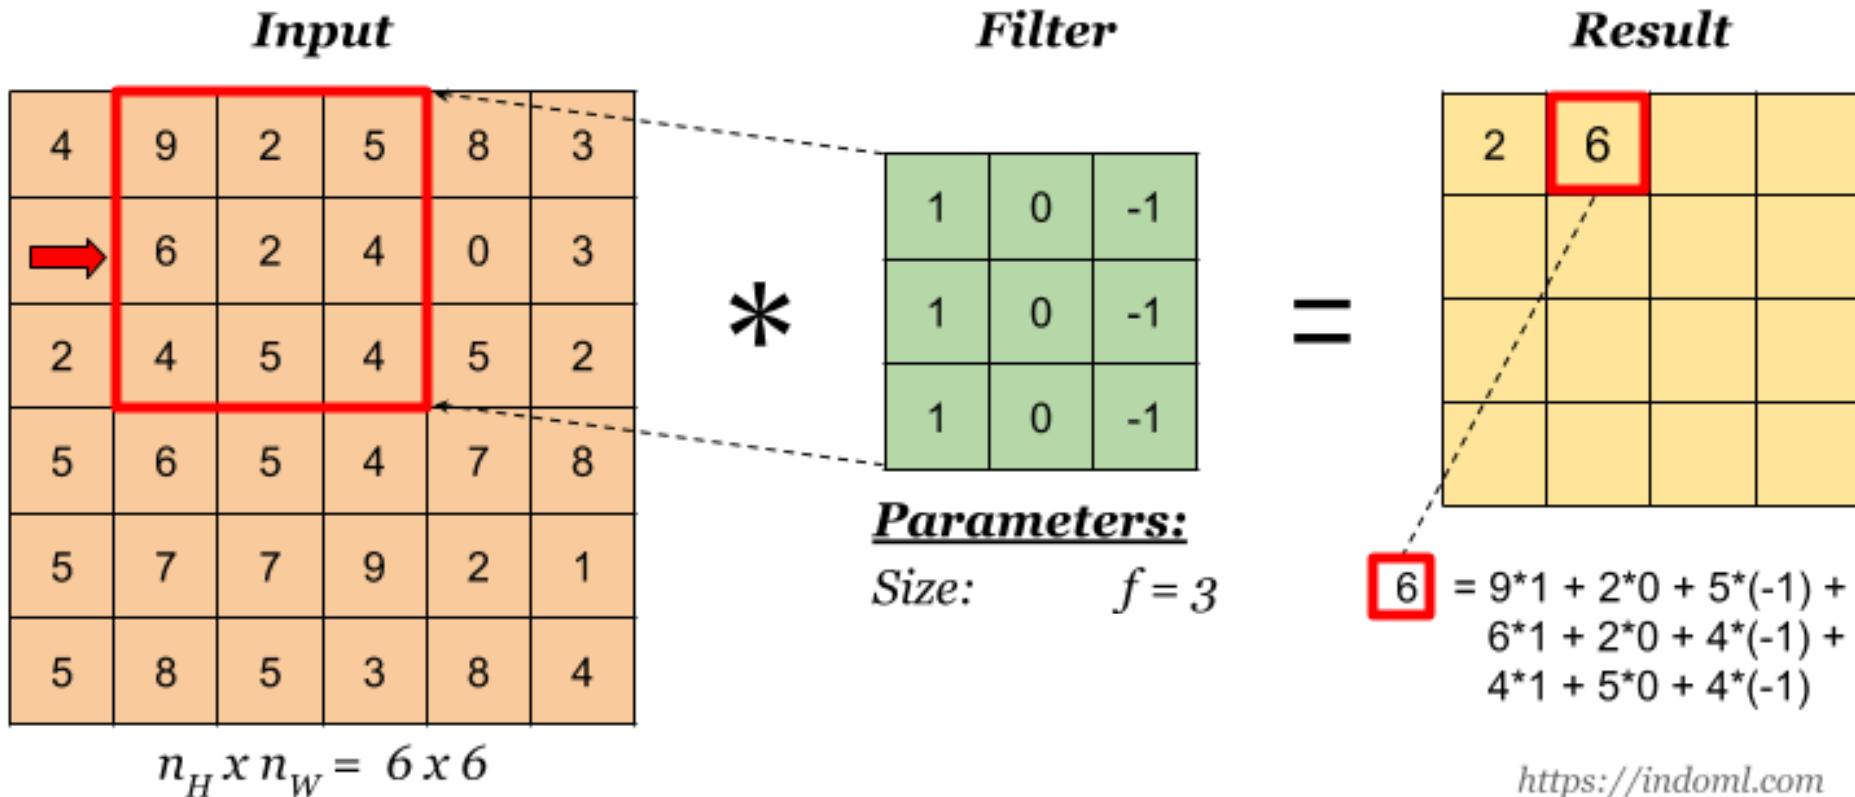
# Brief CNNs: convolution

Elementwise multiplication of input signal and filter



# Brief CNNs: convolution

Elementwise multiplication of input signal and filter



What if we do not want shrinking of the input size?

# DL: convolution operation



What do these operations do?

**Identity**

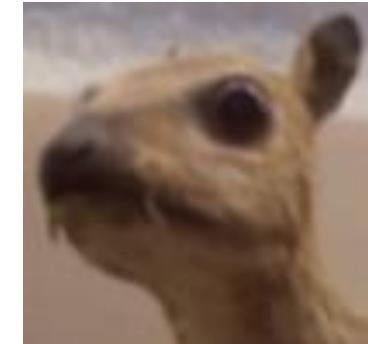
$$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$$

**Sharpen**

$$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$$

**Blur**

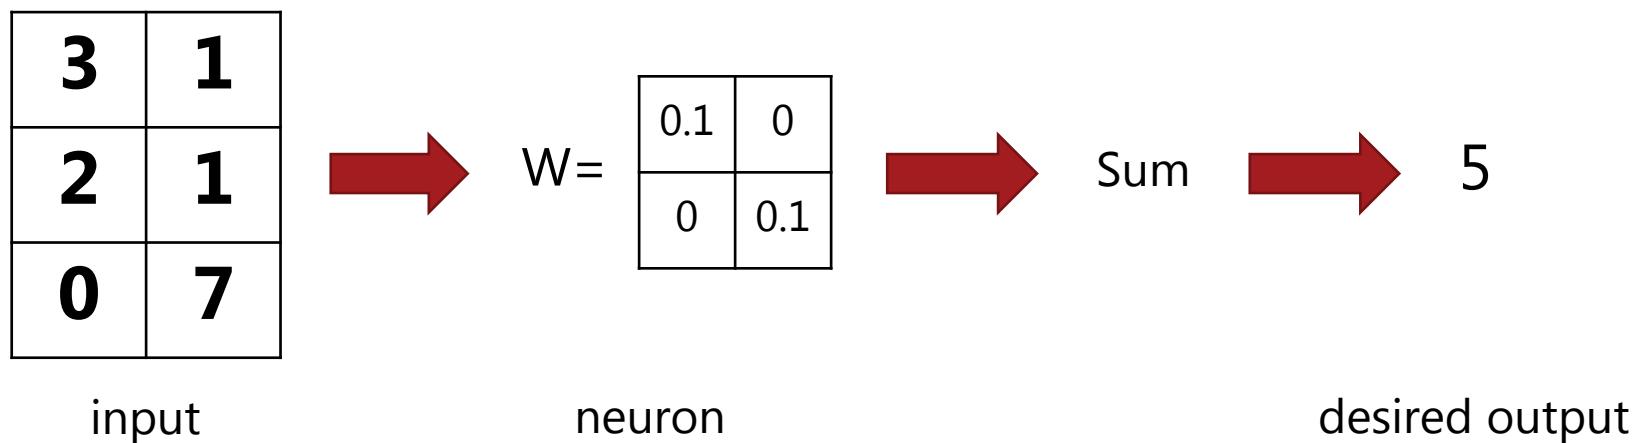
$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$



# Convolution operation: backpropagation

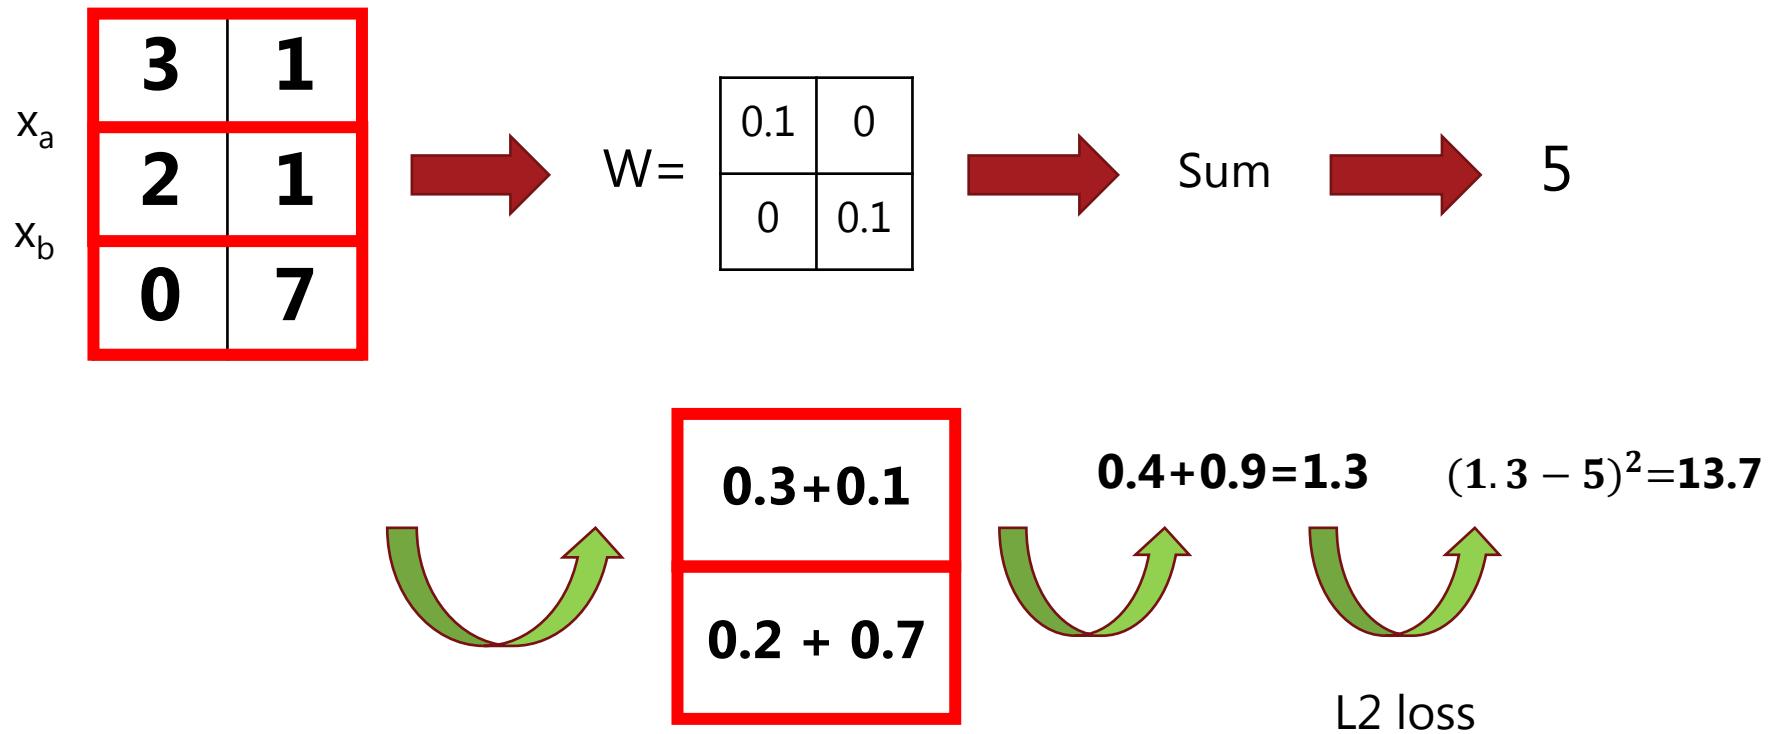
Network design:

- Input  $x$  is 3x2 array
- Neuron is 2x2 convolution with no bias
- Output  $y$  is 1D value
- Activation function  $\sigma$  is identity function



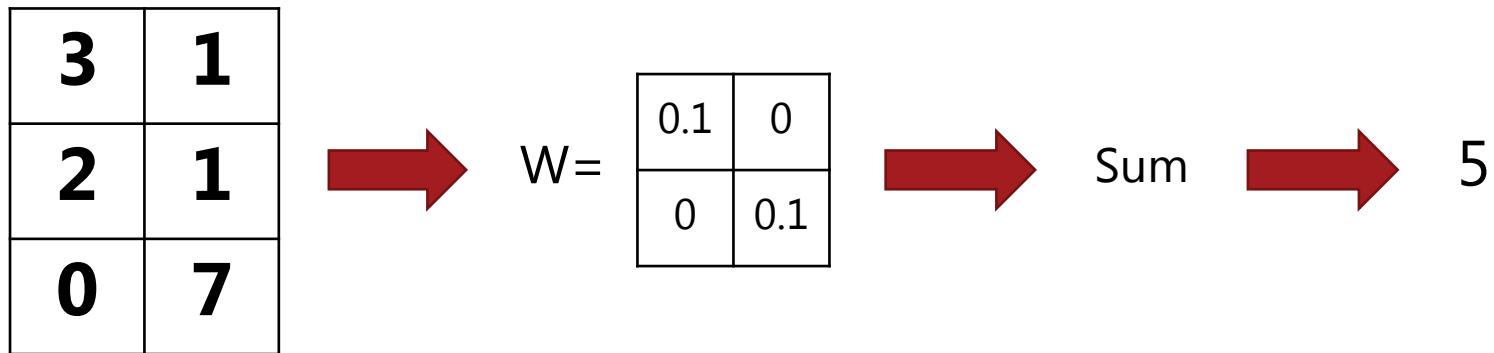
# Convolution operation: backpropagation

Backward pass:



# Convolution operation: backpropagation

Backward pass:



$$\begin{aligned}\frac{\partial Loss}{\partial w_{0,0}} &= 2 \sum (W(x_a) + W(x_b) - y_i) \\ &= 2x_{0,0} \sum (W(x_a) + W(x_b) - y_i) \\ &\quad + 2x_{1,0} \sum (W(x_a) + W(x_b) - y_i)\end{aligned}$$



13.7

# Convolution operation: backpropagation

Backward pass:

$$\begin{aligned} x_{0,0} &= \begin{array}{|c|c|} \hline 3 & 1 \\ \hline 2 & 1 \\ \hline 0 & 7 \\ \hline \end{array} \\ x_{1,0} &= \begin{array}{|c|c|} \hline 2 & 1 \\ \hline 0 & 7 \\ \hline \end{array} \end{aligned}$$

$$\begin{aligned} W(x_a) &= \begin{array}{|c|} \hline 0.3 + 0.1 \\ \hline \end{array} \\ W(x_b) &= \begin{array}{|c|} \hline 0.2 + 0.7 \\ \hline \end{array} \end{aligned}$$

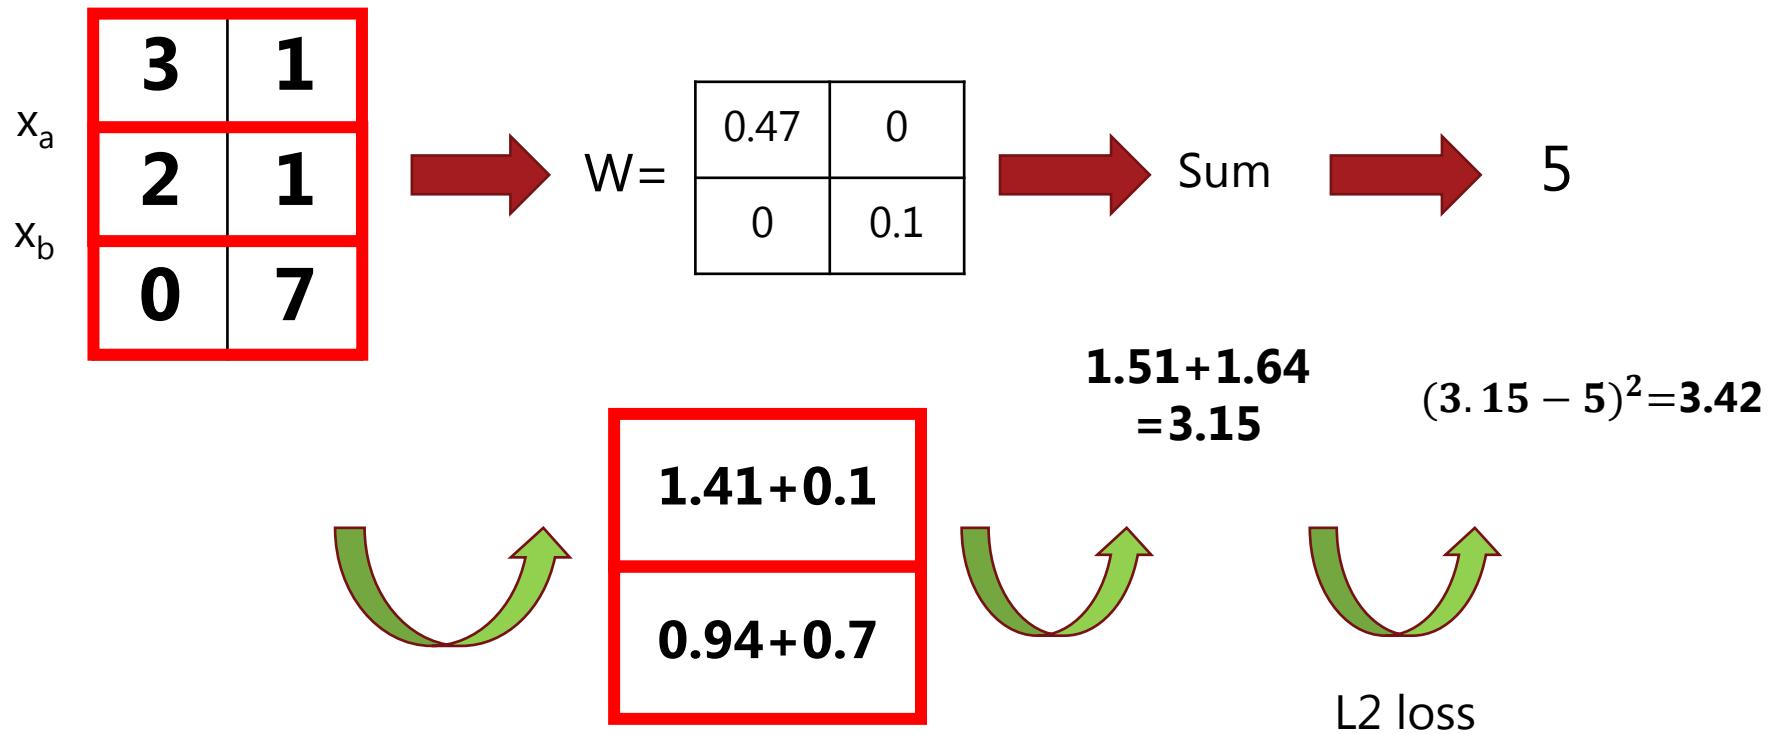
$$\frac{\partial Loss}{\partial w_{0,0}} = 2x_{0,0} \sum (W(x_a) + W(x_b) - y_i) + 2x_{1,0} \sum (W(x_a) + W(x_b) - y_i) =$$

$$w_{0,0} \leftarrow w_{0,0} - \gamma \frac{\partial Loss}{\partial W} = 0.1 - 0.01(-37) = 0.47$$

$$W = \begin{array}{|c|c|} \hline 0.1 & 0 \\ \hline 0 & 0.1 \\ \hline \end{array}$$

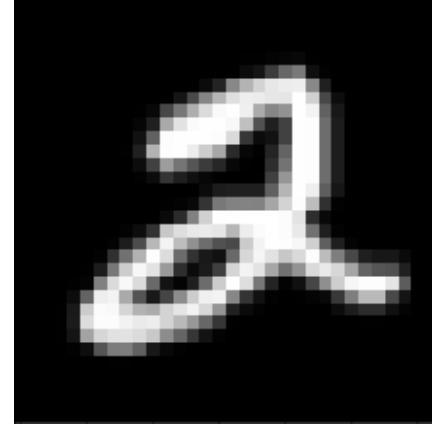
# Convolution operation: backpropagation

Forward pass:



# Convolution

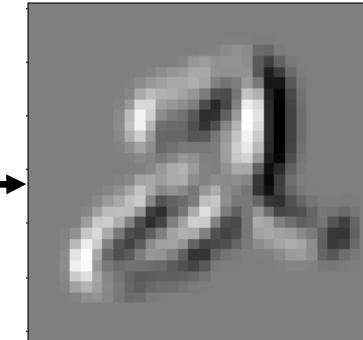
This is an example of convolution layer with filter size 3x3 and number of features 3



Input

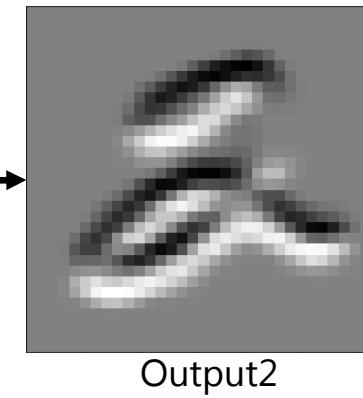
-1	0	1
-1	0	1
-1	0	1

Filter1



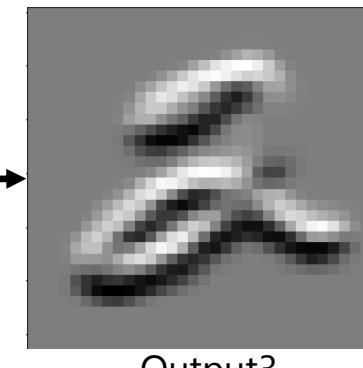
1	1	1
0	0	0
-1	-1	-1

Filter2



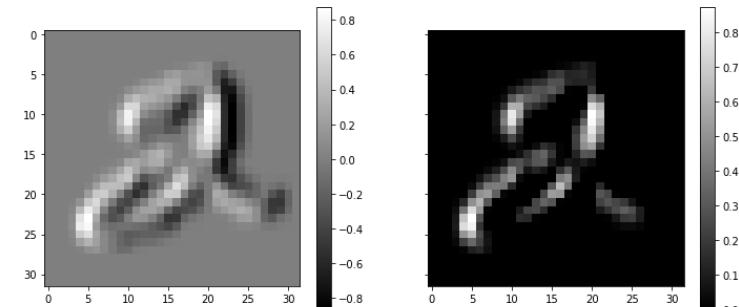
-1	-1	-1
0	0	0
1	1	1

Filter3



# Convolution: activation function

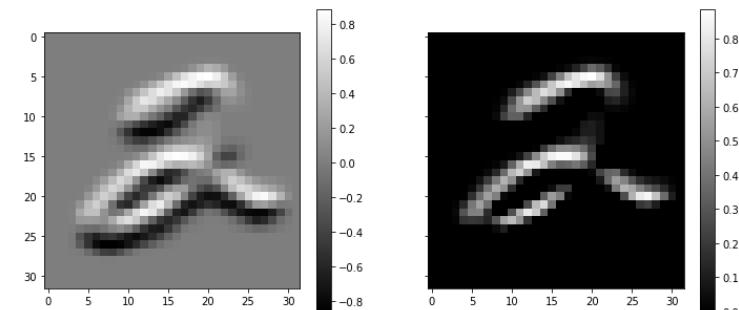
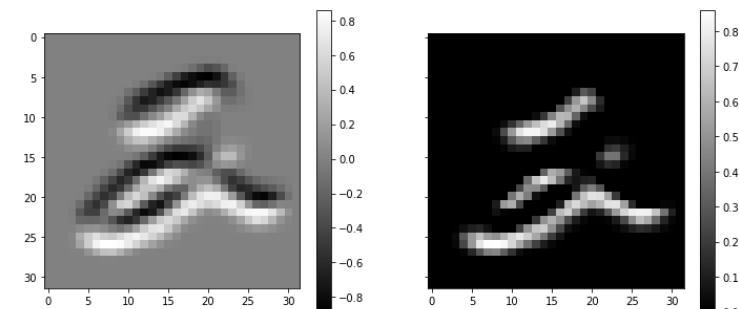
Raw convolution layers can be transformed into one convolution layer, so there will be no depth in the network\*



Activation function addresses this issue:

- $\text{Tanh}(x)$
- $\text{Sigmoid}(x)$
- $\text{ReLU}(x)$

$$\text{ReLU}(x) = \max(0, x)$$



\*<https://medium.com/the-theory-of-everything/understanding-activation-functions-in-neural-networks-9491262884e0>

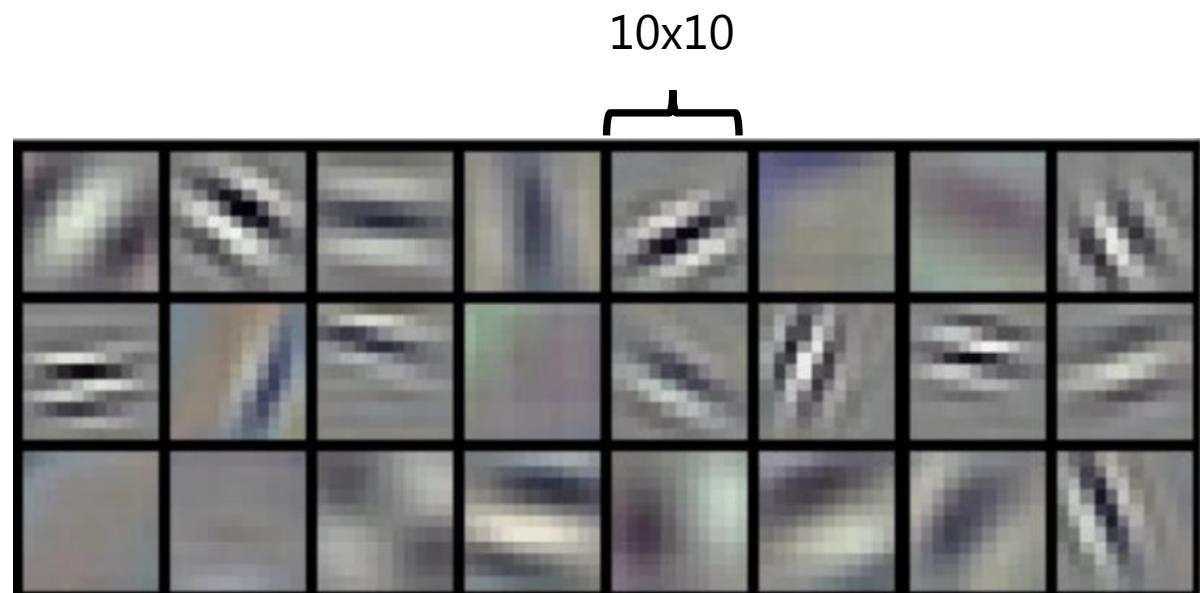
# Convolution in keras

```
conv = Conv2D(24, 10, activation = 'relu', padding = 'same')(prev_Layer)
```

24 – number of features

10 – size of each feature

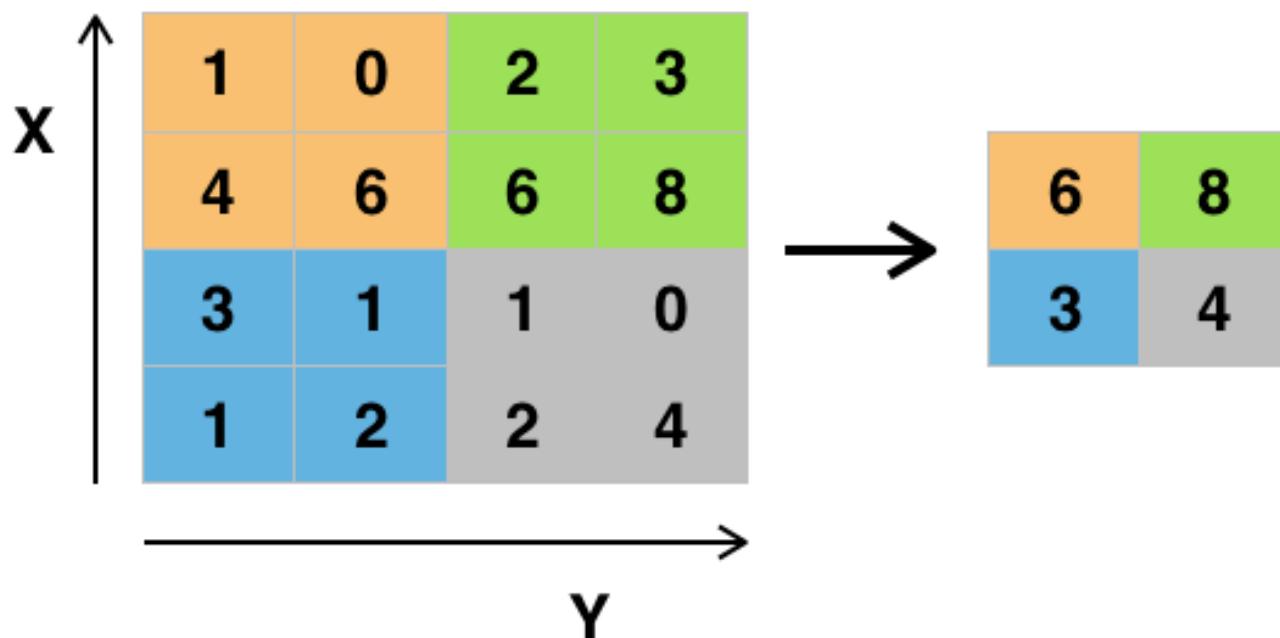
Relu - nonlinearity



# Pooling

Pooling layers:

- Reduce dimensionality of the input
- Introduce some translation invariance



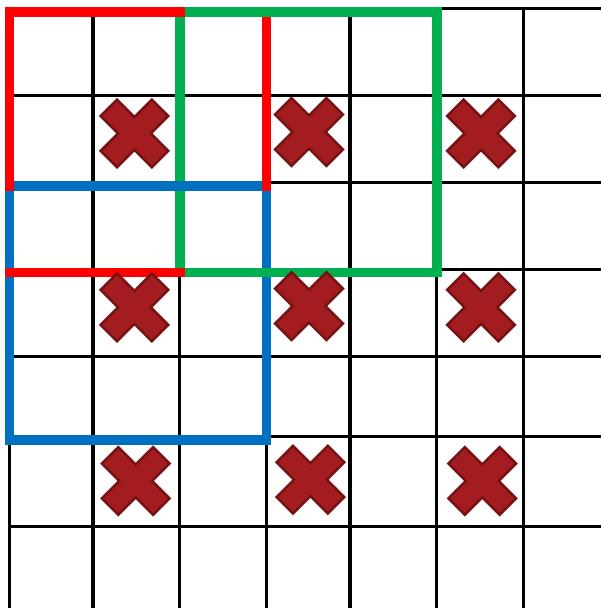
2x2 maxpooling

# Convolution with stride

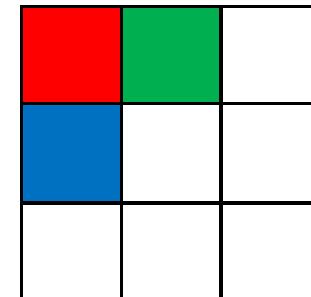
Convolution with stride >1:

- Reduce dimensionality of the input
- Introduce some translation invariance

7 x 7 Input Volume

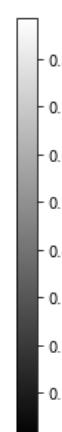
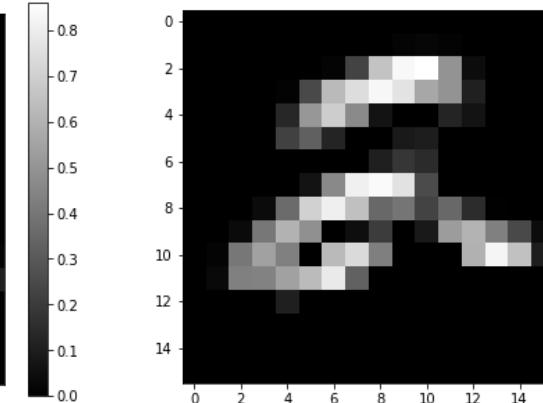
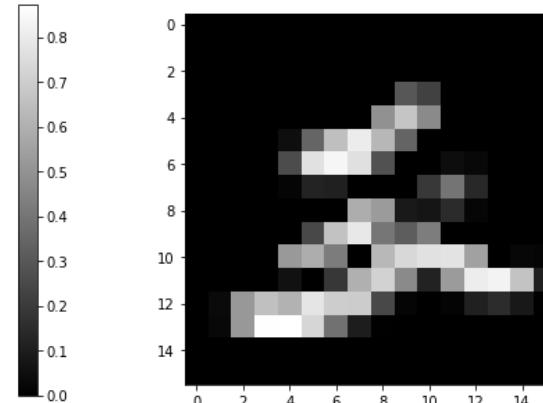
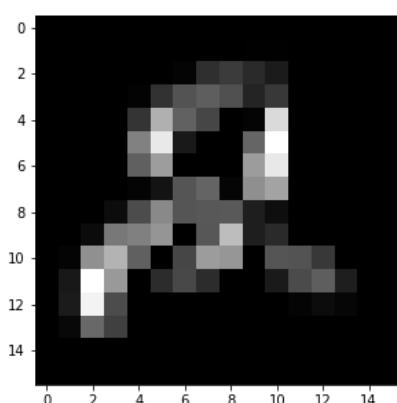
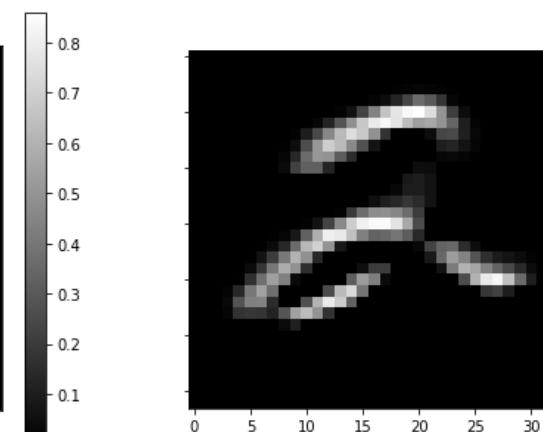
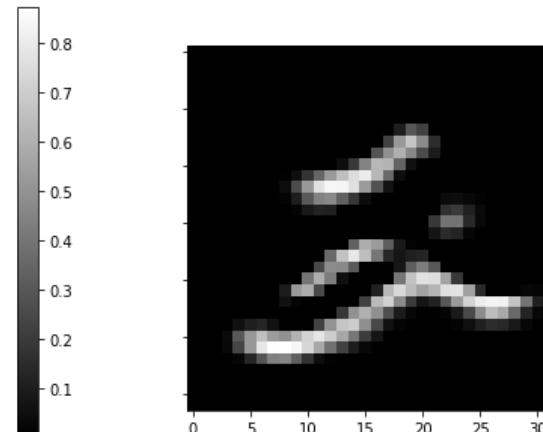
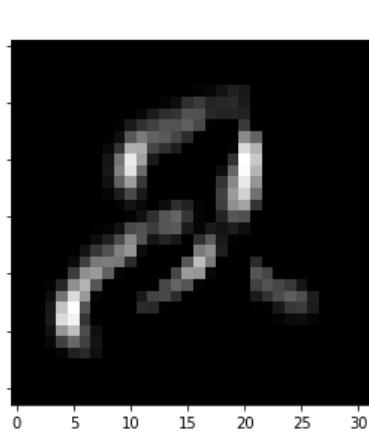


3 x 3 Output Volume



# Pooling in keras

pool = MaxPooling2D(pool\_size=(2, 2))(prev\_Layer)



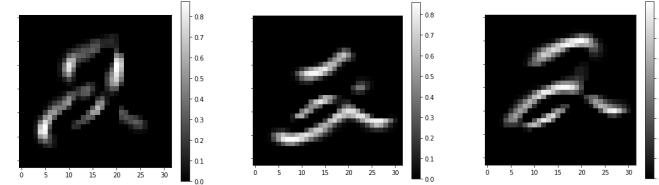
# Stacking layers in keras

Let's say we have one testing input image of size 32x32:

input shape = (None, 32, 32, 1)

conv1 = Conv2D(x, 3, activation = 'relu', stride = 1, padding = 'same')(input)

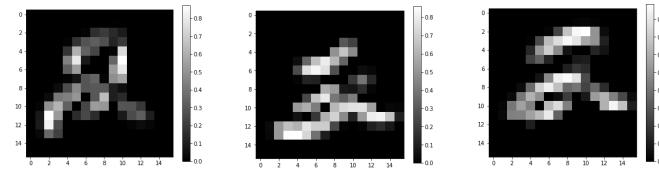
output shape = (None, 32, 32, x)



input shape = (None, 32, 32, x)

pool1 = MaxPooling2D(pool\_size=(2, 2))(conv1)

input shape = (None, 16, 16, x)



input shape = (None, 16, 16, x)

conv2 = Conv2D(y, 3, activation = 'relu', stride = 1, padding = 'same')(pool1)

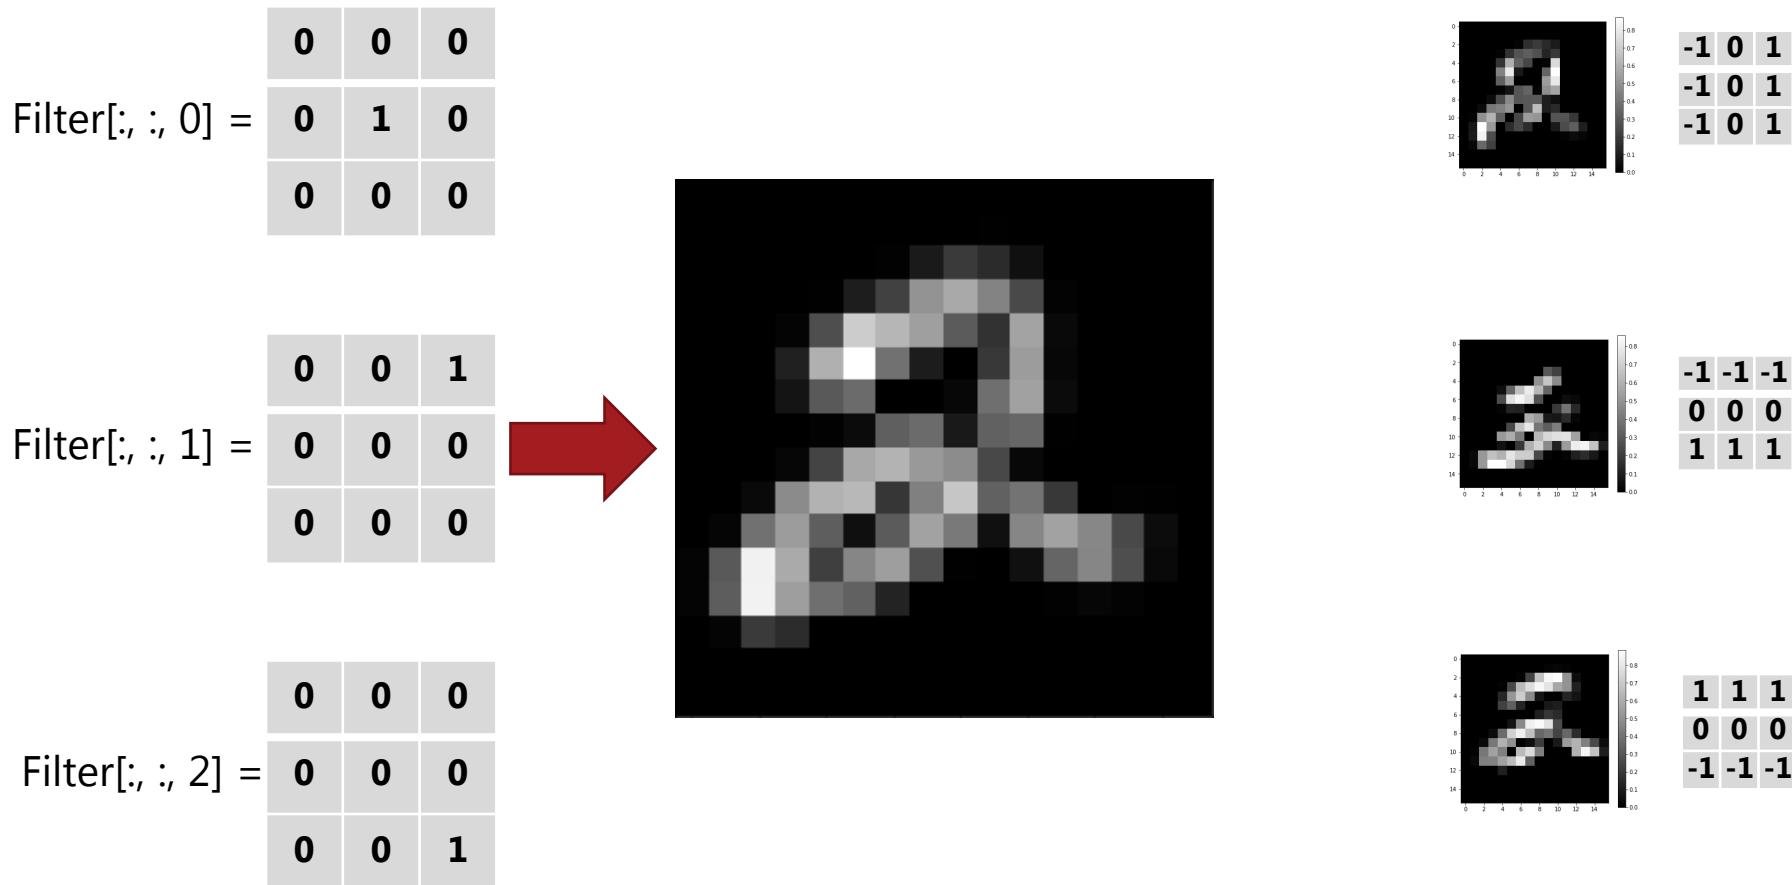
input shape = (None, 16, 16, y)

How do conv2 filters look like?

# Stacking layers in keras

- Each filter in conv2 works ALL x maps of pool1

Filter in conv2 of size (3x3x3)



# L2 regularization

- Regularization makes your network more robust

$$E = \underbrace{\dots \dots \dots \dots}_{\text{plain error}} + \frac{\lambda}{2} * \sum w_i^2 \underbrace{\quad}_{\text{weight penalty}}$$

Weight penalty  
 $0.5\lambda (3 * (-1)^2 + 3 * (1)^2) = 3\lambda$

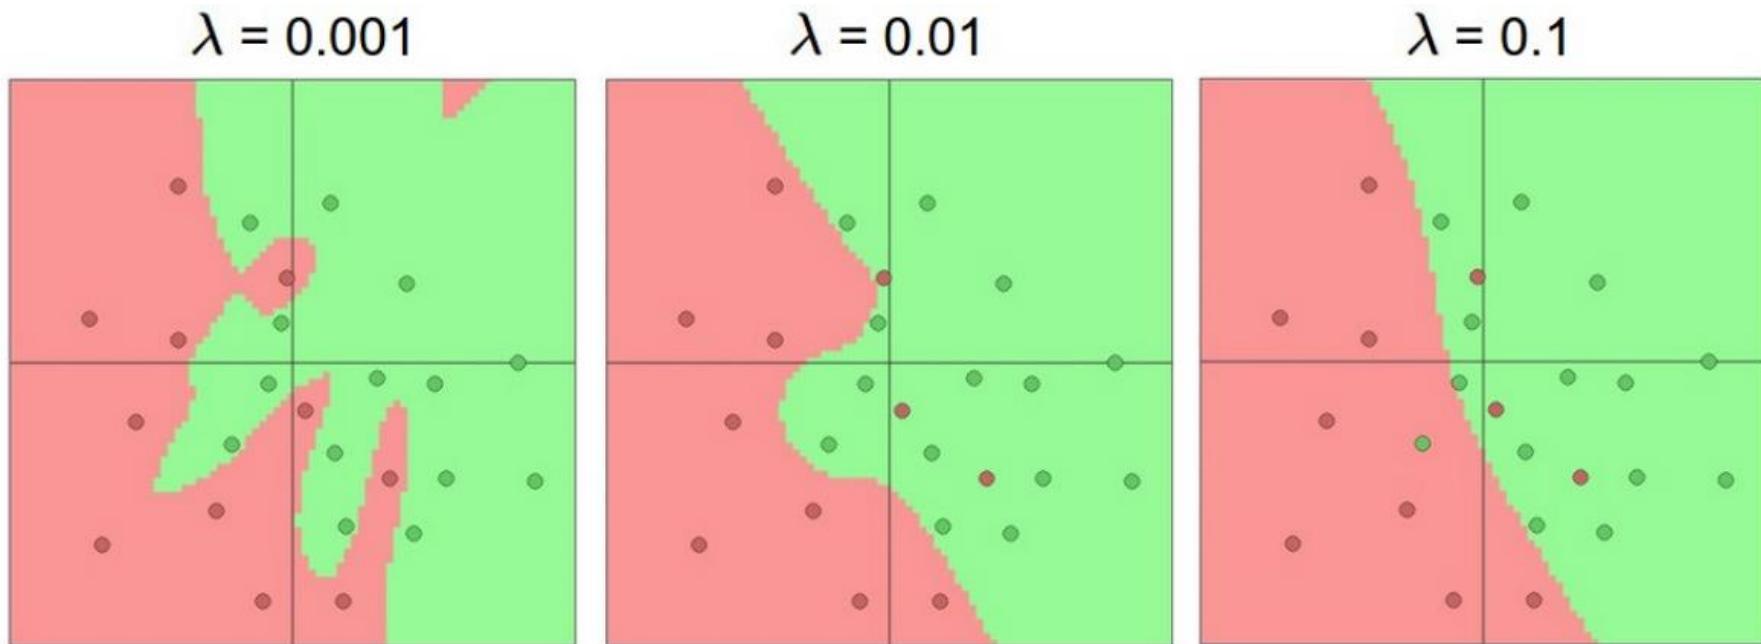
-1	0	1
-1	0	1
-1	0	1

Weight penalty  
 $0.5\lambda (6 * (-0.5)^2 + 3 * (1)^2) = 2.25\lambda$

-0.5	-0.5	1
-0.5	-0.5	1
-0.5	-0.5	1

# L2 regularization

- Regularization makes your network more robust



# Deep learning in image analysis

Deep learning started its expansion from visual object classification

This work can be extrapolated to classification of medical images

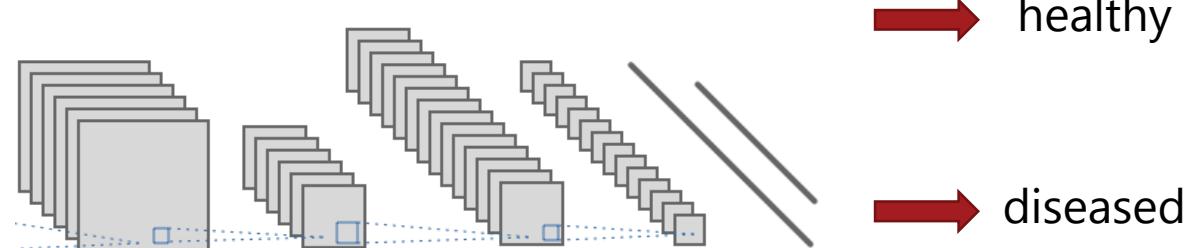


- 1,000 object classes (categories).
- Images:
  - 1.2 M train
  - 100k test.



# Deep learning in medical imaging: image → classes

Classification into healthy and pathological cases

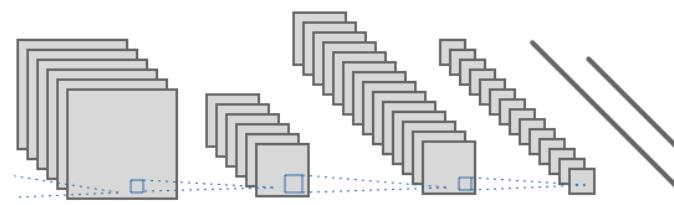


CNN for classification

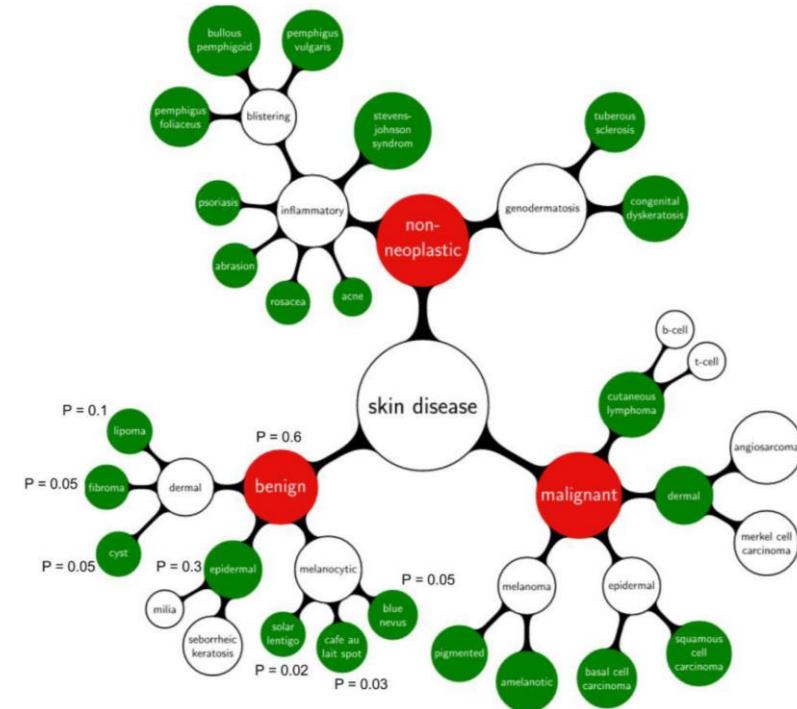
# Deep learning in medical imaging: image → classes

Classification labels may form complex taxonomy:

- Aggregation of rare outputs
- Class imbalance problem



CNN for classification



\*Dermatologist-level classification of skin cancer with deep neural networks

# Deep learning in medical imaging: image → classes

Preservation of the input image size is not needed for classification task

## Deep Visualization Toolbox

[yosinski.com/deepvis](http://yosinski.com/deepvis)

#deepvis



Jason Yosinski



Jeff Clune



Anh Nguyen



Thomas Fuchs



Hod Lipson



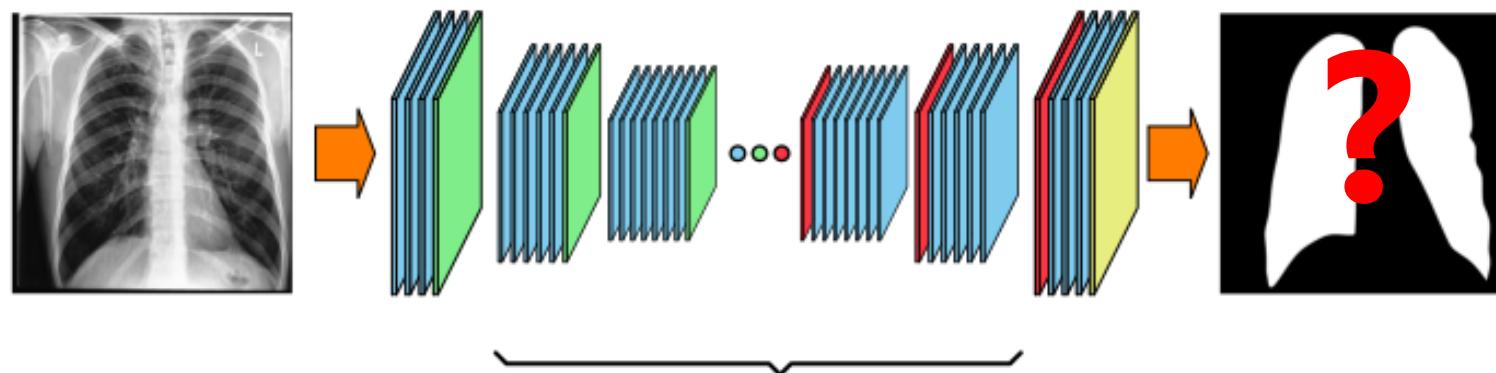
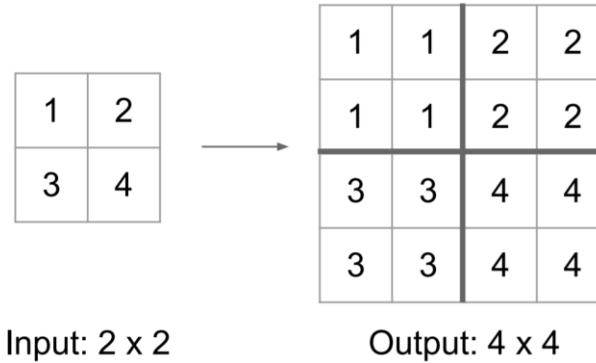
Cornell University



**Jet Propulsion Laboratory**  
California Institute of Technology

# Deep learning in medical imaging: image → image

Upsampling layer

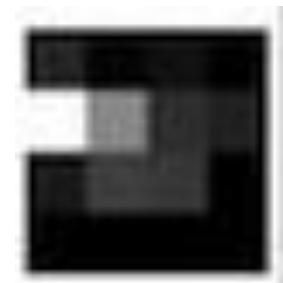
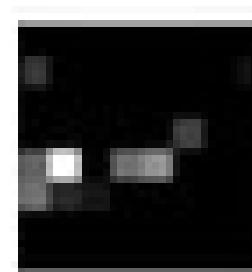
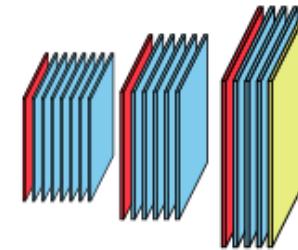
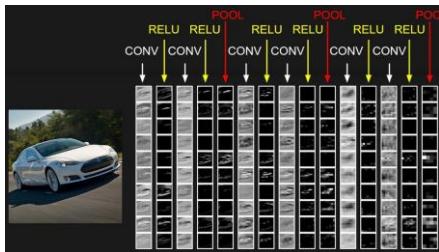


- Conv+BatchNorm+ReLU Layer
- Pooling Layer
- Upsampling Layer

# Deep learning in medical imaging: image → image

What happens with fine image details after upsampling?

Can we expect that network layers after upsampling will recover fine object border?

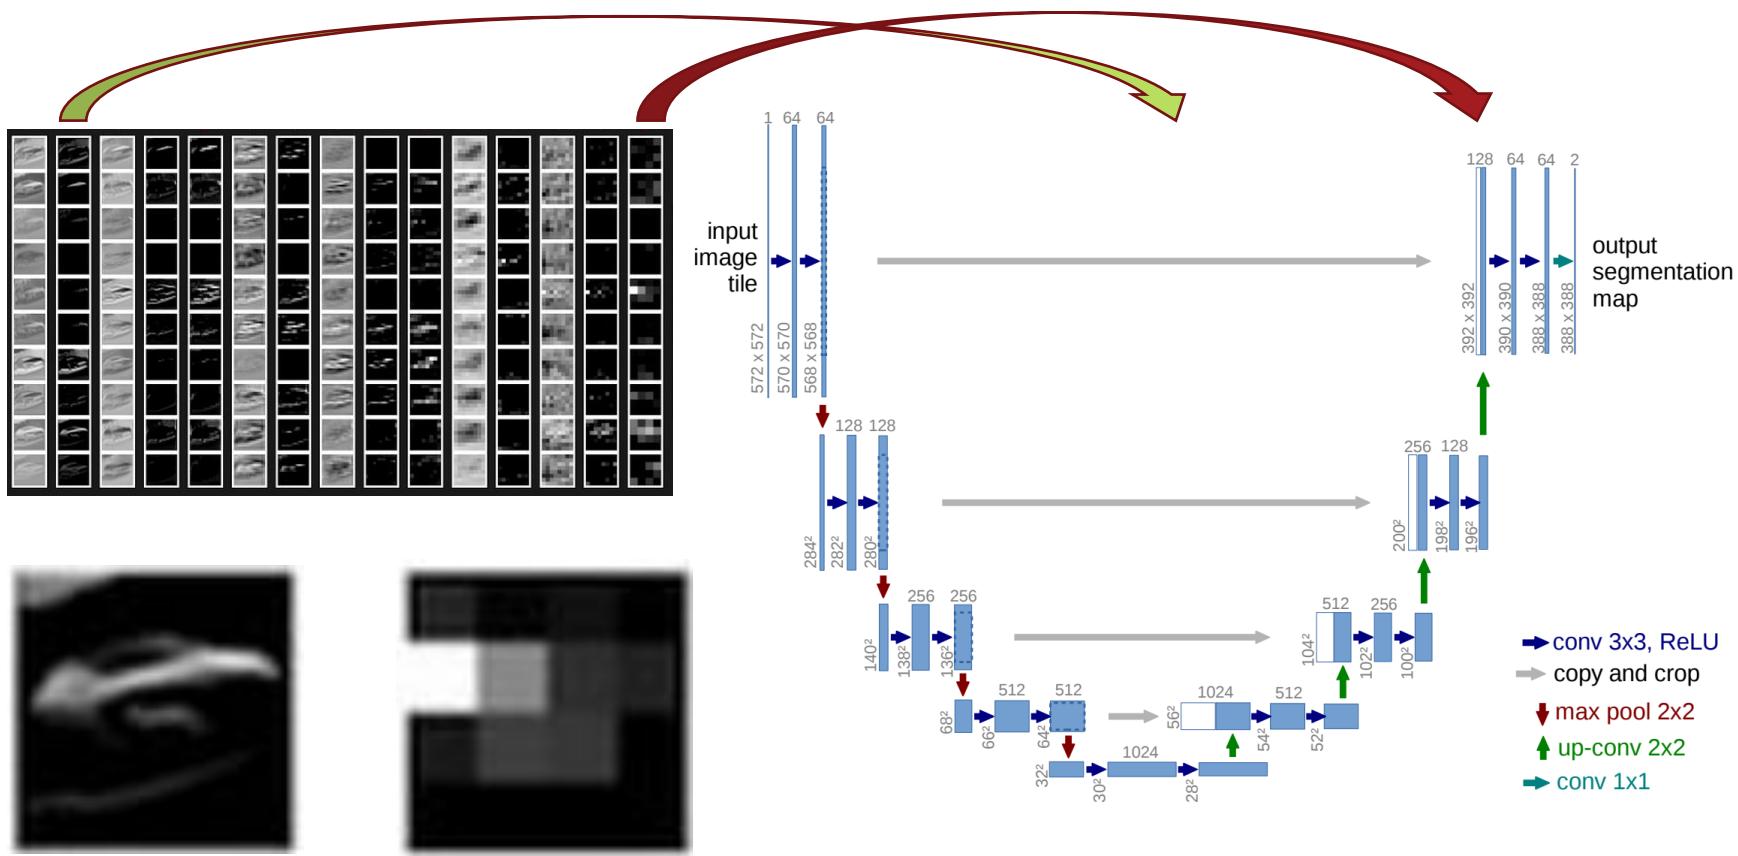


Can we recover side window borders from feature maps that look like this?

# Deep learning in medical imaging: image → image

UNet architecture for segmentation:

- Skip connections pass fine details from early layers to late layers

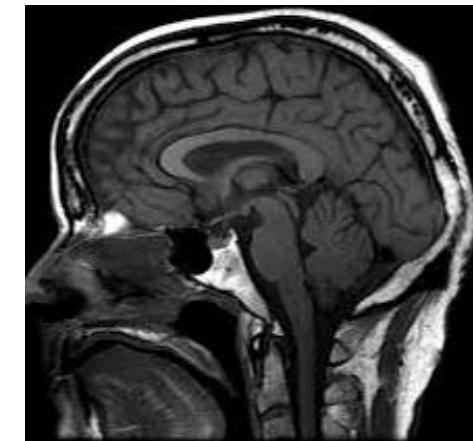
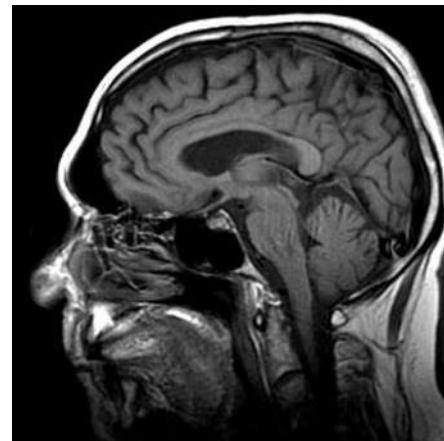


Can we recover side window borders from  
feature maps that look like this?

# Medical image analysis vs Computer vision

Medical images are better standardized:

- Similar imaging protocols (resolution, slice thickness, etc.)
- Standardized patient positioning



Brain MRIs are way more similar

Natural images of the same object look  
very different

# Medical image analysis vs Computer vision

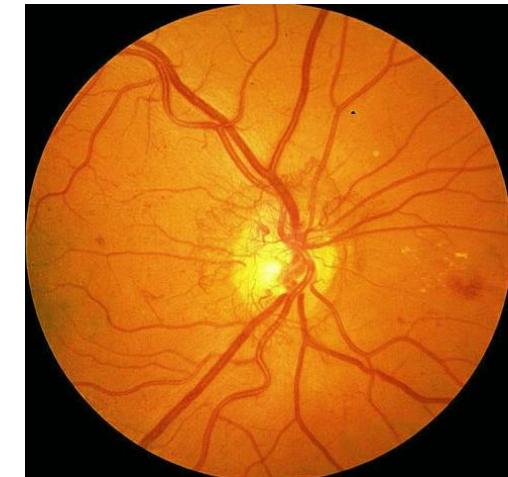
The difference between natural images of different types is also better pronounced than for medical images



Cats are very different from planes



Healthy retina

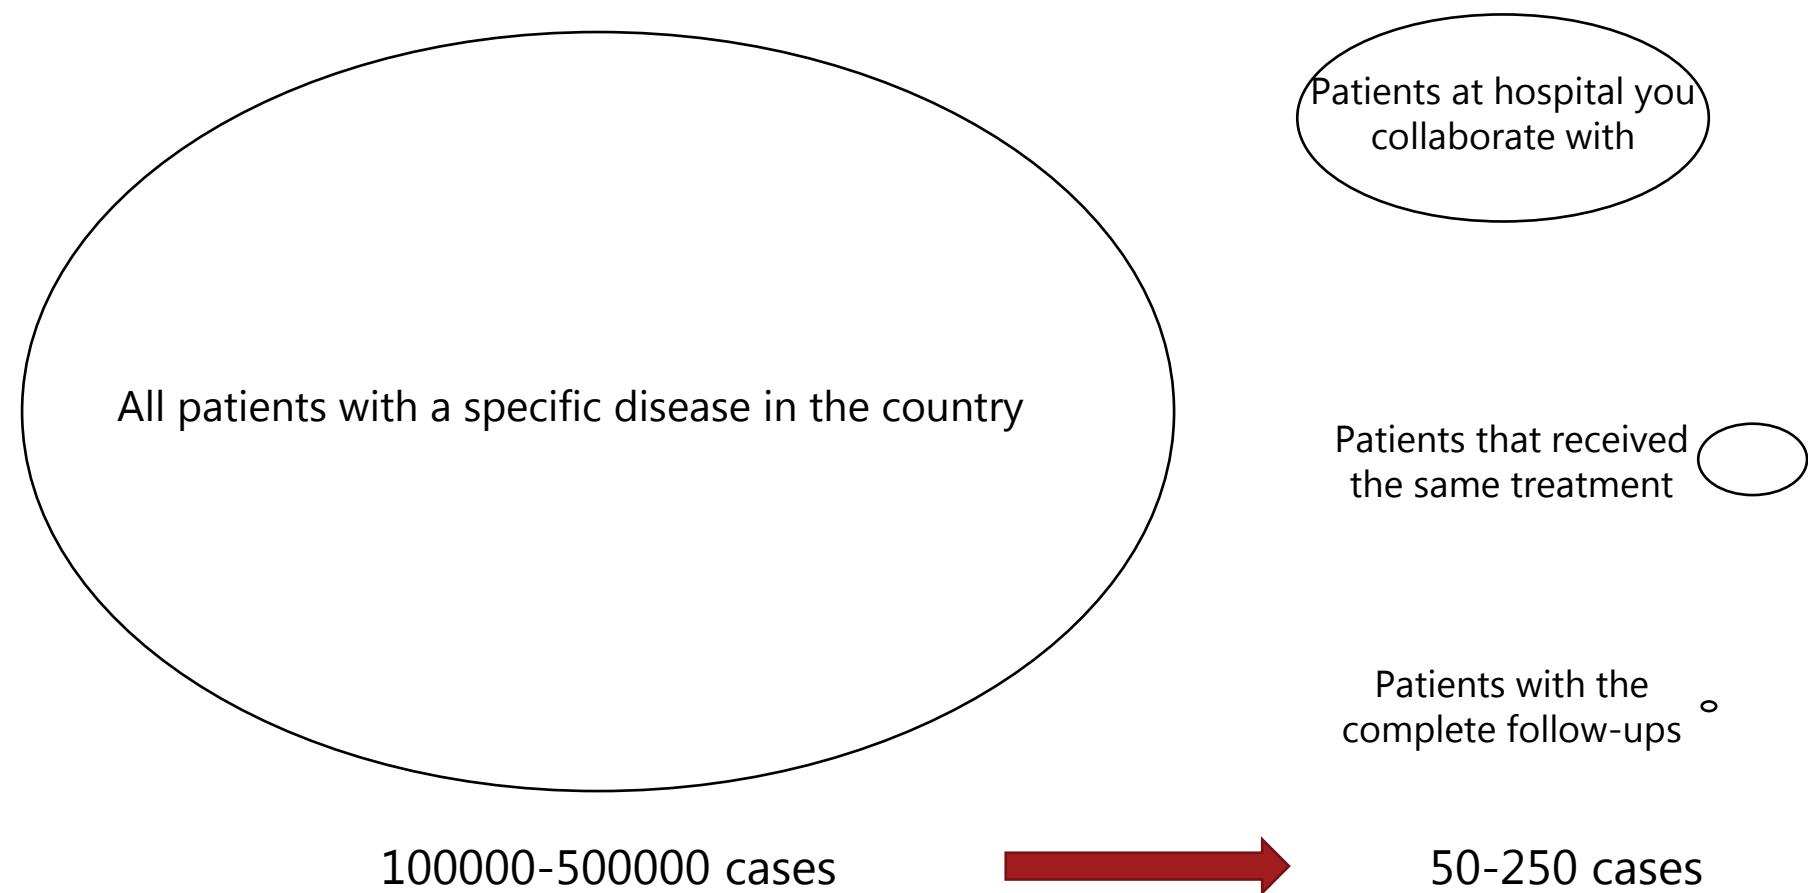


Diabetic retina

# Medical image analysis vs Computer vision

The number of available medical images is by several orders of magnitude lower than of natural images:

- Google 'images of cats' and get 25000 images in the first link



# Training-validation-testing

Training images – training features of the model

Validation images - training hyperparameters of the model

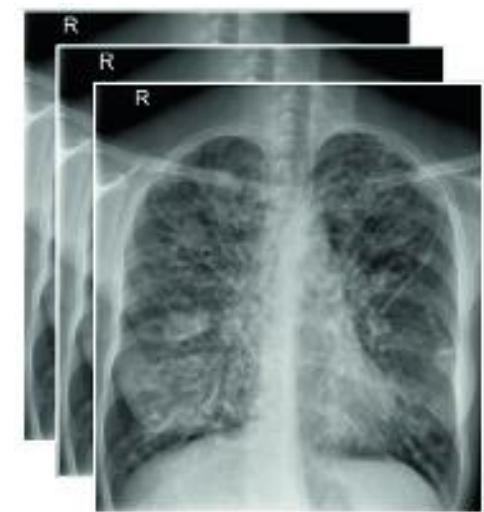
Testing images – used at the very end, when the model is locked and will not be anymore edited



Training images



Validation images



Testing images

# Loss: Binary cross-entropy

```
model.compile(optimizer = Adam(lr = 1e-4), loss = 'binary_crossentropy',  
metrics = ['accuracy'])
```

$$H_p(q) = -\frac{1}{N} \sum_{i=1}^N y_i \cdot \log(p(y_i)) + (1 - y_i) \cdot \log(1 - p(y_i))$$

True label	Prediction	Loss
1	0.3	1.2

$$-(1 * \ln(0.3) + (1-1)*\ln(1-0.3)) = \\ -\ln(0.3) = 1.204$$

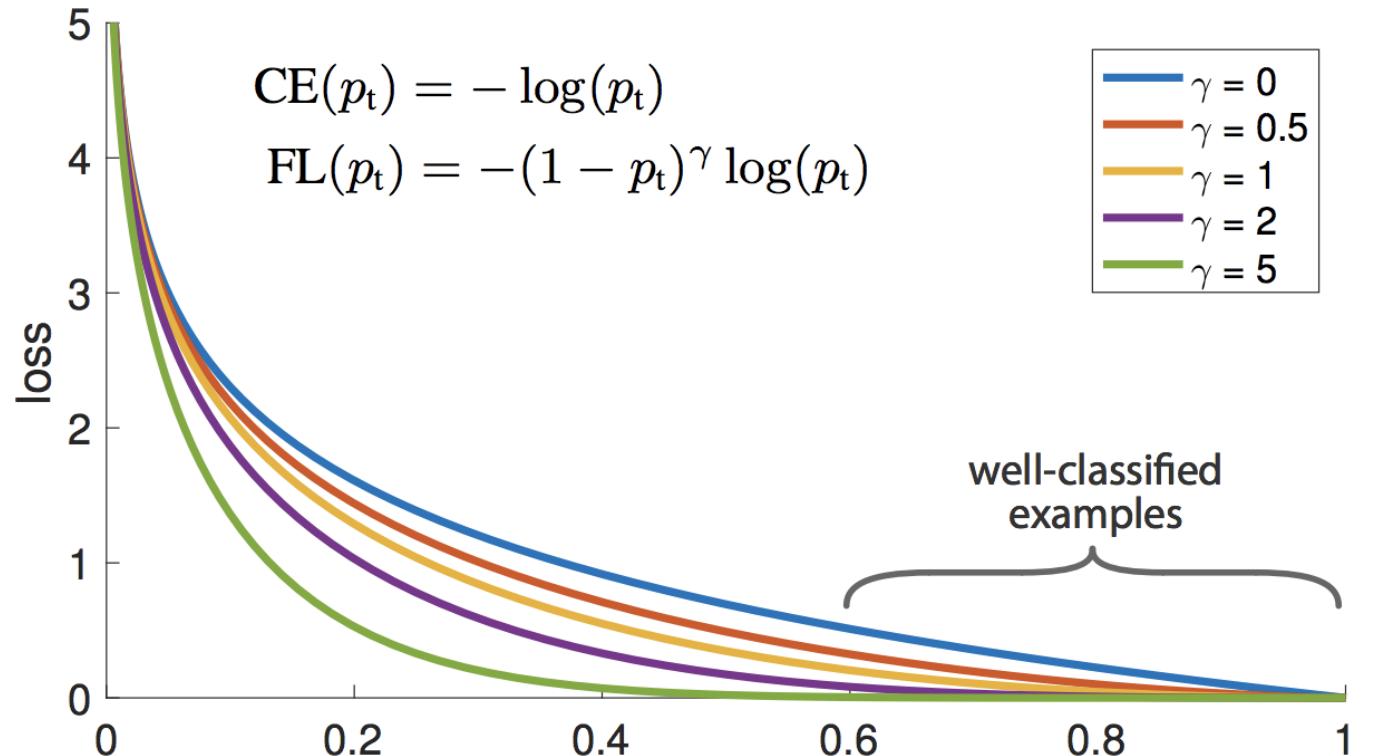
True label	Prediction	Loss
1	0.9	0.1

$$-(1 * \ln(0.9) + (1-1)*\ln(1-0.1)) = \\ -\ln(0.9) = 0.105$$

# Focal loss

Class imbalance (large size of the background) is a problem for binary cross-entropy

Focal loss focuses on difficult-to-classify cases



# Dice loss

We can use Dice coefficient to compute loss

No problem with background size

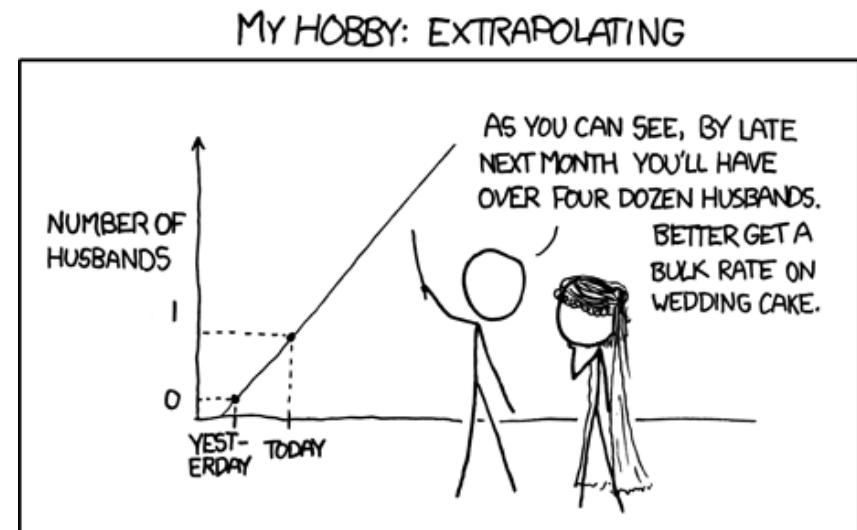
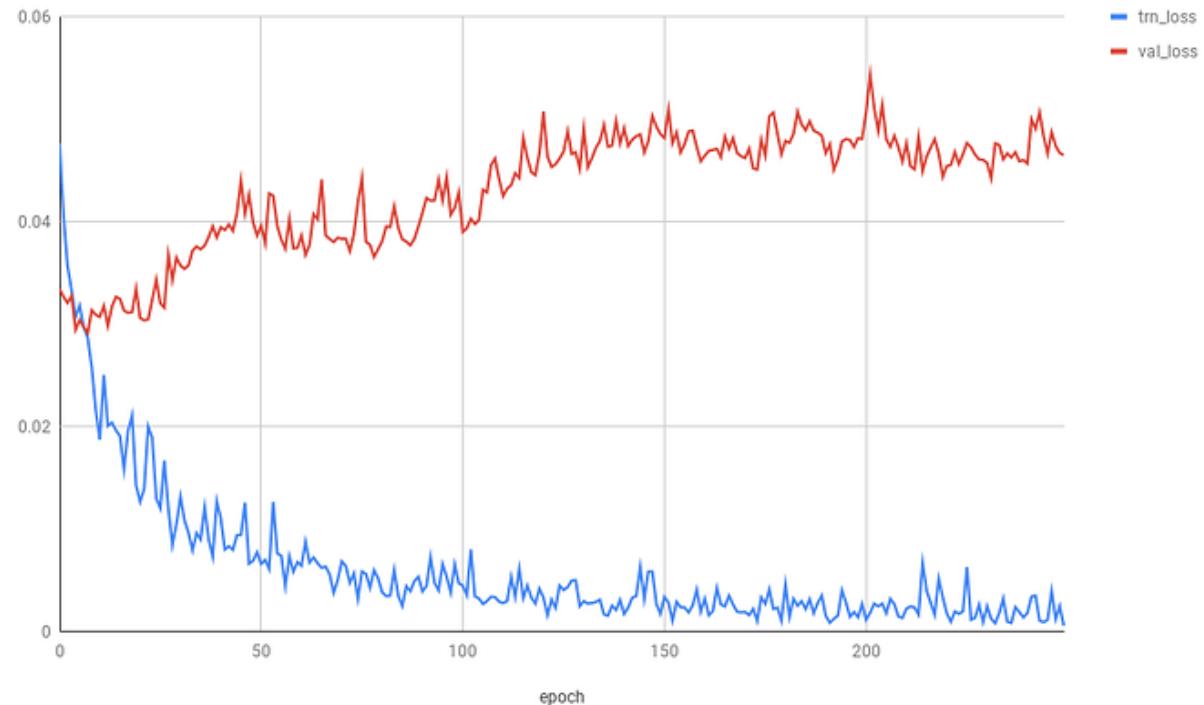
```
def __dice_coefficient(y_true, y_pred, smooth=1.):
    y_true_f = keras.flatten(y_true)
    y_pred_f = keras.flatten(y_pred)
    intersection = keras.sum(y_true_f * y_pred_f)
    return (2. * intersection + smooth) / (keras.sum(y_true_f) + keras.sum(y_pred_f) + smooth)

def __dice_coefficient_loss(y_true, y_pred):
    return 1 - NN_Architectures.__dice_coefficient(y_true, y_pred)
```

# Overfitting

Ways to reduce overfitting:

- Reducing the number of features in convolution layers
- Regularization
  - Dropout
  - L2 regularization
- Data augmentation



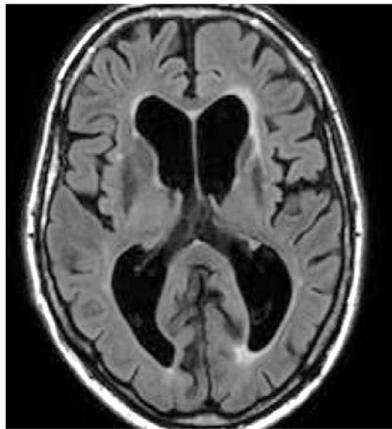
# Questions?

# Medical image augmentation

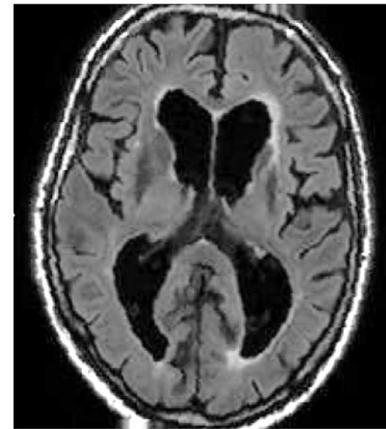
Image augmentation is useful:

- To enrich training database
- To address the class imbalance problem

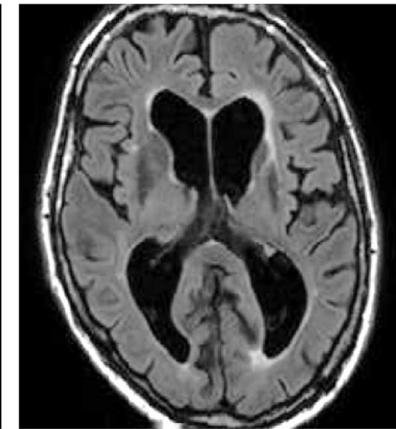
Affine transformations (anatomical restrictions may apply)



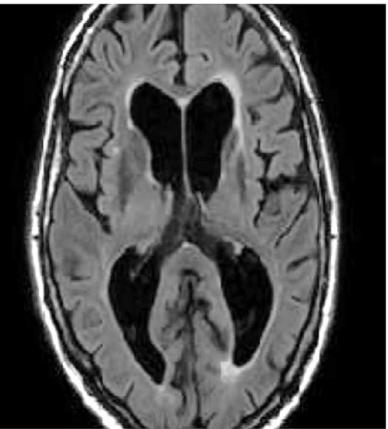
original slice



rotation



shear mapping



scaling

# Medical image augmentation

Intensity transformations:

- Noise
- Blur
- Contrast
- etc



Original image



Random brightness



Random gamma value



Random Gaussian noise

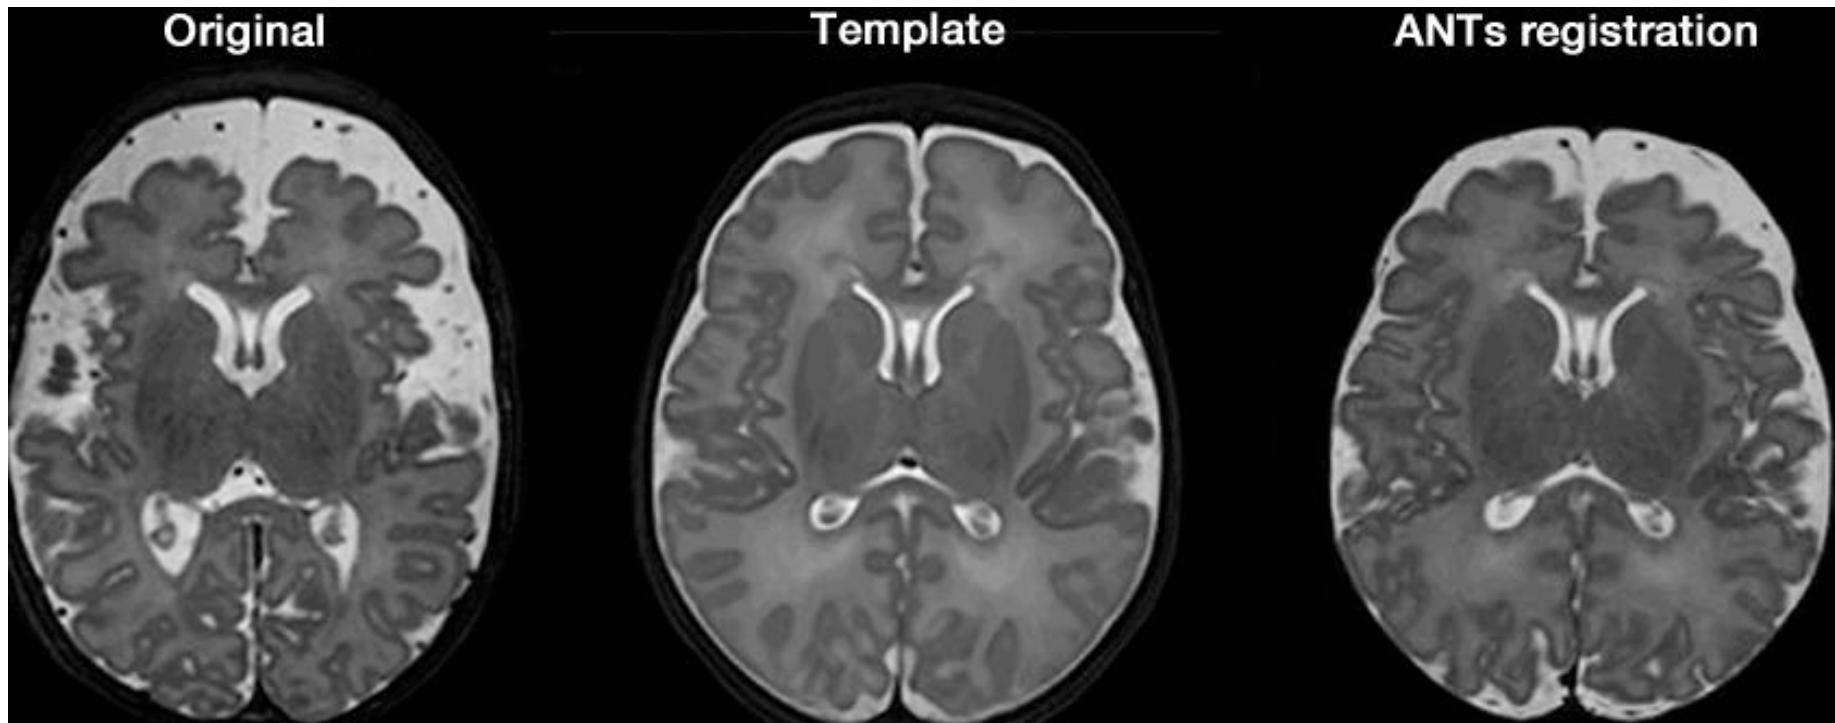


Random Gaussian blur

# Medical image augmentation

Registration-based augmentation:

- Very useful for medical in contrast to natural images



# Medical image augmentation

Style transform simultaneously preserves both:

- Content of the target image  
(neuron activations – critical-to-network information that leads to content image classification)
- Style from the style image  
(pairwise activation covariances – the neurons that get activated together, i.e. consistent appearance patterns)



# Medical image augmentation

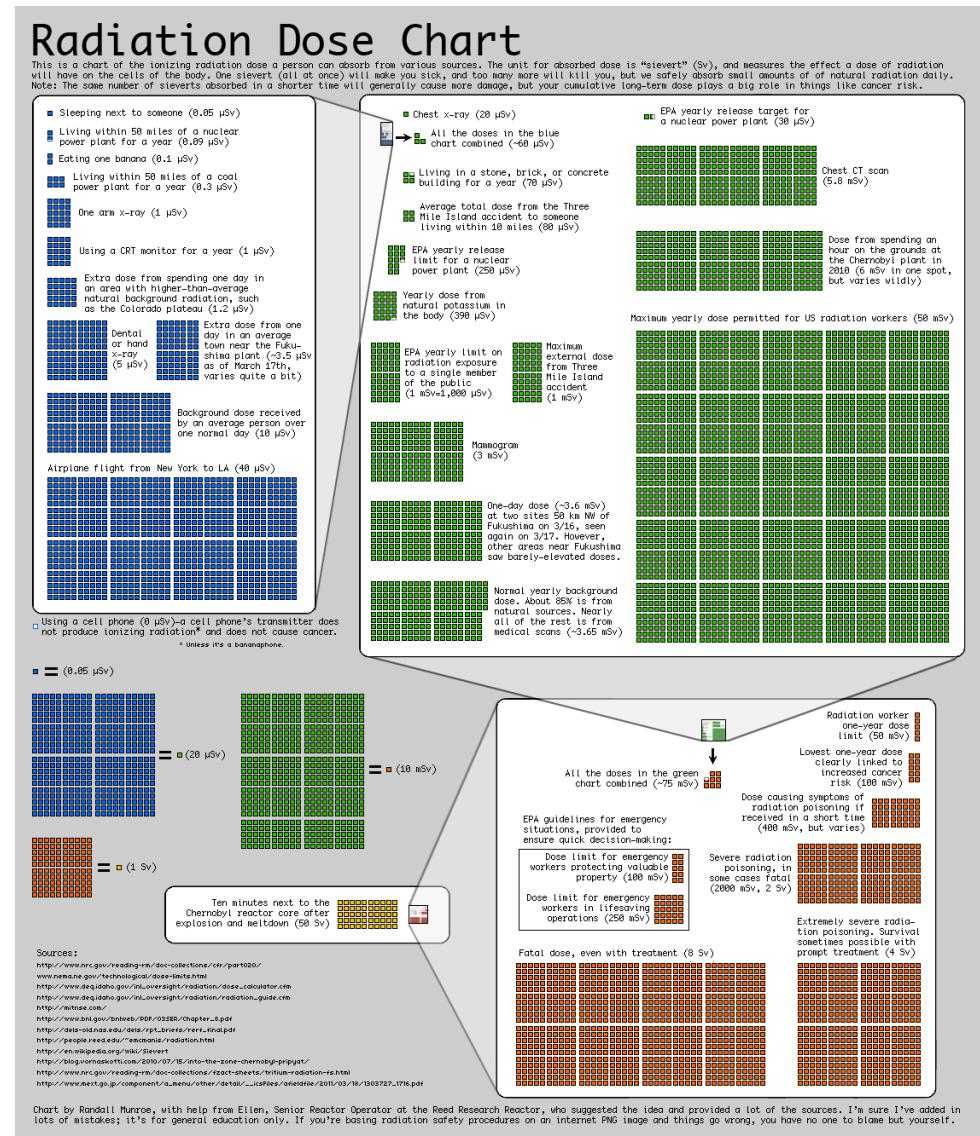
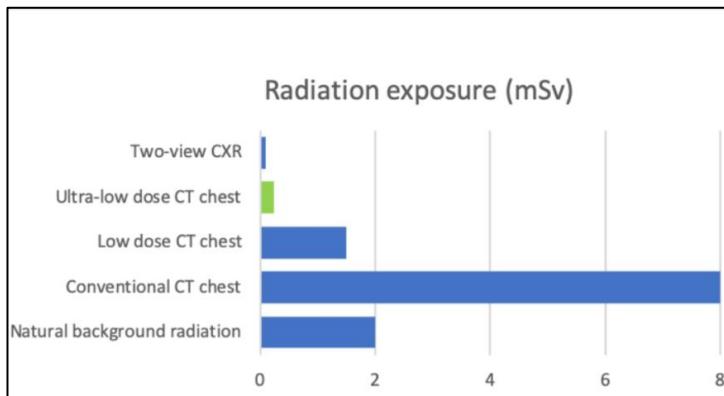
Style transform for medical image augmentation



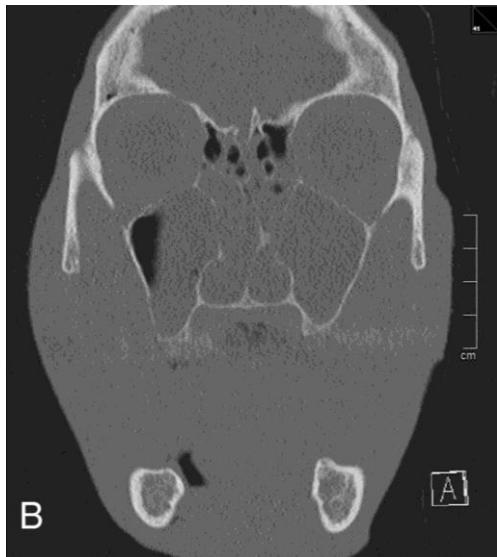
What if content image is for a healthy case,  
while the style image is for a pathological case?

# Image enhancement: dose/quality problem

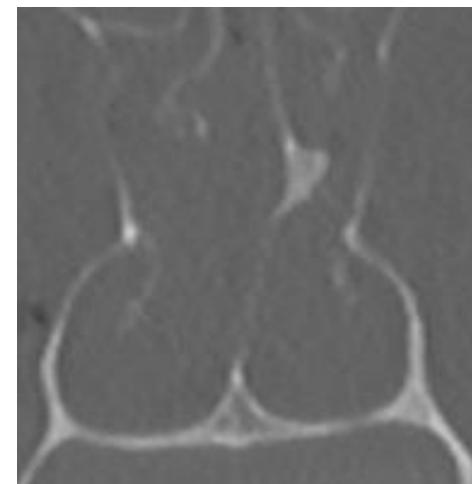
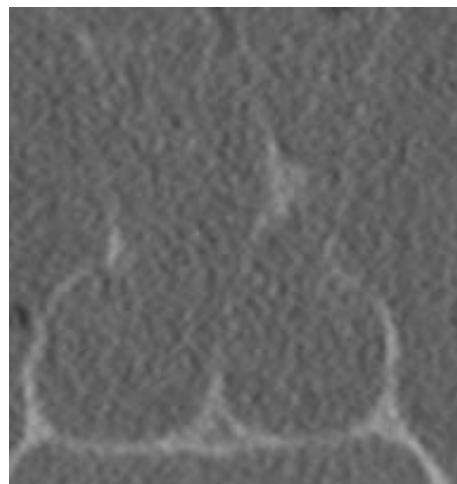
- The dose of CT imaging has dramatically reduced during the last 10 years
- The number of CTs per capita increased 3 times in the last 10 years



# Image enhancement: dose/quality problem

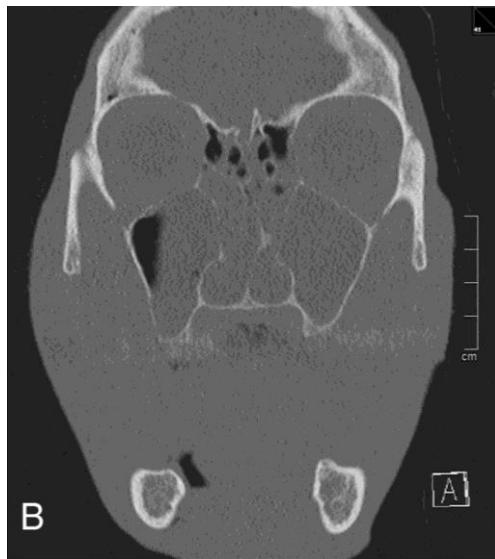


Dose reduction results in  
reduction of the image quality

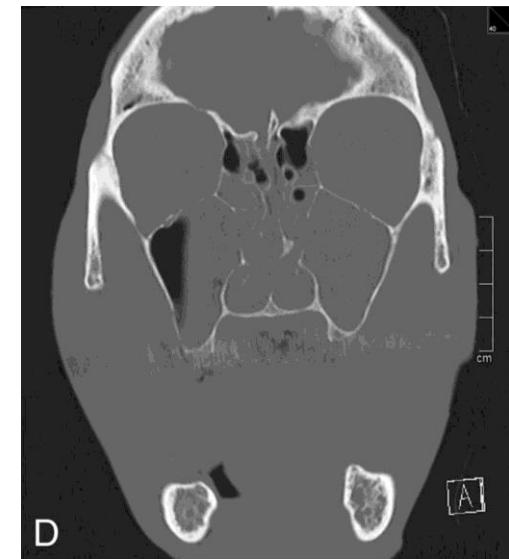
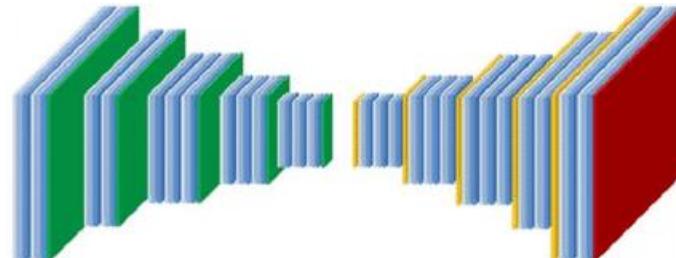


# Image enhancement: dose/quality problem

- Input: low-resolution images
- Output: corresponding high-resolution images
- The input can be generated from high-resolution images by image blurring and adding noise

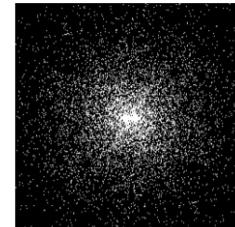


Encoder-decoder NN

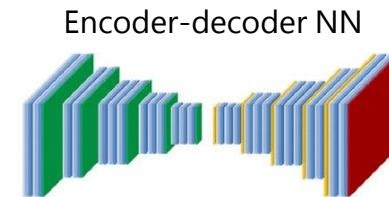


# Image enhancement

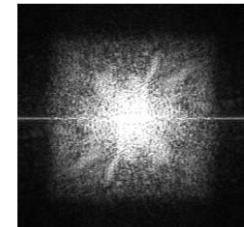
We can use raw data instead of reconstructed medical images



Undersampled spectrum

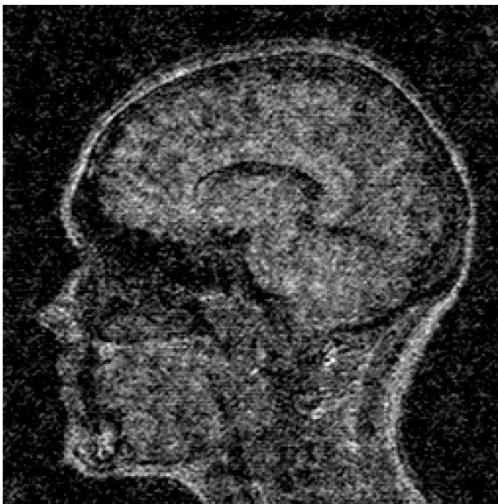


Encoder-decoder NN

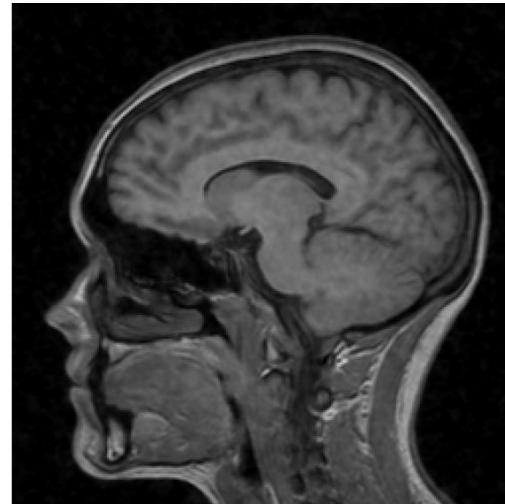


High-resolution spectrum

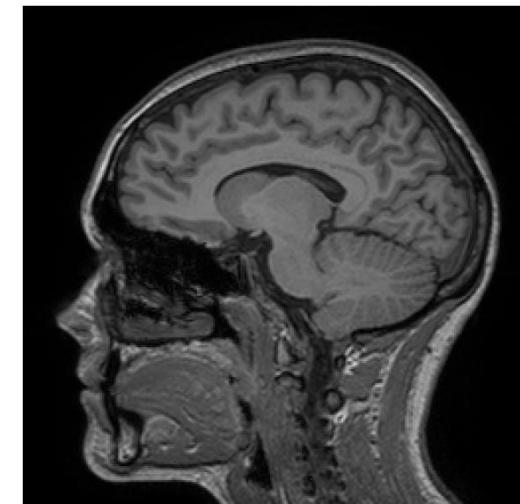
MRI from  
undersampled raw data



Reconstructed MRI



Ground truth



# Image enhancement: improved resolution

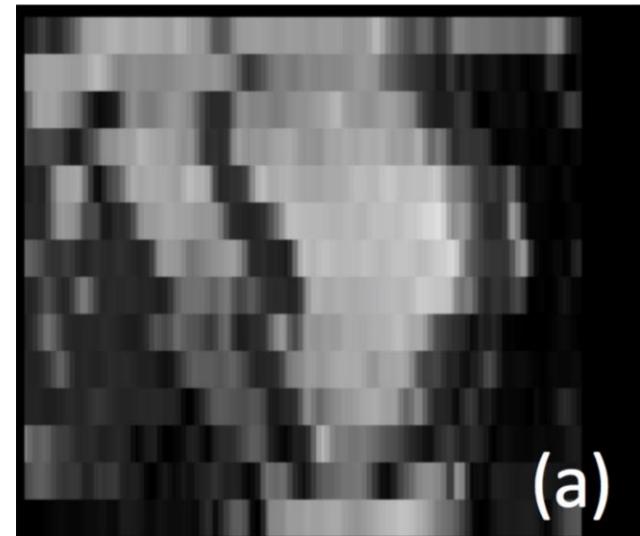
3D imaging takes time:

- Patients move slightly during acquisition
- Movement of some organs cannot be controlled, e.g. heart
- More motion –more blur



How to image moving organs?

- Reduce overall image quality
- Sacrifice one image dimension

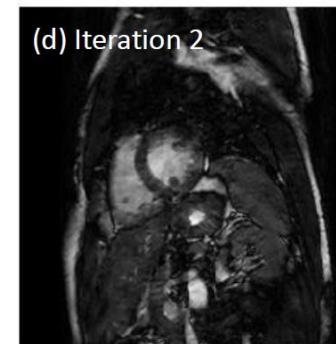
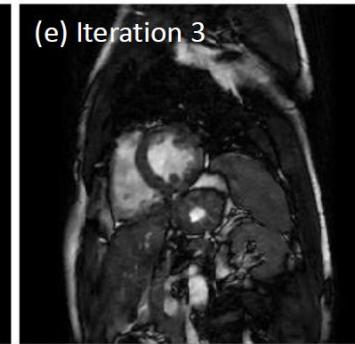
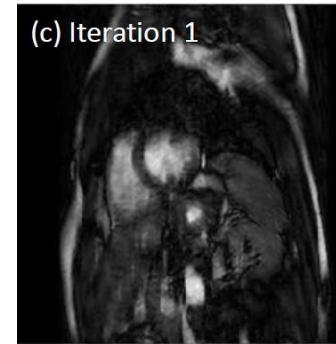
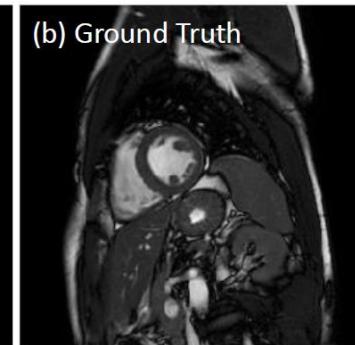
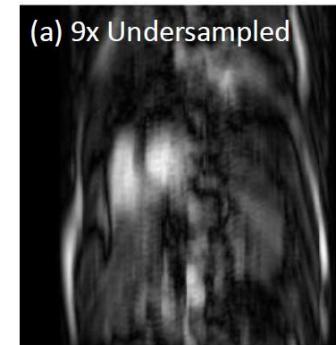
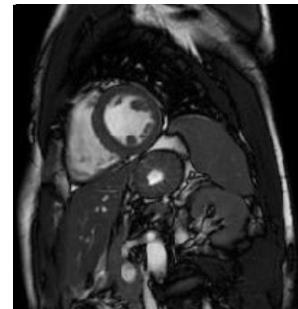
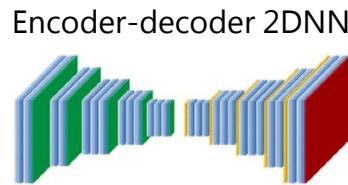
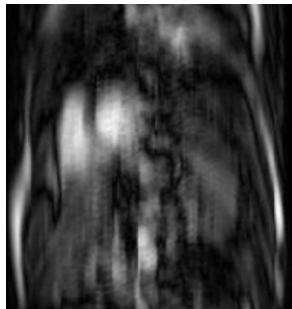


One option – sacrifice one dimension

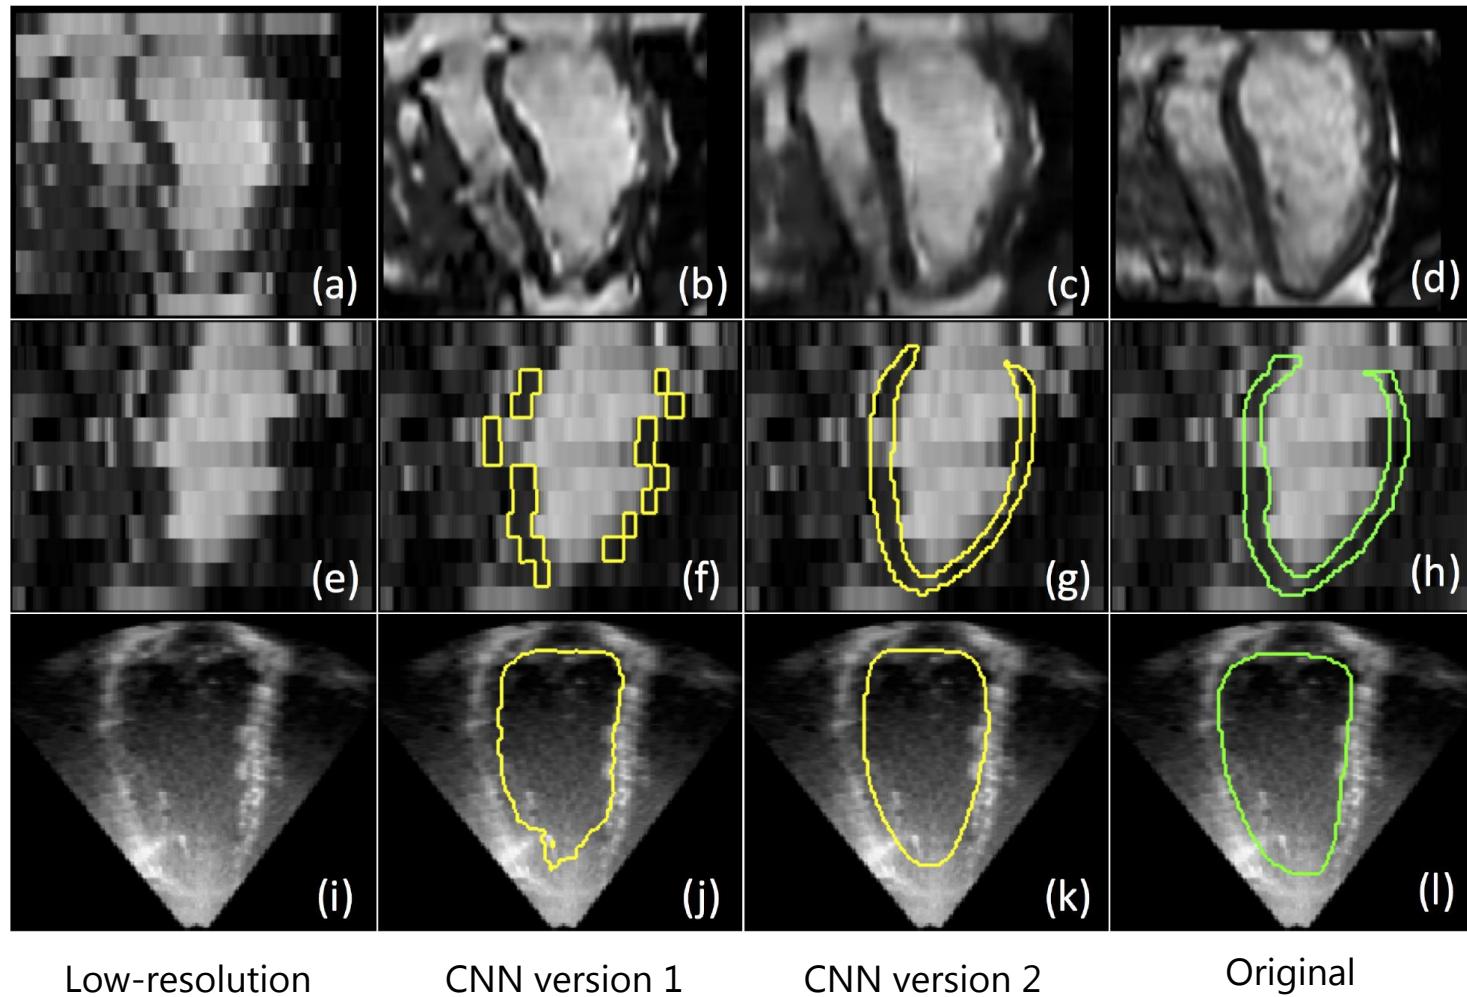
# Image enhancement : improved resolution

One solution:

- Acquire axially(coronally/sagittally)-reconstructed images or high-resolution images
- Artificially vandalize them
  - Undersample spectrum
  - Reduce resolution in one dimension
- Train the network to recover the original images



# Image enhancement: improved resolution

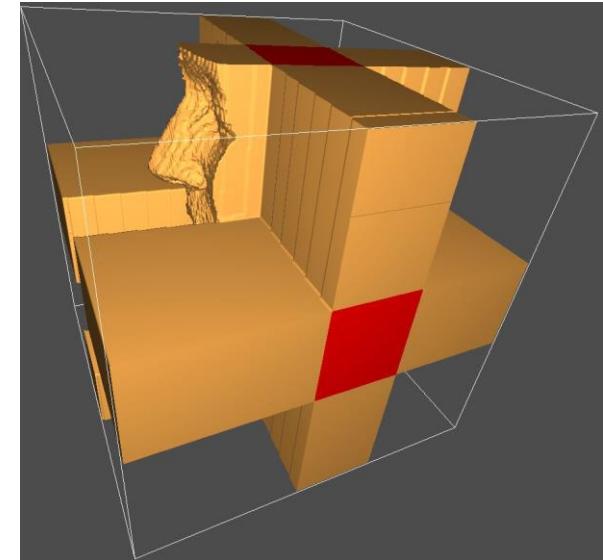
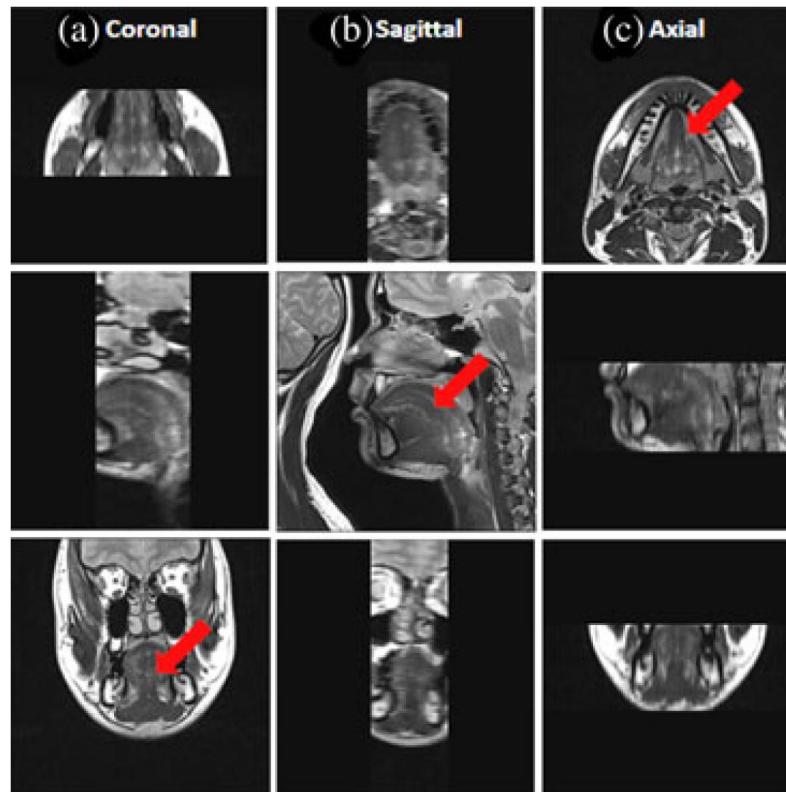


\* Anatomically Constrained Neural Networks (ACNN): Application to Cardiac Image Enhancement and Segmentation

# Image enhancement: improved resolution

Tongue image analysis:

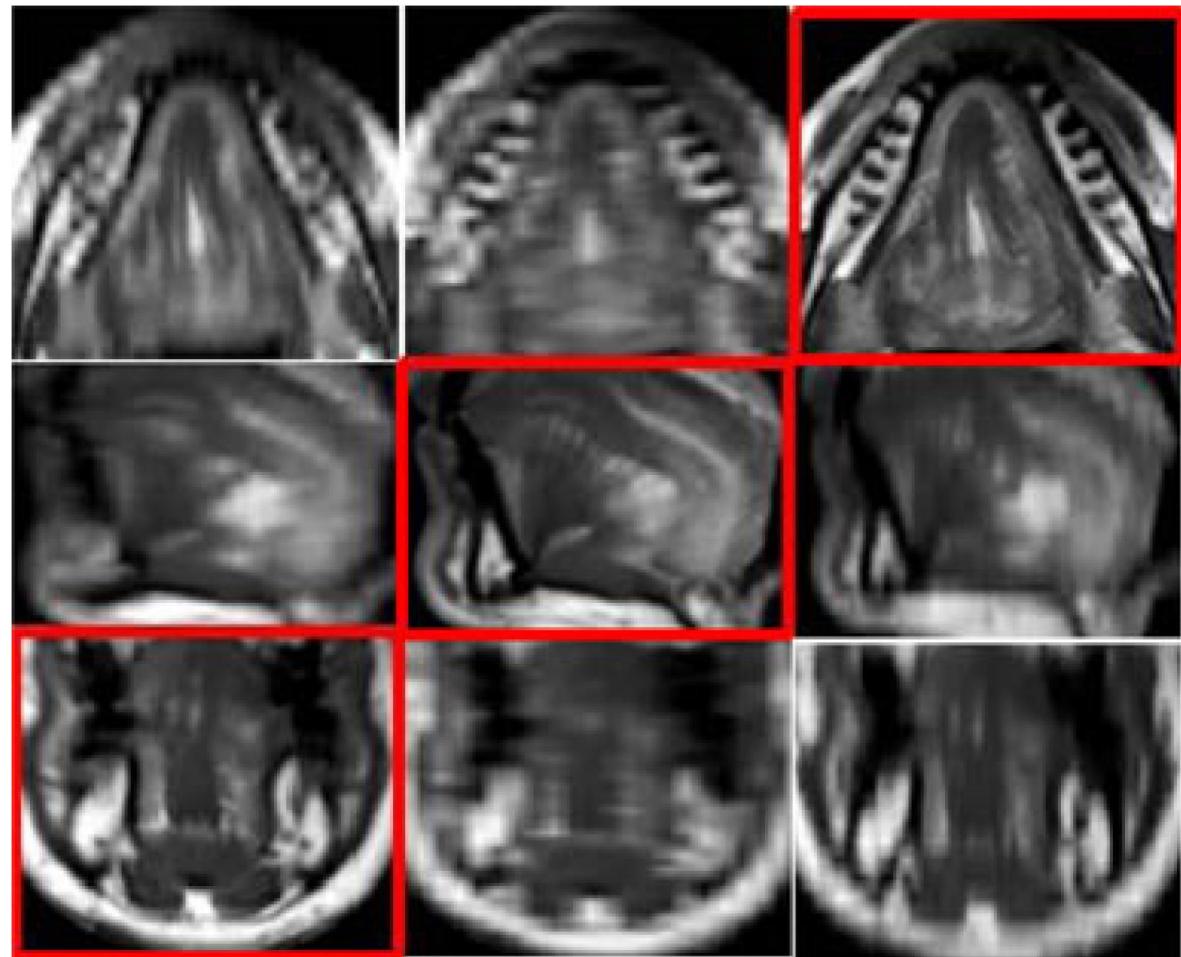
- Training: we have 3 orthogonal images that, however, do not cover the whole 3D volume



# Image enhancement: improved resolution

Tongue image analysis:

- Use image registration to align images
- Merge aligned images into one
- Train CNN to recover the merged image from individual orthogonal image



\* Reconstruction of high-resolution tongue volumes from MRI

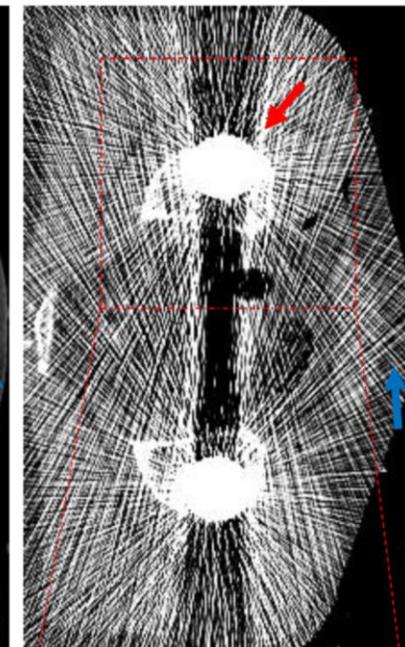
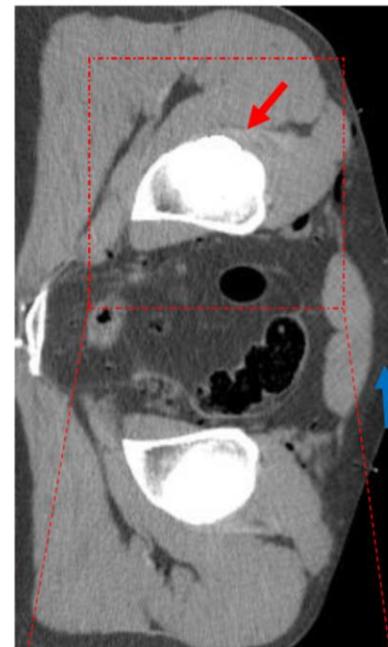
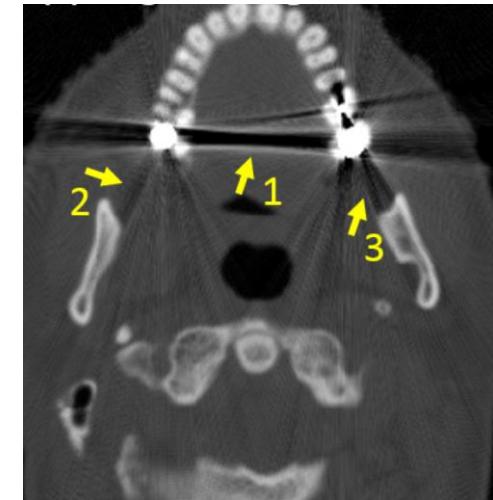
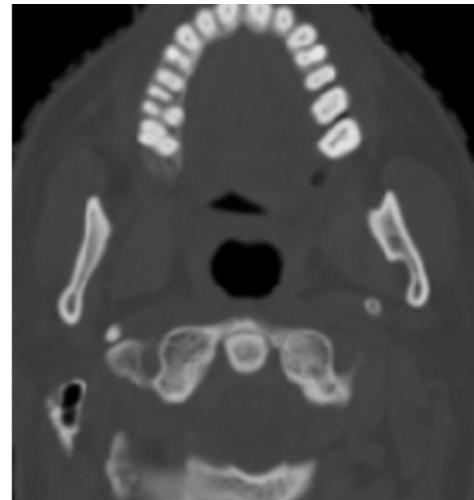
# Image enhancement: artifacts removal



# Image enhancement: artifacts removal

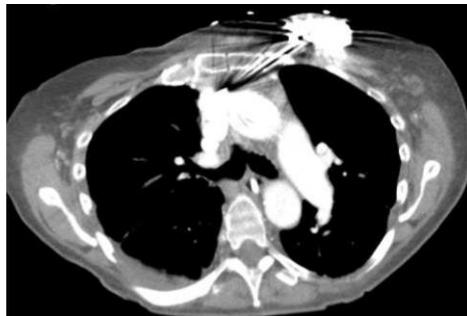
Image metal artifacts:

- Dental metal crowns
- Femoral head implants
- Pedicle screws
- etc.



# Image enhancement: artifacts removal

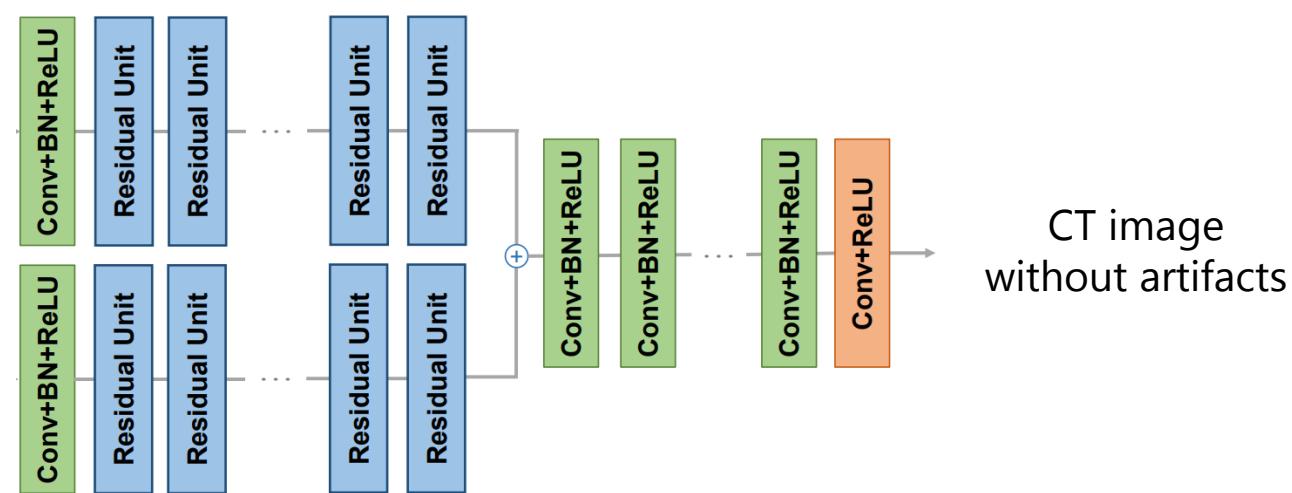
We can use only CT, only sinogram or both sinogram and CT image with artifacts



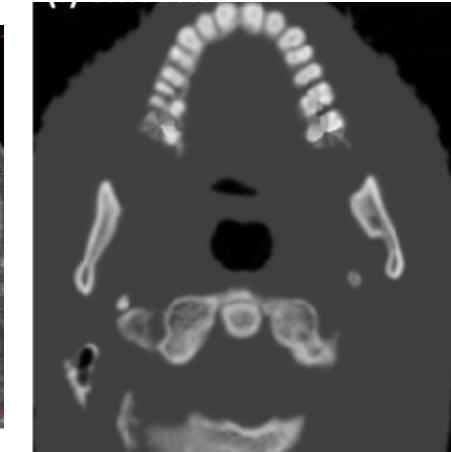
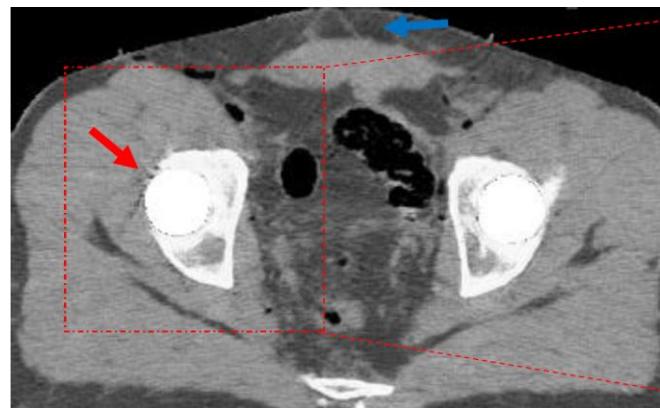
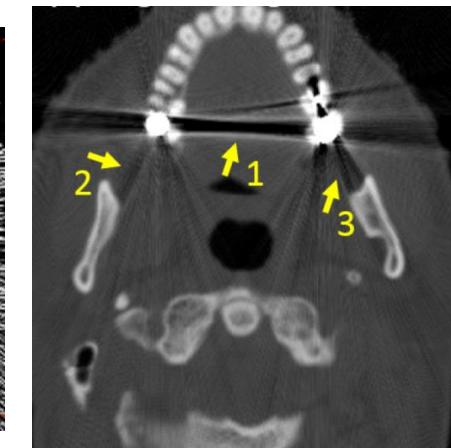
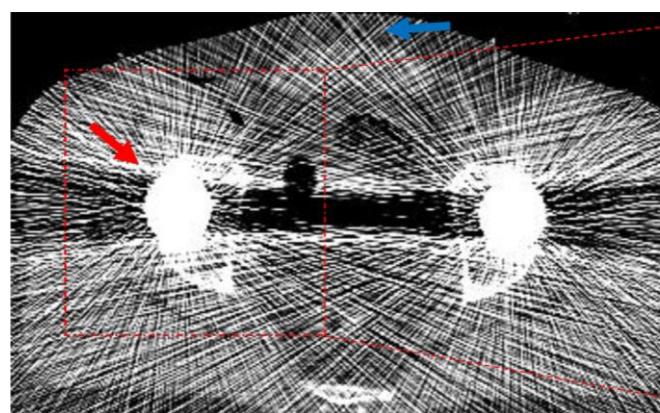
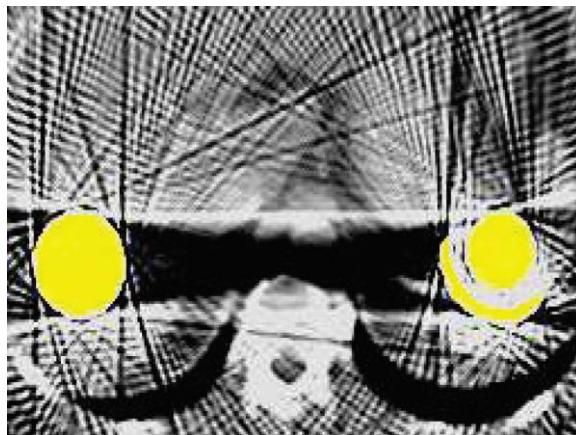
## Original Image



## Sinogram



# Image enhancement: artifacts removal



\*DuDoNet: Dual Domain Network for CT Metal Artifact Reduction  
Deep Neural Network for CT Metal Artifact Reduction with a Perceptual Loss Function  
Convolutional Neural Network based Metal Artifact Reduction in X-ray Computed Tomography

# Image synthesis

Artificially generate one image modality from another on for the same patient

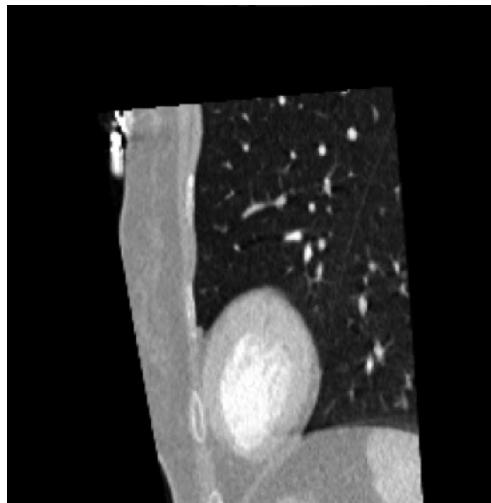
How can this be useful?



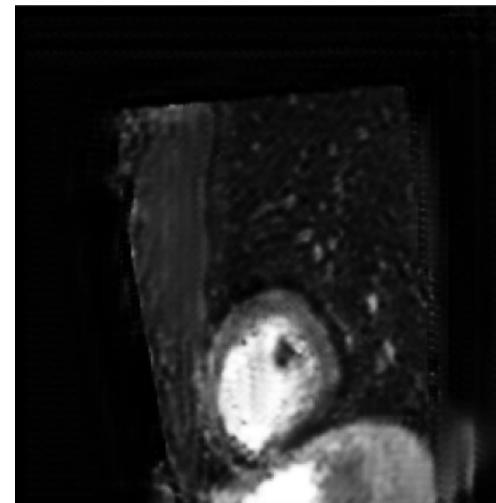
# Image synthesis

Example:

- Let's say we have an algorithm for segmentation /classification /detection that requires both CT and MR
- A new patient has only MR
- Using image synthesis we can generate artificial CT



real CT



synthetic MR

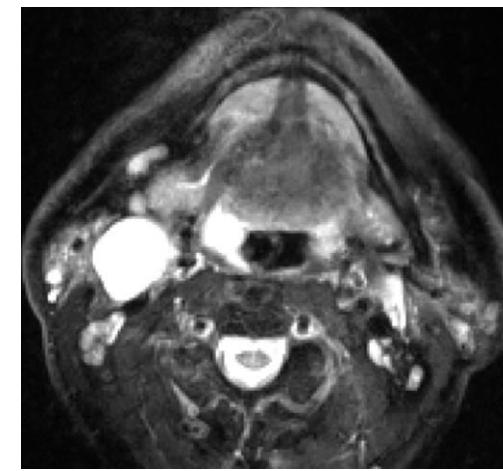
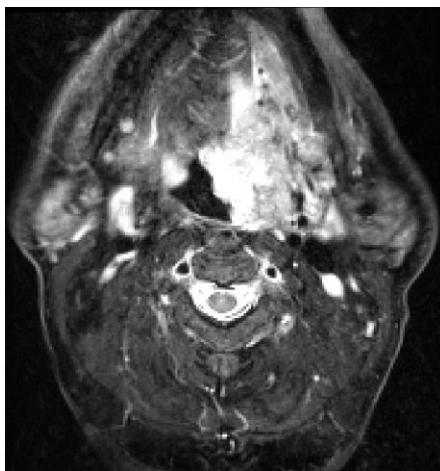
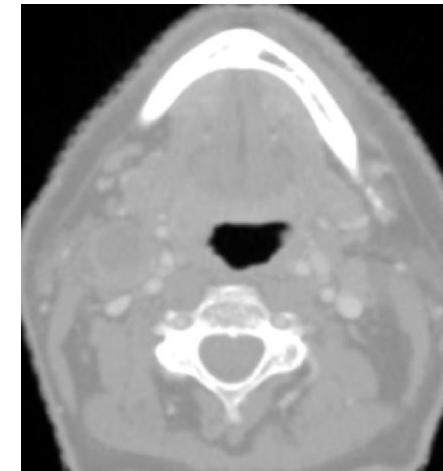
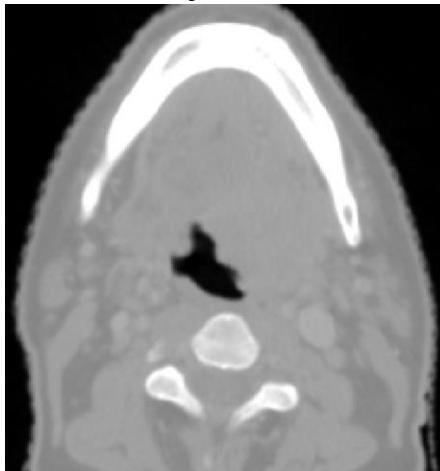


ventricle mask

# Segmentation through synthesis

Example:

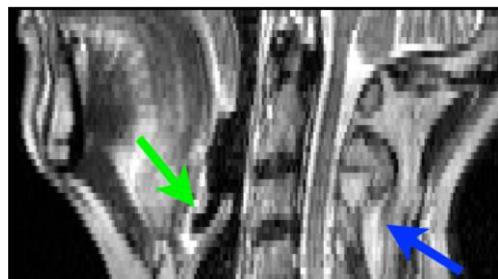
- Most of head and neck tumors are very poorly visible on CT images but very visible on MR T2 images



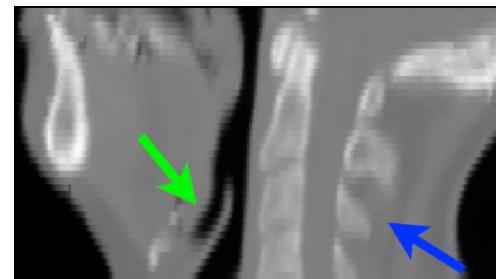
# Segmentation through synthesis

- Register MRs to CT images for healthy subjects
- Train network to synthesize MR images from healthy CT
- Apply network to synthesize MR from CT with tumor. This MR will have no tumor because network does not know the tumor appearance and tumor is not very visible on CT
- The difference between original and synthetic MR will be the tumor

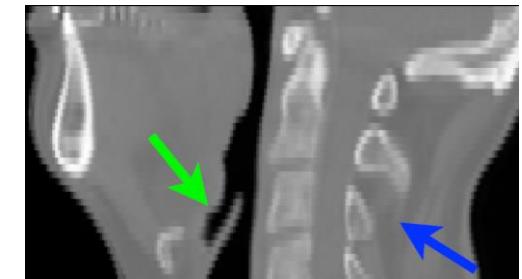
Head-and-neck CT synthesis from MR images



Original MR



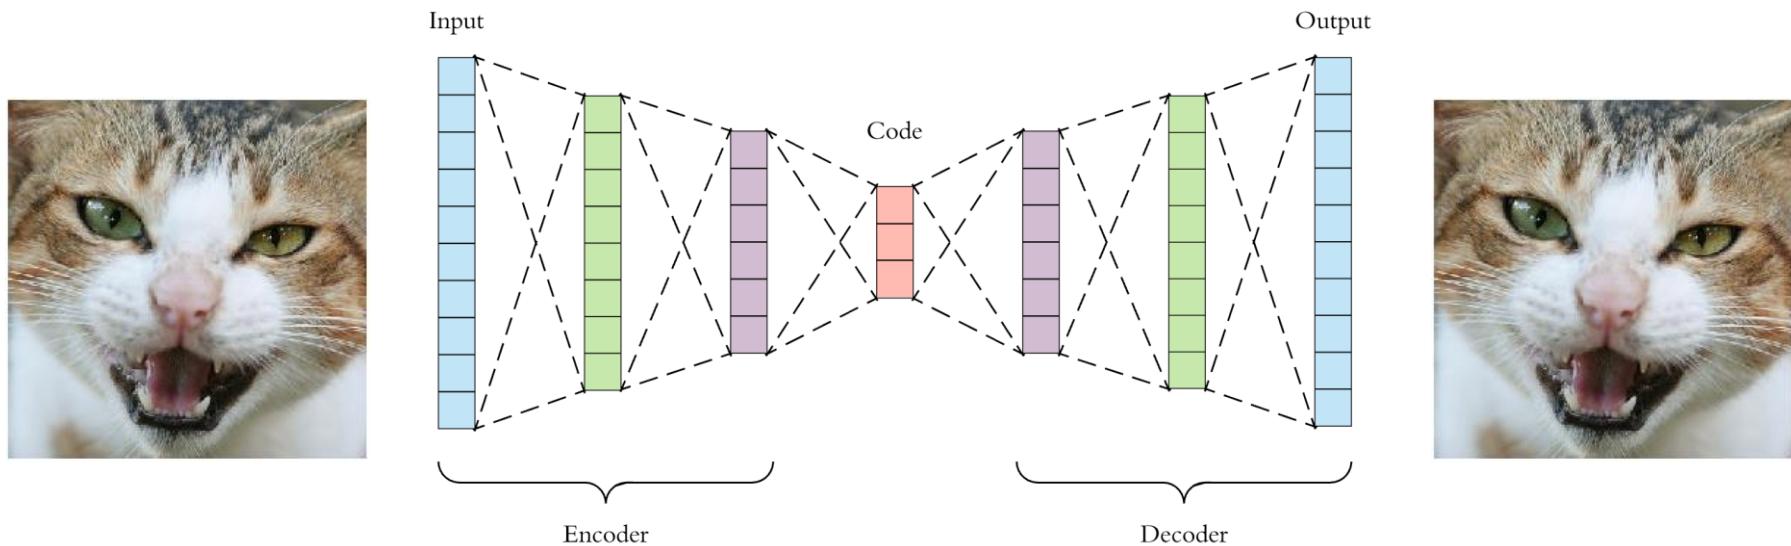
Original CT



Synthetic CT

# Segmentation through synthesis

Autoencoders:



Autoencoders can be used for:

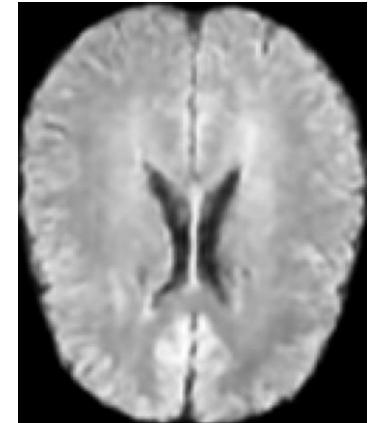
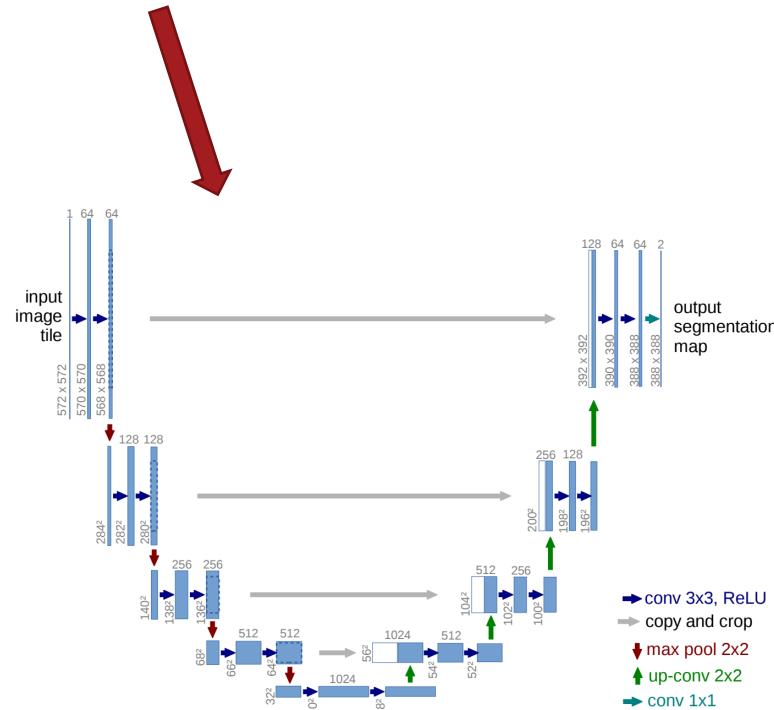
- Data compression
- Data denoising

# Segmentation through synthesis

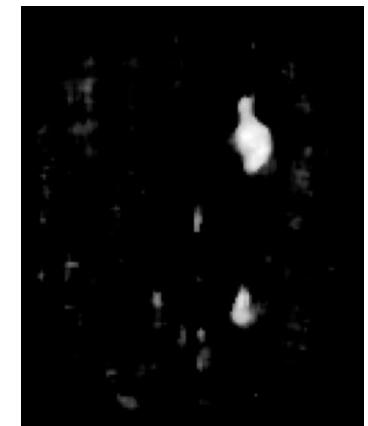
Let's train an autoencoder on healthy brain images



Pathological test image



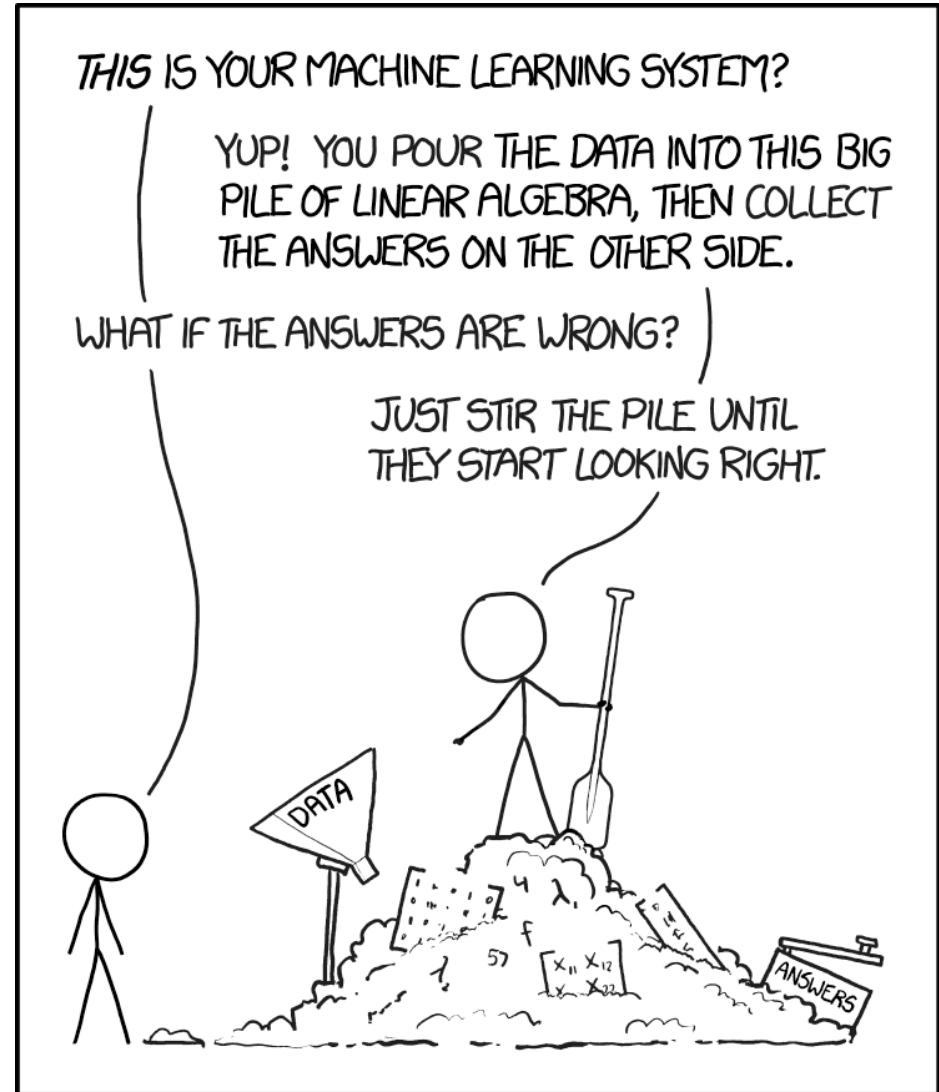
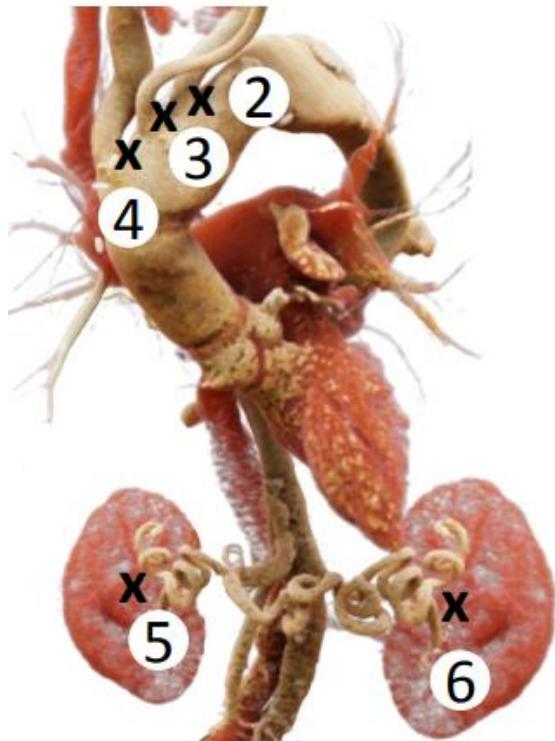
Synthetic image



Difference between  
real and synthetic image

# Landmarking

Can we effectively solve all segmentation/ classification /diagnosis problems in medical imaging following this schema?



# Cephalometry

Can we effectively solve all segmentation/ classification /diagnosis problems in medical imaging following this schema?

Can you guess what kind of measurements were used to compute the agreement?



100% of agreement



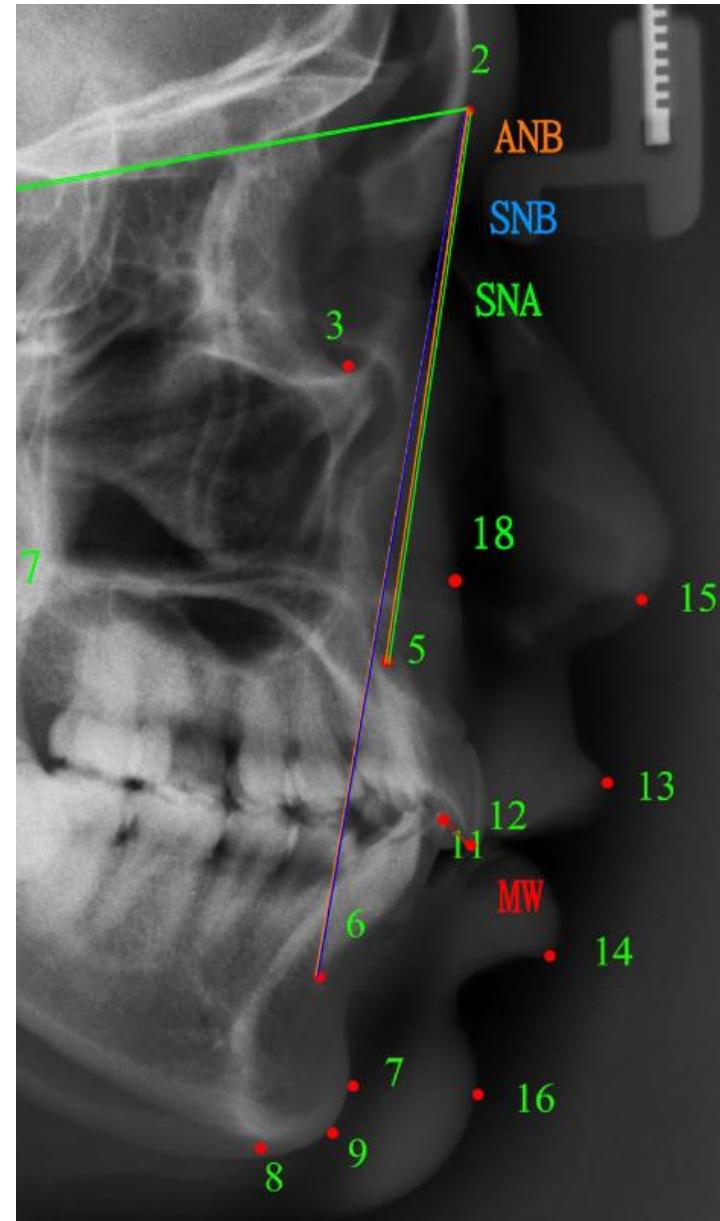
38% of agreement

# Cephalometry

Cephalometry is based on computing angles and distances from head X-ray

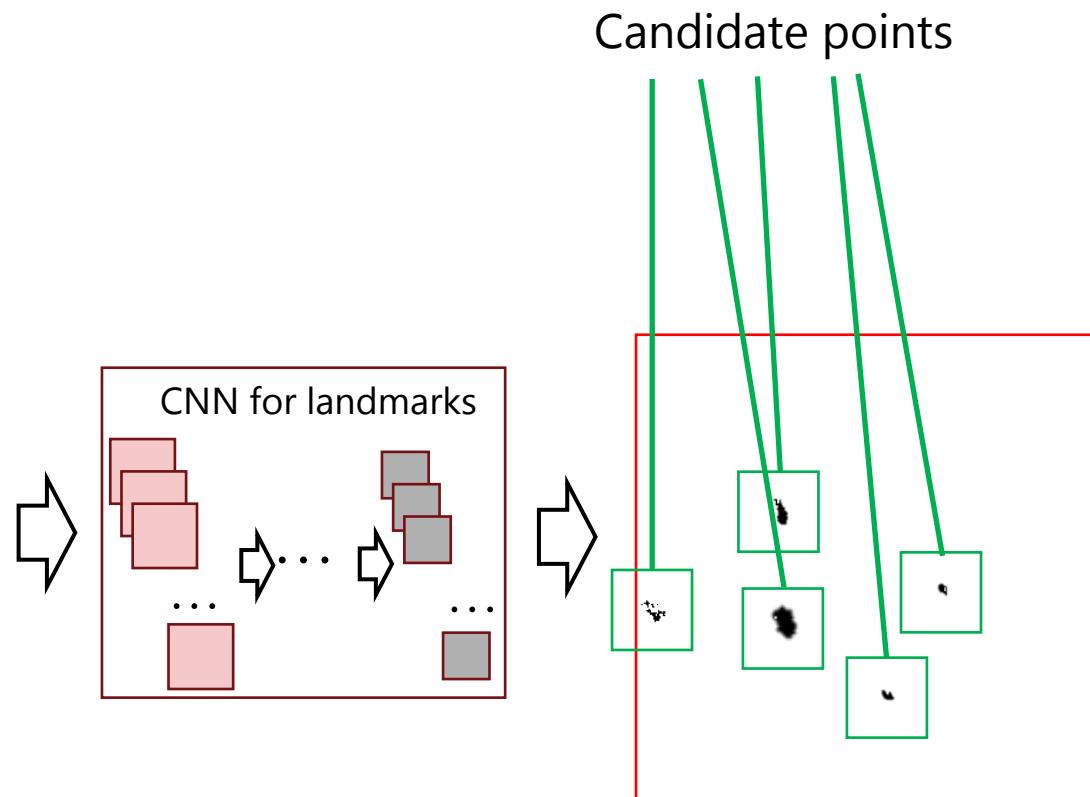
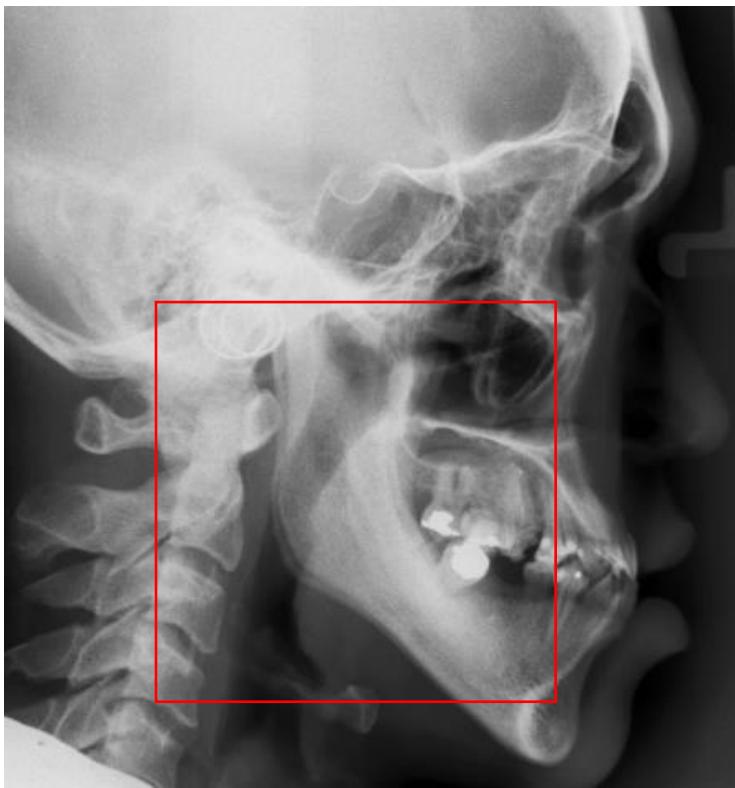
ANB angle:

- $[3.7^\circ - 5.2^\circ]$  – normal range
- $> 5.2^\circ$  - pathology type I
- $< 3.7^\circ$  - pathology type II



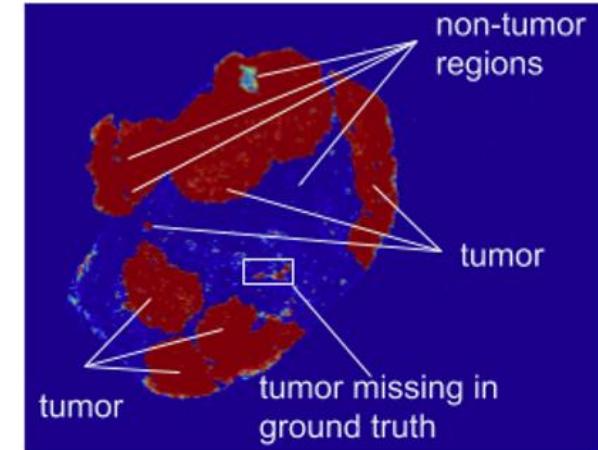
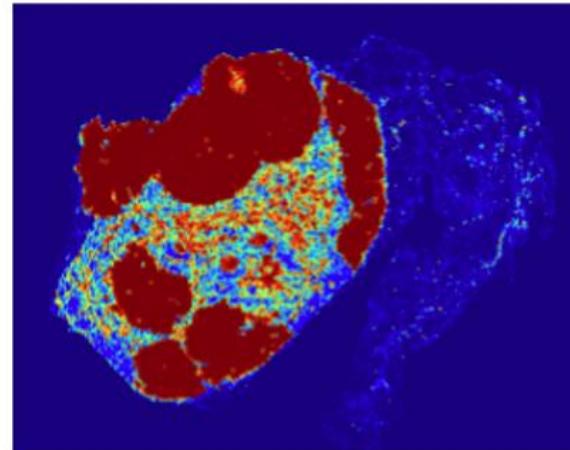
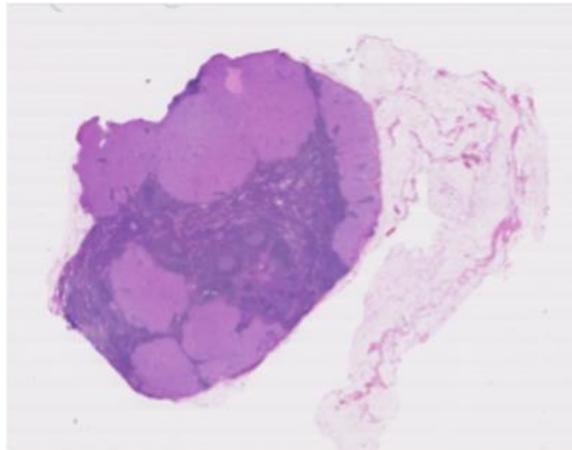
# Cephalometry

A good idea is to use CNN to detect landmarks and not to predict pathologies that are based on landmarks

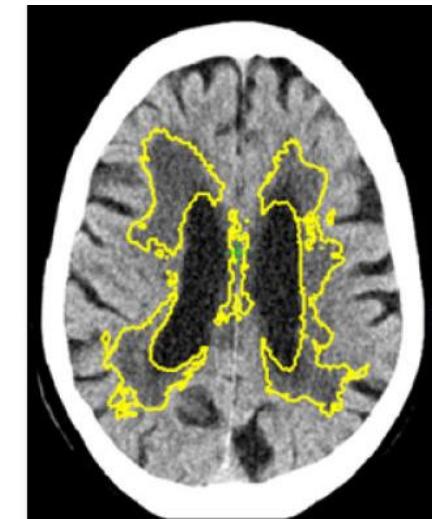
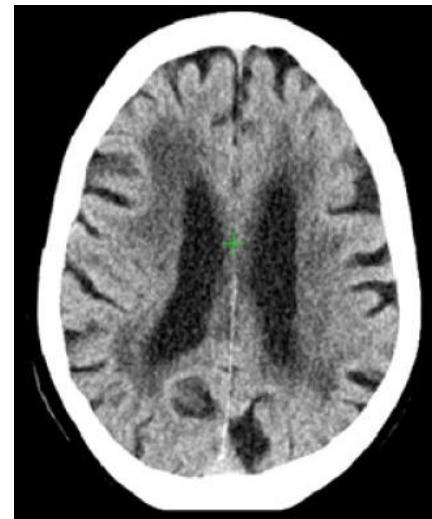


# Disease diagnosis

CNN for detection and segmentation pathologies



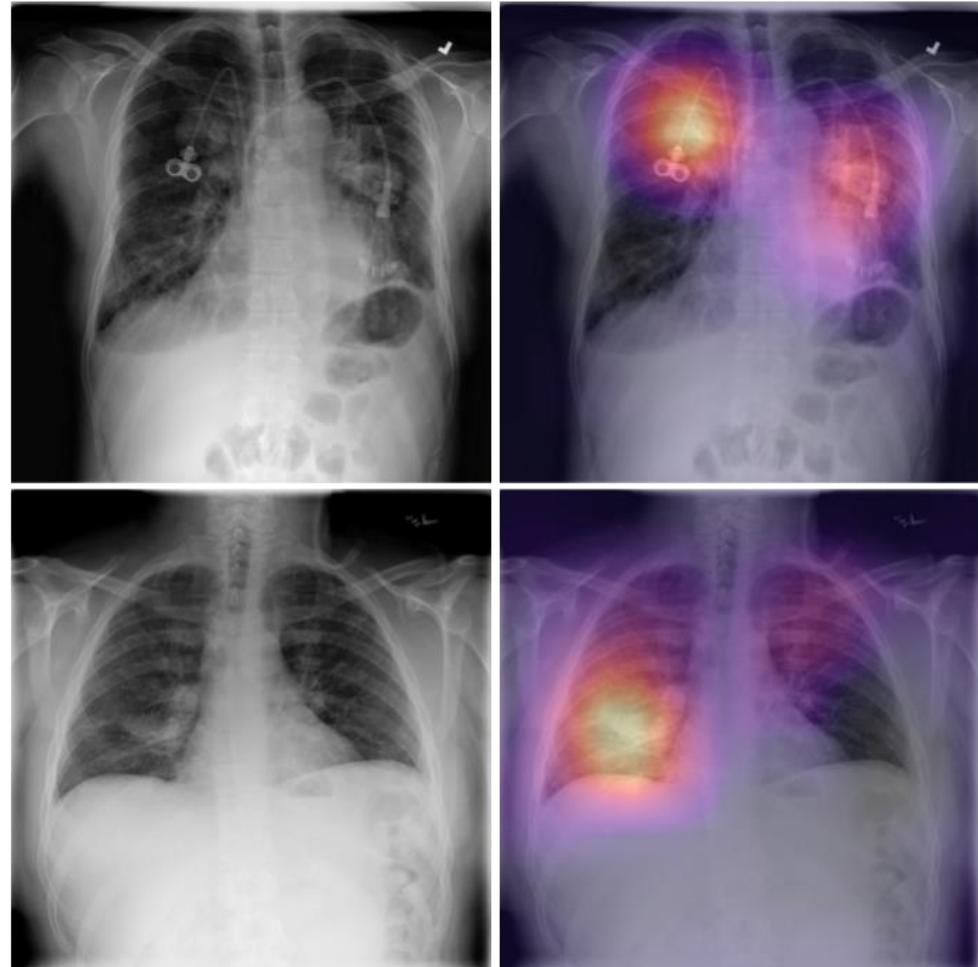
<https://ai.googleblog.com/2017/03/assisting-pathologists-in-detecting.html>



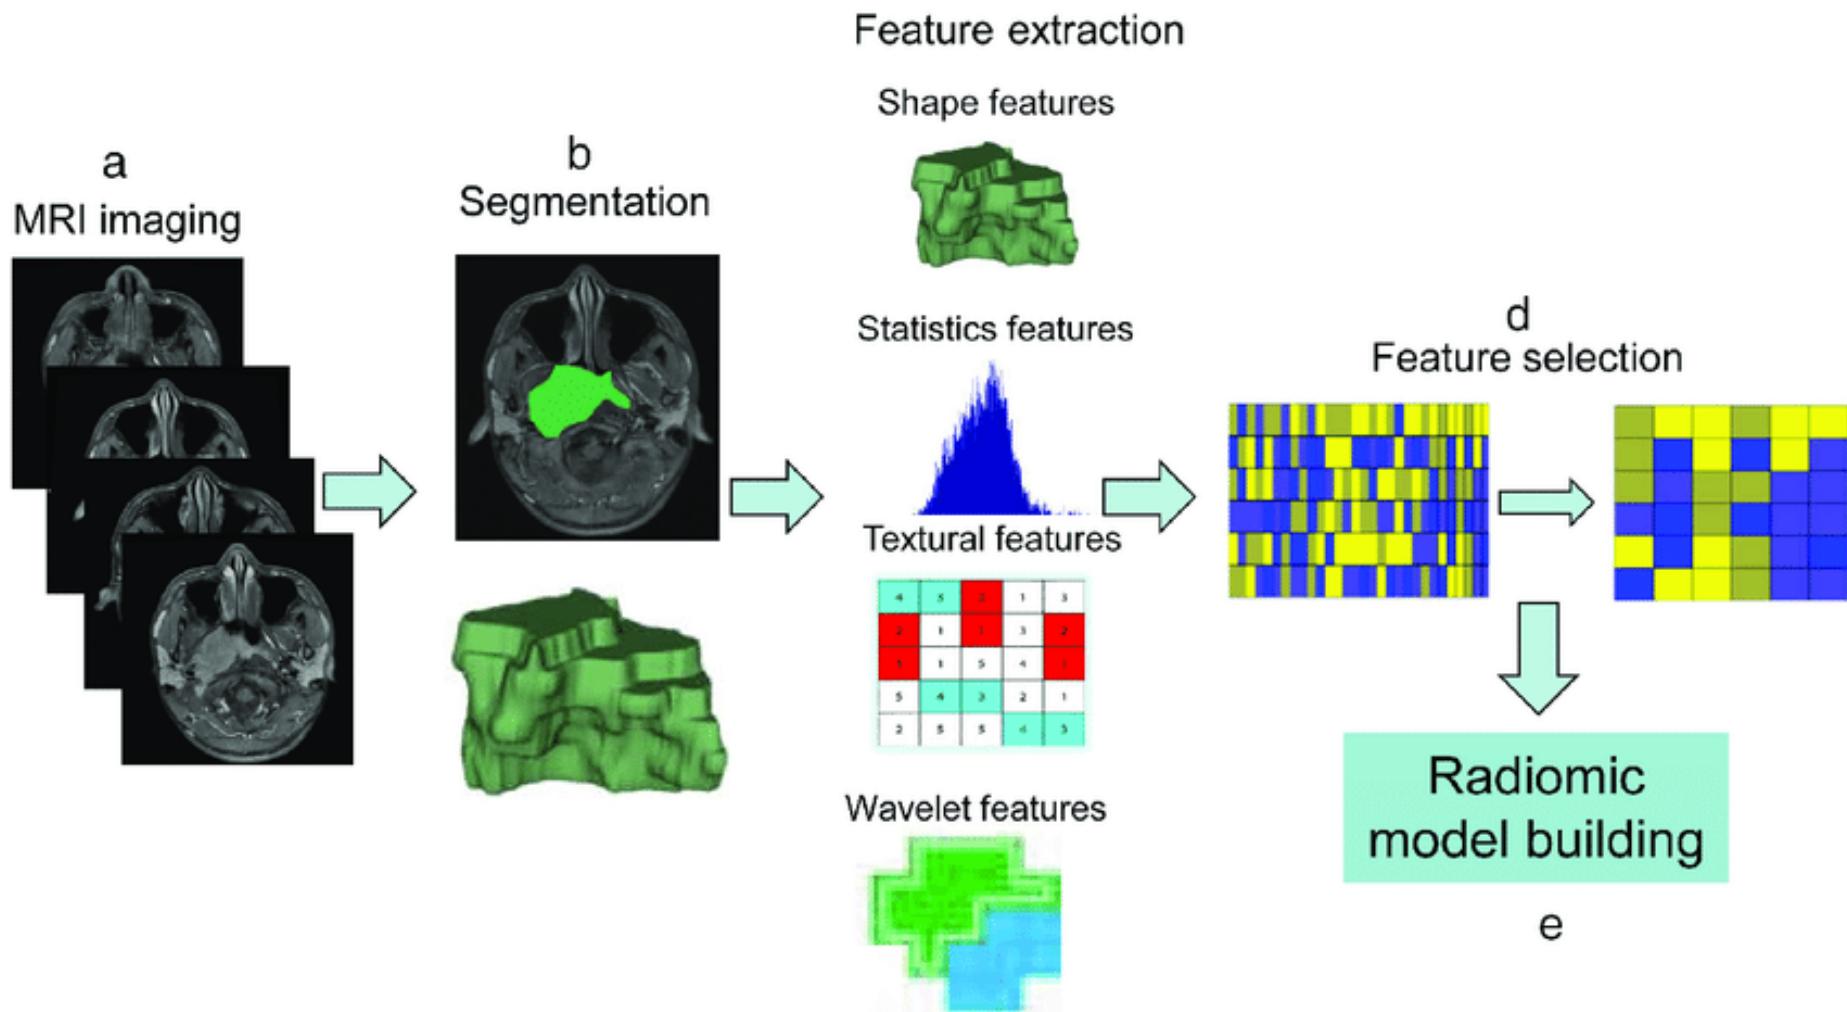
<https://www.imperial.ac.uk/news/186108/artificial-intelligence-improves-stroke-dementia-diagnosis/>

# Disease diagnosis

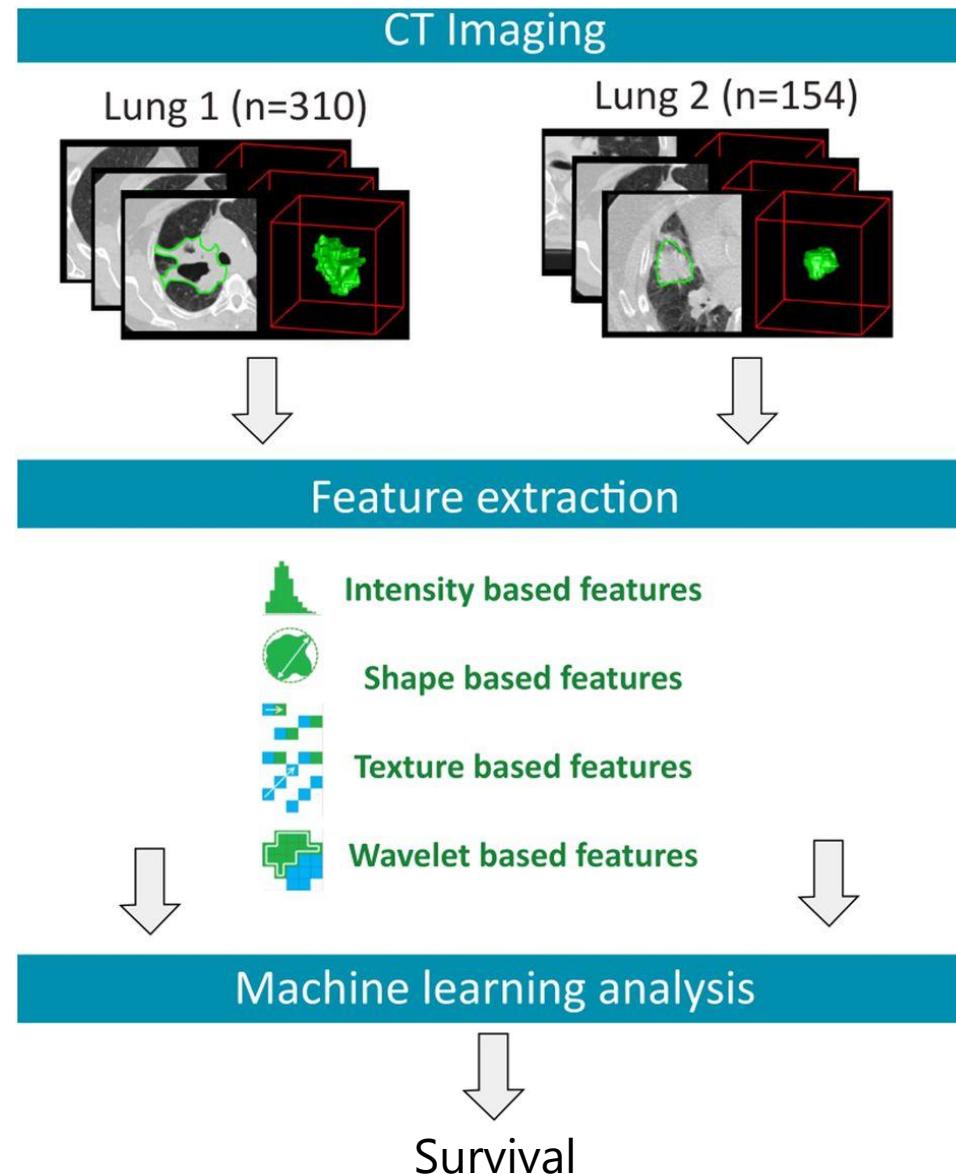
- Diagnosis with explicit segmentation
- CNNs do not know where to look, but know which pathology should be present
- Heatmaps from CNNs tells us if CNNs look at meaningful spots



# Radiomics

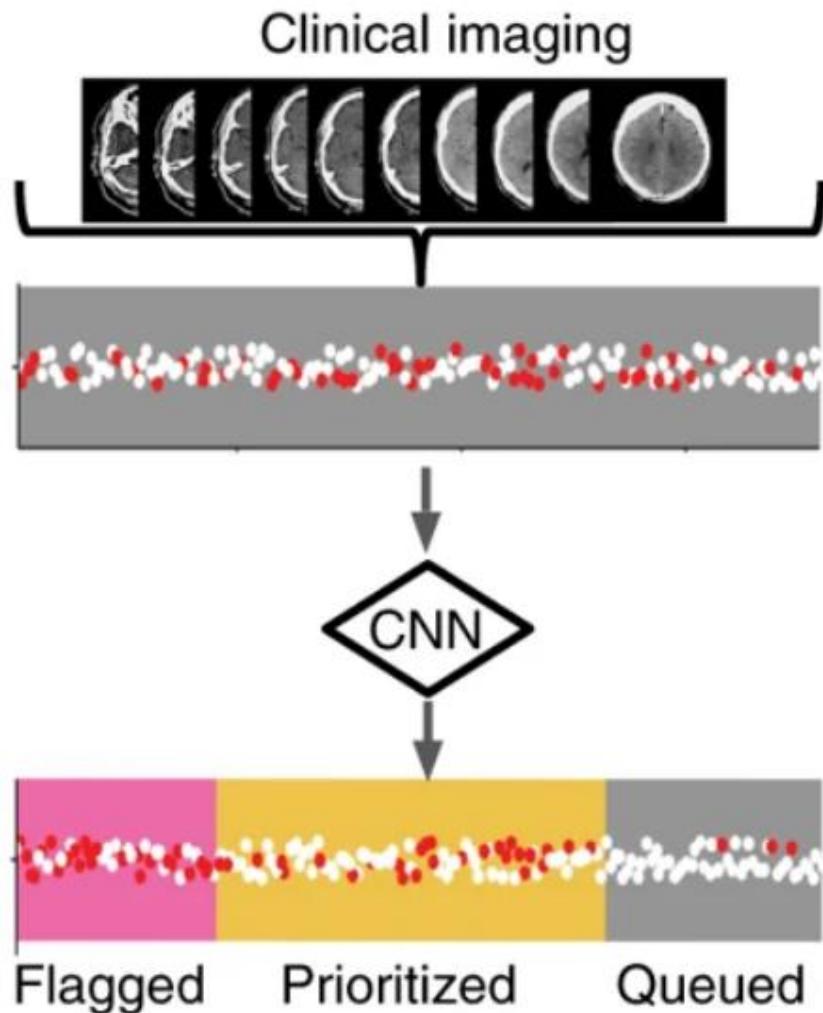


# Radiomics



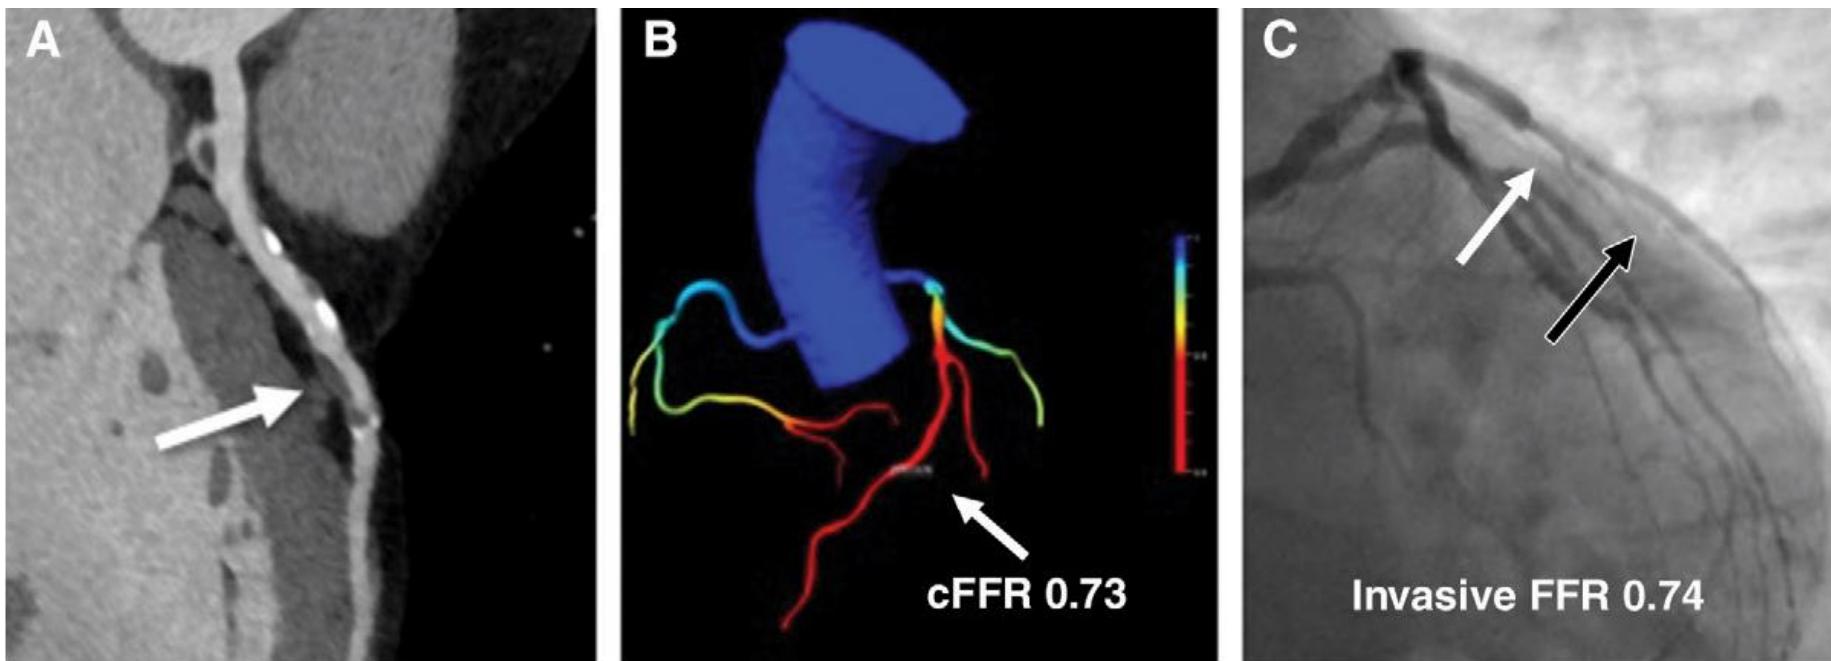
# Triaging

- Instead of automated diagnosis, machine learning can be used to sort patients
- Urgent patients will be shown first to physicians
- Safer, final decision is on human and not machine



# Fractional Flow Reserve

- Use machine learning to segment heart vessels
- Estimate the blood flow through vessels using physical models of machine learning

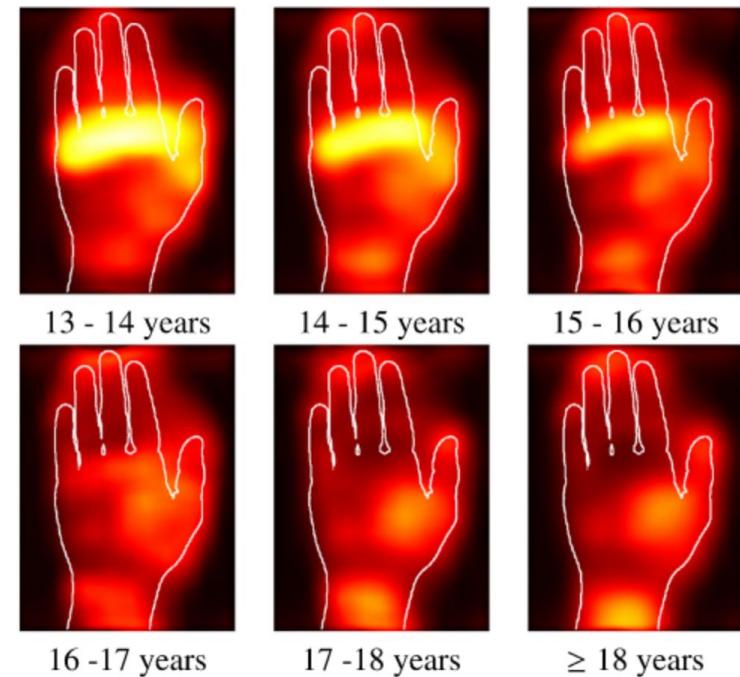
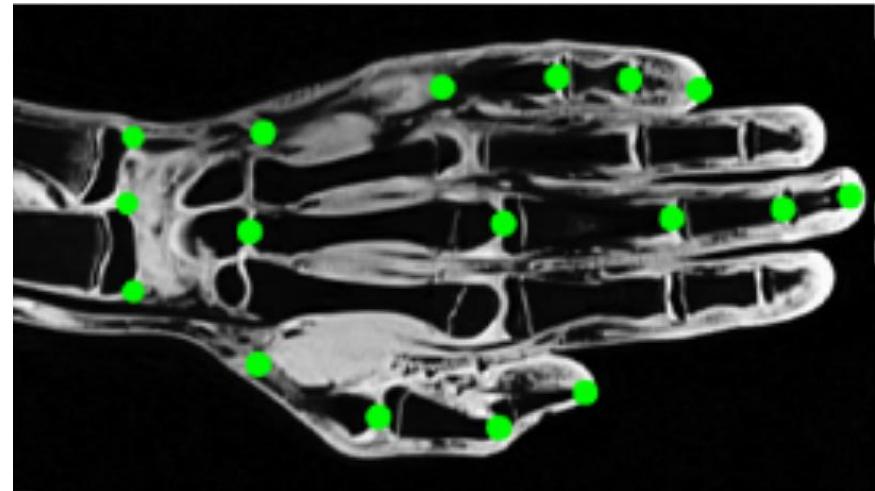


# Age estimation

Problem:

- A person claims he is younger than 15 years?
- No reliable document is available

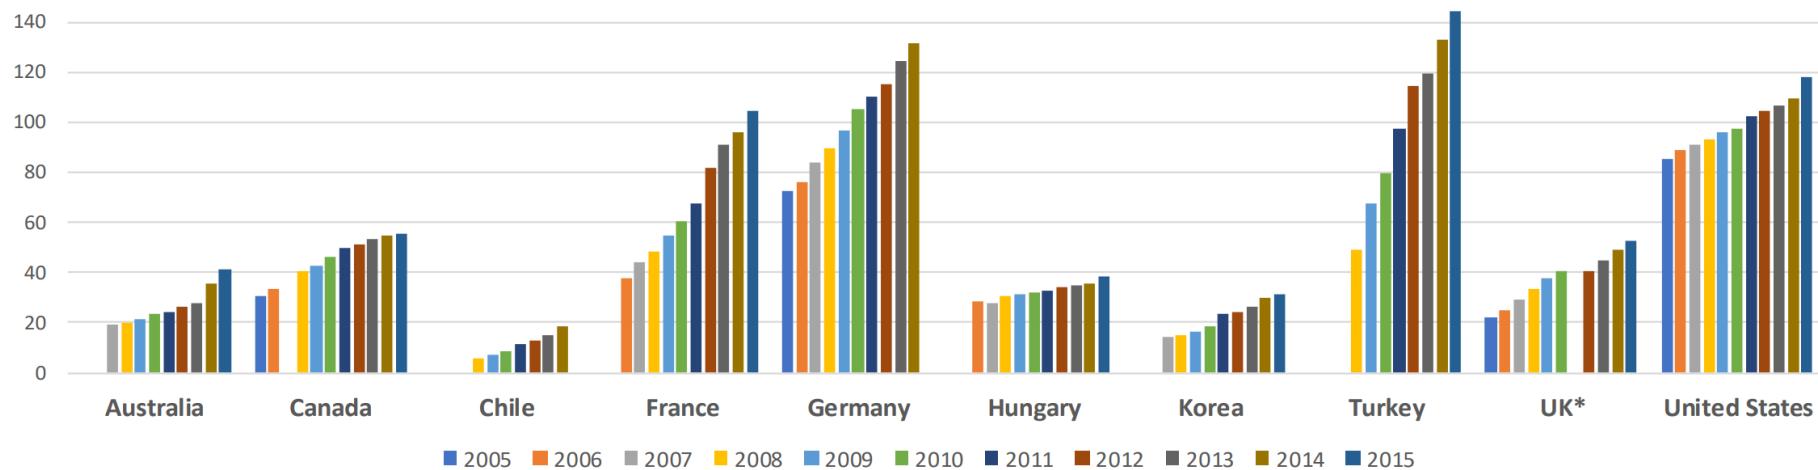
How to approximate the real age?



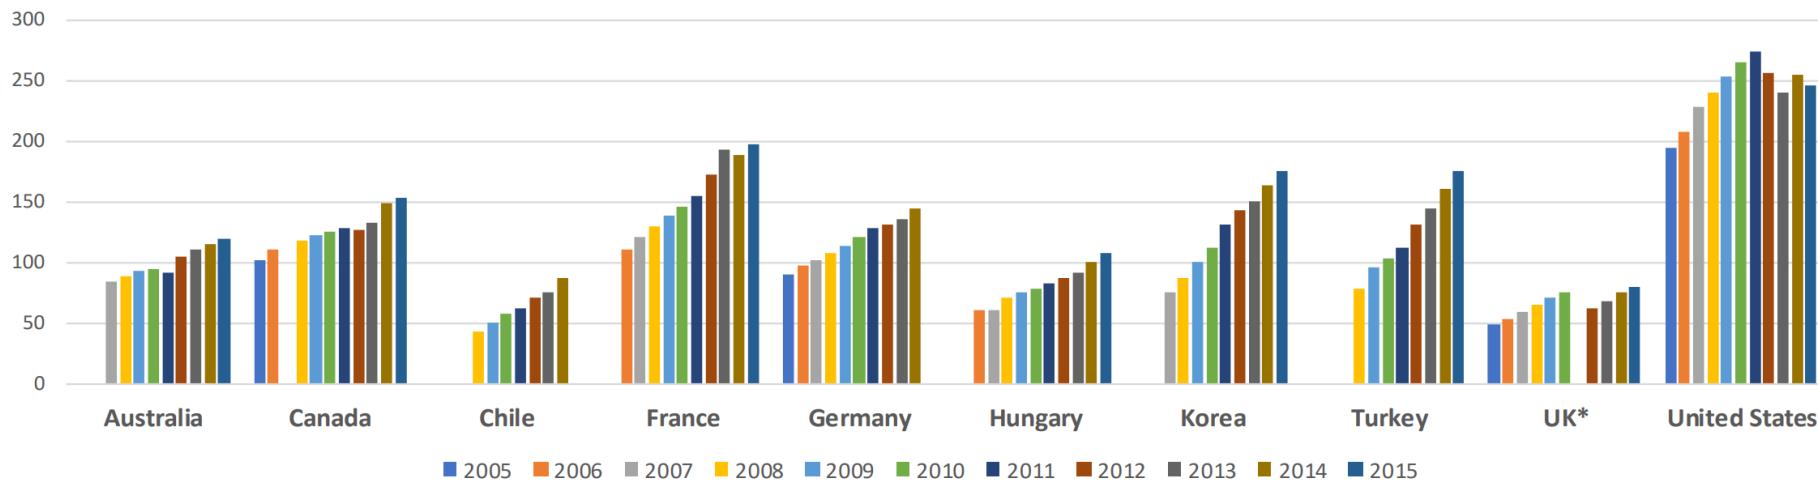
\*Automated age estimation from MRI volumes of the hand, 2019

# Potential for computer-aided medical image analyses

Diagnostic exams, Magnetic Resonance Imaging exams, Per 1 000 population



Diagnostic exams, Computed Tomography exams, Per 1 000 population



# Questions?