

Generalization Bounds

Yevgeny Seldin

Generalization Bounds

- Generalization bound for a single hypothesis
- Learning by selection
- Generalization lower bound
- Generalization bound for a finite \mathcal{H} (finite selection)
- Approximation-Estimation (bias-variance) trade-off
- Occam's razor: generalization bound for a countable \mathcal{H}
- Application example (Occam): binary decision trees

Generalization bound: single hypothesis

$$\begin{array}{c} \boxed{S = \{(X_1, Y_1), \dots, (X_n, Y_n)\}} \\ \downarrow \\ h \longrightarrow \hat{L}(h, S) = \frac{1}{n} \sum_{i=1}^n \ell(h(X_i), Y_i) \end{array}$$

Hoeffding:

$$\mathbb{P} \left(\mathbb{E} \left[\frac{1}{n} \sum_{i=1}^n Z_i \right] - \frac{1}{n} \sum_{i=1}^n Z_i \geq \sqrt{\frac{\ln \frac{1}{\delta}}{2n}} \right) \leq \delta$$

- Hoeffding:

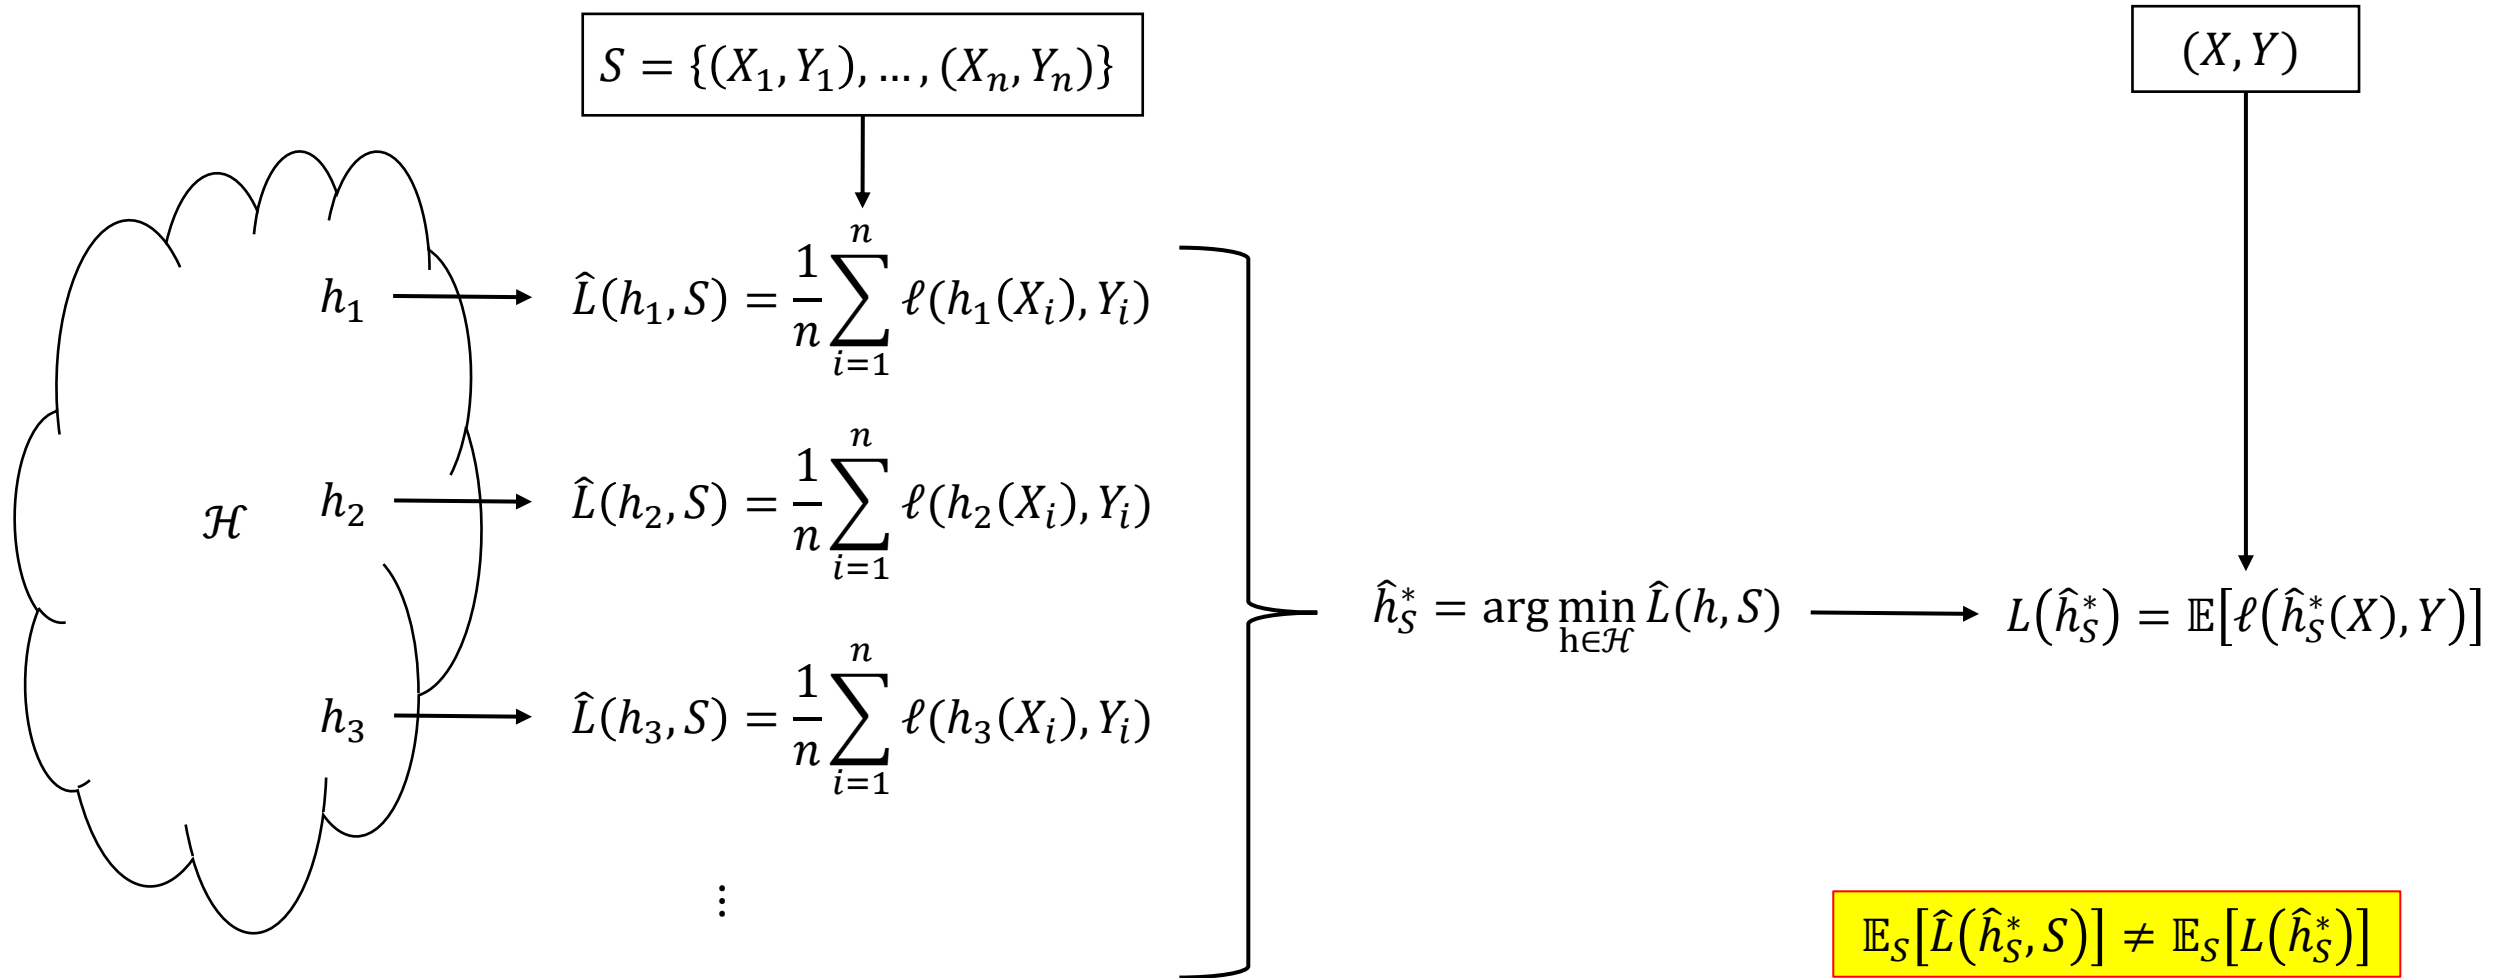
$$\begin{aligned} & \mathbb{P} \left(L(h) - \hat{L}(h, S) \geq \sqrt{\frac{\ln \frac{1}{\delta}}{2n}} \right) \leq \delta \\ \Rightarrow & \mathbb{P} \left(L(h) - \hat{L}(h, S) \leq \sqrt{\frac{\ln \frac{1}{\delta}}{2n}} \right) \geq 1 - \delta \end{aligned}$$

- In words: with probability at least $1 - \delta$:

$$L(h) \leq \hat{L}(h, S) + \sqrt{\frac{\ln \frac{1}{\delta}}{2n}}$$

(the probability is over observing $\hat{L}(h, S)$, not over $L(h)$)


Learning by Selection



Lower bound for learning by selection from finite \mathcal{H}

- Lower bound

- $|\mathcal{X}| = 2n$
- $|\mathcal{H}| = 2^{2n}$
- $p(x)$ – uniform
- y – random w.p. $\frac{1}{2}$



- $\mathbb{E}[\hat{L}(\hat{h}_S^*, S)] = 0$
- $\mathbb{E}[L(\hat{h}_S^*)] \geq \frac{1}{2} \cdot 0 + \frac{1}{2} \cdot \frac{1}{2} \geq \frac{1}{4}$

- Corollary:

$$\mathbb{P}\left(\exists h \in \mathcal{H}: L(h) - \hat{L}(h, S) \geq \frac{1}{8}\right) \geq \frac{1}{8}$$

- Proof by contradiction:

Assume that

$$\mathbb{P}\left(\exists h \in \mathcal{H}: L(h) - \hat{L}(h, S) \geq \frac{1}{8}\right) < \frac{1}{8}$$

Then

$$\mathbb{E}[L(\hat{h}_S^*)] \leq \frac{1}{8} \cdot 1 + \left(1 - \frac{1}{8}\right) \left(\underbrace{\hat{L}(\hat{h}_S^*, S)}_{=0} + \frac{1}{8}\right) < \frac{1}{4}$$

Generalization bound for finite \mathcal{H}

- Theorem: Let \mathcal{H} be finite with $|\mathcal{H}| = M$. Then

$$\mathbb{P}\left(\exists h \in \mathcal{H}: L(h) \geq \hat{L}(h, S) + \sqrt{\frac{\ln \frac{M}{\delta}}{2n}}\right) \leq \delta$$

- Corollary: $\mathbb{P}\left(L(\hat{h}_S^*) \geq \hat{L}(\hat{h}_S^*, S) + \sqrt{\frac{\ln \frac{M}{\delta}}{2n}}\right) \leq \delta$
- Equivalently: $\mathbb{P}\left(L(\hat{h}_S^*) \leq \hat{L}(\hat{h}_S^*, S) + \sqrt{\frac{\ln \frac{M}{\delta}}{2n}}\right) \geq 1 - \delta$
- In words: with probability at least $1 - \delta$:

$$L(\hat{h}_S^*) \leq \hat{L}(\hat{h}_S^*, S) + \sqrt{\frac{\ln \frac{M}{\delta}}{2n}}$$

- For single h we had:

$$L(h) \leq \hat{L}(h, S) + \sqrt{\frac{\ln \frac{1}{\delta}}{2n}}$$

- Price of selection: $\ln M$

- Lower bound: for $|\mathcal{H}| = M = 2^{2n}$
 $\mathbb{P}\left(\exists h \in \mathcal{H}: L(h) - \hat{L}(h, S) \geq \frac{1}{8}\right) \geq \frac{1}{8}$

- No contradiction: $\sqrt{\frac{\ln \frac{M}{\delta}}{2n}} \approx 1$

- For $M \ll e^n$ we get a meaningful bound

Proof

$\mathbb{E}_S[\hat{L}(\hat{h}_S^*, S)] \neq \mathbb{E}_S[L(\hat{h}_S^*)]$
We cannot apply Hoeffding!

We break the dependence

$$\mathbb{P}\left(L(\hat{h}_S^*) \geq \hat{L}(\hat{h}_S^*, S) + \sqrt{\frac{\ln \frac{\textcolor{red}{M}}{\delta}}{2n}}\right) \leq \mathbb{P}\left(\exists h \in \mathcal{H}: L(h) \geq \hat{L}(h, S) + \sqrt{\frac{\ln \frac{\textcolor{red}{M}}{\delta}}{2n}}\right)$$

(Union bound)

$$\leq \sum_{h \in \mathcal{H}} \mathbb{P}\left(L(h) \geq \hat{L}(h, S) + \sqrt{\frac{\ln \frac{\textcolor{red}{M}}{\delta}}{2n}}\right)$$

(Hoeffding with $\delta' = \frac{\delta}{M}$)


$$\leq \sum_{h \in \mathcal{H}} \frac{\delta}{\textcolor{red}{M}} = \delta$$

In the background:

The space $(\mathcal{X} \times \mathcal{Y})^n$ of all possible samples S of size n

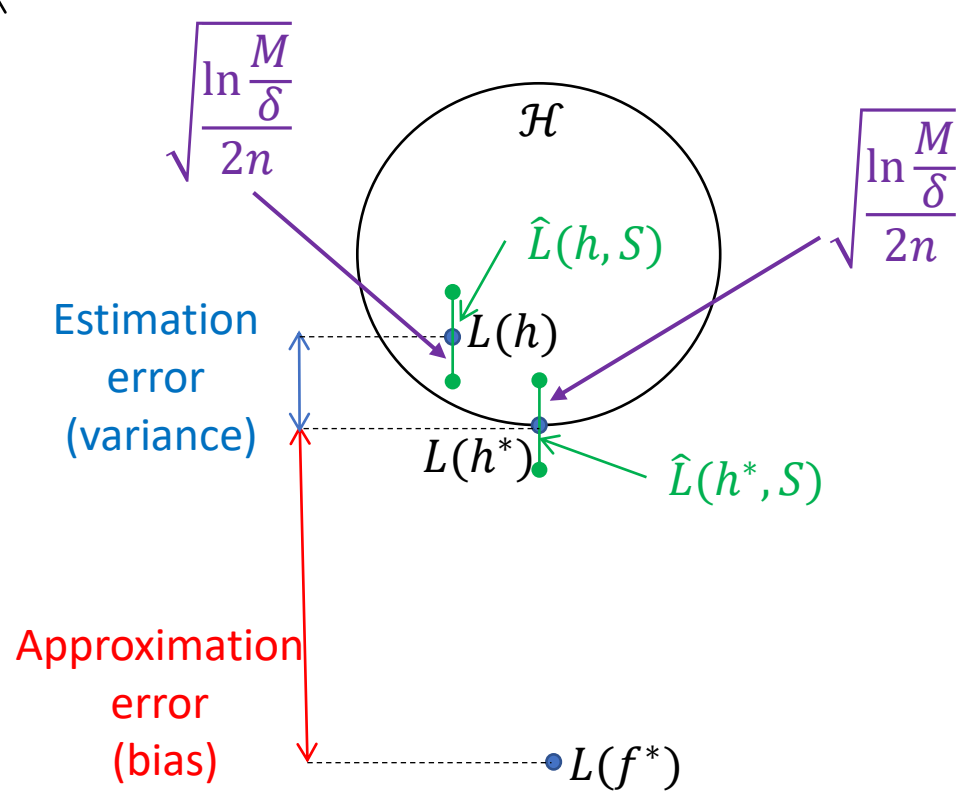
Each h gets $\frac{1}{M}$ share of the confidence budget δ

The total probability mass of violations of the inequality is bounded by δ




Approximation-Estimation (bias-variance) trade-off

Error



Selection from a small \mathcal{H}



Selection from a large \mathcal{H}

Estimation error $L(h) - L(h^*)$
can be up to $2 \sqrt{\frac{\ln \frac{M}{\delta}}{2n}}$

Occam's razor – Generalization bound for countable \mathcal{H}

- Theorem (Occam's razor): Let $\pi(h)$ be nonnegative and **independent of S** and satisfy $\sum_{h \in \mathcal{H}} \pi(h) \leq 1$. Then:

$$\mathbb{P} \left(\exists h \in \mathcal{H} : L(h) \geq \hat{L}(h, S) + \sqrt{\frac{\ln \frac{1}{\pi(h)\delta}}{2n}} \right) \leq \delta.$$

- Proof:

$$\mathbb{P} \left(\exists h \in \mathcal{H} : L(h) \geq \hat{L}(h, S) + \sqrt{\frac{\ln \frac{1}{\pi(h)\delta}}{2n}} \right)$$


(Union bound)

$$\leq \sum_{h \in \mathcal{H}} \mathbb{P} \left(L(h) \geq \hat{L}(h, S) + \sqrt{\frac{\ln \frac{1}{\pi(h)\delta}}{2n}} \right)$$

(Hoeffding, π is independent of S !)

$$\leq \sum_{h \in \mathcal{H}} \pi(h) \delta \leq \delta$$

In the background:
uneven distribution
of the confidence
budget δ according
to $\pi(h)$



Occam's razor selection

$$\mathbb{P} \left(\exists h \in \mathcal{H} : L(h) \geq \hat{L}(h, S) + \sqrt{\frac{\ln \frac{1}{\pi(h)\delta}}{2n}} \right) \leq \delta$$

$$h^* = \arg \min_h \underbrace{\hat{L}(h, S)}_{\substack{\text{Empirical} \\ \text{Performance}}} + \underbrace{\sqrt{\frac{\ln \frac{1}{\pi(h)\delta}}{2n}}}_{\text{Complexity}}$$

With probability at least $1 - \delta$: $L(\hat{h}_S^*) \leq \hat{L}(\hat{h}_S^*, S) + \sqrt{\frac{\ln \frac{1}{\pi(h)\delta}}{2n}}$

Application example: binary decision trees

$$\pi(\mathcal{H}_0) = \frac{1}{2}$$




$$|\mathcal{H}_0| = 2 = 2^{2^0}$$

$$\pi(\mathcal{H}_1) = \frac{1}{2^2} = \frac{1}{4}$$

$$|\mathcal{H}_1| = 4 = 2^{2^1}$$

$$\pi(\mathcal{H}_2) = \frac{1}{2^3} = \frac{1}{8}$$

$$|\mathcal{H}_2| = 2^{2^2}$$



...

- $d(h)$ - depth of tree h
- $\pi(h) = \pi(\mathcal{H}_{d(h)}) \frac{1}{|\mathcal{H}_{d(h)}|} = \frac{1}{2^{d(h)+1}} \frac{1}{2^{2^{d(h)}}}$
- $\sum_{h \in \mathcal{H}} \pi(h) = \sum_{d=0}^{\infty} \sum_{h \in \mathcal{H}_d} \pi(h) = \sum_{d=0}^{\infty} \sum_{h \in \mathcal{H}} \frac{1}{2^{d(h)+1}} \frac{1}{2^{2^{d(h)}}} = \sum_{d=0}^{\infty} \frac{1}{2^{d(h)+1}} \underbrace{\sum_{h \in \mathcal{H}_d} \frac{1}{2^{2^{d(h)}}}}_{=1} = \sum_{d=0}^{\infty} \frac{1}{2^{d(h)+1}} = 1$

- With probability $\geq 1 - \delta$, for all $h \in \mathcal{H}$: $L(h) \leq \hat{L}(h, S) + \sqrt{\frac{\ln \frac{1}{\pi(h)\delta}}{2n}} = \hat{L}(h, S) + \sqrt{\frac{\ln(2)(2^{d(h)}) + d(h)+1 + \ln \frac{1}{\delta}}{2n}}$

Alternative $\pi(\mathcal{H}_d) = \frac{1}{(d+1)(d+2)}$ dominant

$$L(h) \leq \hat{L}(h, S) + \sqrt{\frac{\ln(2)2^{d(h)} + \ln \frac{(d+1)(d+2)}{\delta}}{2n}}$$

- Permutation-symmetric trees get the same prior
- In absence of prior knowledge, no reason to discriminate (structurally symmetric prior)
- The size of \mathcal{H}_d gives the dominant term
- Why no contradiction with the lower bound?