

Clustering

Mohammad Sadegh Talebi

Department of Computer Science



Introduction

Clustering: Grouping **unlabeled** data points by **similarity**

- putting similar items in the same group, and
- dissimilar items end up in different groups

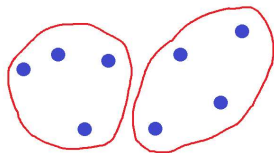


Introduction

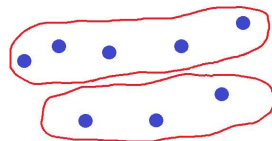
Clustering: Grouping **unlabeled** data points by **similarity**

- putting similar items in the same group, and
- dissimilar items end up in different groups

Although sounds natural, clustering is an ill-defined problem.



VS.

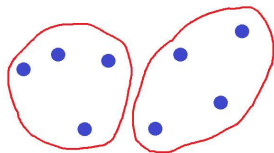


Introduction

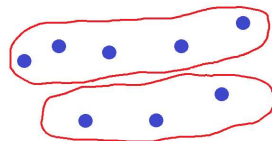
Clustering: Grouping **unlabeled** data points by **similarity**

- putting similar items in the same group, and
- dissimilar items end up in different groups

Although sounds natural, clustering is an ill-defined problem.



vs.



Clustering methods are **unsupervised learning** techniques —No access to a labeling teacher \implies Lack of ground truth.



Motivation

Clustering is perhaps the most widely used tool for exploratory data analysis.

Why clustering?

- Gaining information about internal structure of data
- Modeling over smaller subsets of data
- Data reduction
- Outlier detection



Clustering: Issues

Clustering is **subjective**:

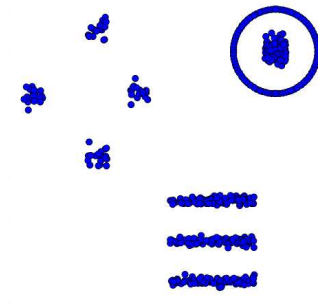
- Proper clustering depends upon the context
- Even the number of clusters is not well-defined in all tasks



Clustering: Issues

Clustering is **subjective**:

- Proper clustering depends upon the context
- Even the number of clusters is not well-defined in all tasks



How many clusters do *you* see here?



Clustering: Issues

- Clustering relies on similarity (or proximity) notions, which are **not transitive**.
 - But clustering is an **equivalence relation** (and thus, transitive).
- A relation R is **transitive** if when R relates u to v , and v to w , then it also relates u to w .
 - A relation R is an **equivalence** if it is reflexive, symmetric, and transitive.



Ingredients

For clustering, we need:



Ingredients

For clustering, we need:

(i) Proximity measure

- A **similarity score function** to quantify how similar x is to y
- Or a **distance function** to quantify how dissimilar data items are



Ingredients

For clustering, we need:

(i) Proximity measure

- A **similarity score function** to quantify how similar x is to y
- Or a **distance function** to quantify how dissimilar data items are

(ii) Criterion to assess the **quality** of clustering

- Helps us compare possible clusterings



Ingredients

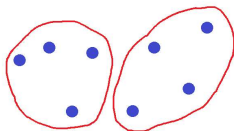
For clustering, we need:

(i) Proximity measure

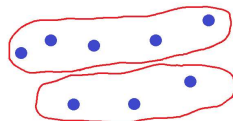
- A **similarity score function** to quantify how similar x is to y
- Or a **distance function** to quantify how dissimilar data items are

(ii) Criterion to assess the **quality** of clustering

- Helps us compare possible clusterings



vs.



Ingredients

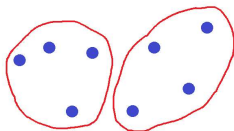
For clustering, we need:

(i) Proximity measure

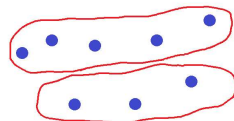
- A **similarity score function** to quantify how similar x is to y
- Or a **distance function** to quantify how dissimilar data items are

(ii) Criterion to assess the **quality** of clustering

- Helps us compare possible clusterings



vs.



(iii) Algorithm to cluster items

- To determine a good clustering in view of the chosen criterion.



A Clustering Model

A rigorous common clustering setup:

Input:

- A set of elements \mathcal{X}
- A **distance function** $d : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}_+$ (or a similarity function $s : \mathcal{X} \times \mathcal{X} \rightarrow [0, 1]$)
- The number k of clusters (optional)

Output:

- A **partition** of \mathcal{X} into subsets $C = (C_1, \dots, C_k)$, namely,

$$\bigcup_{i=1}^k C_i = \mathcal{X} \quad \text{and} \quad C_i \cap C_j = \emptyset, \forall i \neq j.$$

- A measure of the quality of C (optional)



A Clustering Model

A rigorous common clustering setup:

Input:

- A set of elements \mathcal{X}
- A **distance function** $d : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}_+$ (or a similarity function $s : \mathcal{X} \times \mathcal{X} \rightarrow [0, 1]$)
- The number k of clusters (optional)

Output:

- A **partition** of \mathcal{X} into subsets $C = (C_1, \dots, C_k)$, namely,

$$\bigcup_{i=1}^k C_i = \mathcal{X} \quad \text{and} \quad C_i \cap C_j = \emptyset, \forall i \neq j.$$

- A measure of the quality of C (optional)

Often elements in \mathcal{X} need to be encoded as vectors in \mathbb{R}^m .



Clustering Paradigms

Popular paradigms for clustering:

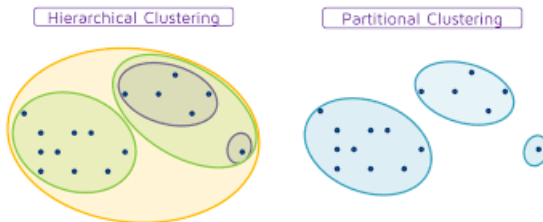
- Hierarchical clustering
 - Agglomerative
 - Divisive



Clustering Paradigms

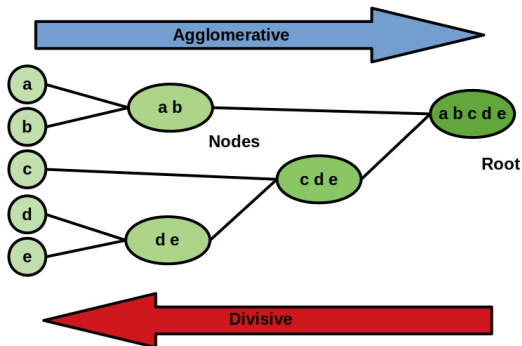
Popular paradigms for clustering:

- Hierarchical clustering
 - Agglomerative
 - Divisive
- Partitional (or flat) clustering
 - Cost minimization based
 - Spectral clustering



Hierarchical Clustering

Agglomerative and divisive clustering



Partitional Clustering via Cost Minimization

This lecture: Partitional clustering based on cost minimization

We study:

- The **K-means** problem, a classical clustering/partitioning problem
- The **k-means** (or Lloyd's) algorithm to approximately solve K-means
- Convergence guarantees of k-means, and its pitfalls
- **k-means++** as an improved seeding for k-means
- Extensions beyond K-means (e.g., K-median)



The K-Means Problem

K-means is a classical problem in computational geometry but also arising in many applications.

The K-Means Problem

Given $\mathcal{X} \subset \mathbb{R}^m$, find $\mathcal{C} = \{c_1, \dots, c_k\} \subset \mathbb{R}^m$ so as to minimize

$$\phi(\mathcal{X}, \mathcal{C}) := \sum_{x \in \mathcal{X}} \min_{c \in \mathcal{C}} \|x - c\|_2^2$$



The K-Means Problem

K-means is a classical problem in computational geometry but also arising in many applications.

The K-Means Problem

Given $\mathcal{X} \subset \mathbb{R}^m$, find $\mathcal{C} = \{c_1, \dots, c_k\} \subset \mathbb{R}^m$ so as to minimize

$$\phi(\mathcal{X}, \mathcal{C}) := \sum_{x \in \mathcal{X}} \min_{c \in \mathcal{C}} \|x - c\|_2^2$$

The K-Means Problem:

$$\min_{\mathcal{C} \subset \mathbb{R}^m, |\mathcal{C}|=k} \phi(\mathcal{X}, \mathcal{C}) = \min_{\mathcal{C} \subset \mathbb{R}^m, |\mathcal{C}|=k} \sum_{x \in \mathcal{X}} \min_{c \in \mathcal{C}} \|x - c\|_2^2$$



The K-Means Problem

K-means is a classical problem in computational geometry but also arising in many applications.

The K-Means Problem

Given $\mathcal{X} \subset \mathbb{R}^m$, find $\mathcal{C} = \{c_1, \dots, c_k\} \subset \mathbb{R}^m$ so as to minimize

$$\phi(\mathcal{X}, \mathcal{C}) := \sum_{x \in \mathcal{X}} \min_{c \in \mathcal{C}} \|x - c\|_2^2$$

The K-Means Problem:
$$\min_{\mathcal{C} \subset \mathbb{R}^m, |\mathcal{C}|=k} \phi(\mathcal{X}, \mathcal{C}) = \min_{\mathcal{C} \subset \mathbb{R}^m, |\mathcal{C}|=k} \sum_{x \in \mathcal{X}} \min_{c \in \mathcal{C}} \|x - c\|_2^2$$

- ϕ is called the cost function. It can be defined using any distance function (e.g., $\|\cdot\|_1$).
- c_1, \dots, c_k **may not** belong to \mathcal{X} .
- An 0-1 Integer Program, which is NP-hard for general k .
- Even NP-hard to approximate



The K-Means Problem: Special Cases

Case 1: $k = 1$

- The solution for $k = 1$ is $\mathcal{C} = \{c_1\}$ where its i -th element is

$$c_{1,i} = \frac{1}{|\mathcal{X}|} \sum_{x \in \mathcal{X}} x_i$$



The K-Means Problem: Special Cases

Case 1: $k = 1$

- The solution for $k = 1$ is $\mathcal{C} = \{c_1\}$ where its i -th element is

$$c_{1,i} = \frac{1}{|\mathcal{X}|} \sum_{x \in \mathcal{X}} x_i$$

Case 2: $\mathcal{C} \subset \mathcal{X}$

- Enumerate all $\binom{n}{k}$ possible configurations and choose the one with minimal ϕ .
- So in this case, for a fixed k , the optimal clustering is found in polynomial time.



The k-means Algorithm

In 1956, Stuart Lloyd proposed [the k-means algorithm](#) (a.k.a. Lloyd's algorithm) to approximate the K-means problem.

k-means:

- Is based on **local search**, i.e., makes local improvements to an arbitrarily chosen initial clustering.
- Is very simple and easy to implement.



k-means

The k-means algorithm partitions the given data into k clusters:

- Each cluster has a cluster center, called **centroid**
- k is specified by the user



k-means

The k-means algorithm partitions the given data into k clusters:

- Each cluster has a cluster center, called **centroid**
- k is specified by the user

Centroid

Consider a finite set $S \subset \mathbb{R}^m$. The **centroid** (or center of mass) of S is

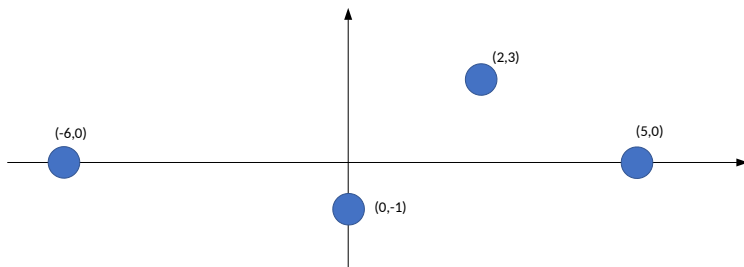
$$c = \frac{1}{|S|} \sum_{x \in S} x$$

The sum of squared distances of the points in S to point x is minimized when x is the centroid.



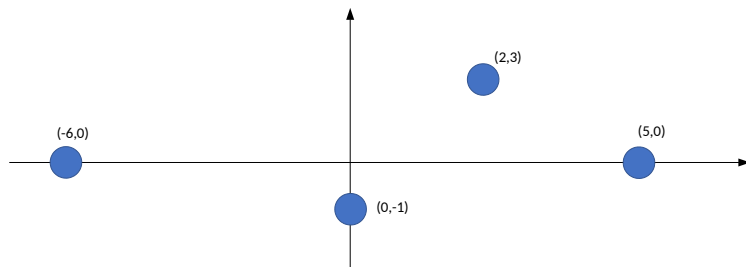
Example

What is the centriod of the following dataset?



Example

What is the centroid of the following dataset?



$$c_1 = \frac{1}{4}(-6 + 0 + 2 + 5) = \frac{1}{4}, \quad c_2 = \frac{1}{4}(0 - 1 + 3 + 0) = \frac{1}{2}$$



So the centroid is $c = (c_1, c_2) = (\frac{1}{4}, \frac{1}{2})$

k-means

input: \mathcal{X}, k

initialization: Select c_1, \dots, c_k arbitrarily.

k-means

input: \mathcal{X}, k

initialization: Select c_1, \dots, c_k arbitrarily.

repeat until convergence:

- **Assign to clusters:** For all $i \in [k]$, set

$$C_i = \left\{ x \in \mathcal{X} : i = \operatorname{argmin}_j \|x - c_j\| \right\}$$

namely, map each point $x \in \mathcal{X}$ to its nearest cluster center (w.r.t. the Euclidean distance).

k-means

input: \mathcal{X}, k

initialization: Select c_1, \dots, c_k arbitrarily.

repeat until convergence:

- **Assign to clusters:** For all $i \in [k]$, set

$$C_i = \left\{ x \in \mathcal{X} : i = \operatorname{argmin}_j \|x - c_j\| \right\}$$

namely, map each point $x \in \mathcal{X}$ to its nearest cluster center (w.r.t. the Euclidean distance).

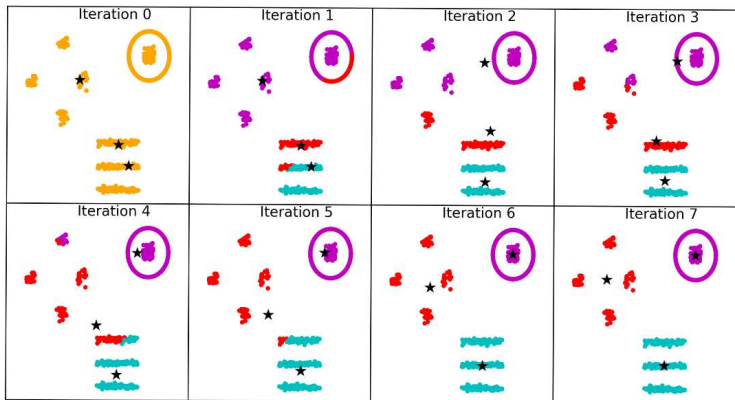
- **Update cluster centers:** For all $j \in [k]$, compute new centroids:

$$c_i = \frac{1}{|C_i|} \sum_{x \in C_i} x$$



k-means: Illustration

Illustration of k-means for $k = 3$:



k-means: Convergence

Lemma

*Each iteration of k-means does not increase ϕ . Thus, k-means always converges after **finitely** many iterations.*



k-means: Convergence

Lemma

*Each iteration of k-means does not increase ϕ . Thus, k-means always converges after **finitely** many iterations.*

The lemma relies on the following facts:

- Let S be a set of points with center of mass $c(S)$, and let z be an arbitrary point. Then,

$$\sum_{x \in S} \|x - z\|^2 - \sum_{x \in S} \|x - c(S)\|^2 = |S| \cdot \|c(S) - z\|^2.$$



k-means: Convergence

Lemma

*Each iteration of k-means does not increase ϕ . Thus, k-means always converges after **finitely** many iterations.*

The lemma relies on the following facts:

- Let S be a set of points with center of mass $c(S)$, and let z be an arbitrary point. Then,

$$\sum_{x \in S} \|x - z\|^2 - \sum_{x \in S} \|x - c(S)\|^2 = |S| \cdot \|c(S) - z\|^2.$$

- k-means makes local improvements to an arbitrary clustering until it is no longer possible to do so.
- There are finitely many feasible clustering configurations.



k-means: Iteration Complexity

- A trivial iteration complexity is $O(k^n)$ with $n =: |\mathcal{X}|$.



k-means: Iteration Complexity

- A trivial iteration complexity is $O(k^n)$ with $n =: |\mathcal{X}|$.
- Lower bounds are important for us to see if there are clustering instances challenging k-means.
- A known lower bound is $2^{\Omega(\sqrt{n})}$ due to Arthur and Vassilvitskii (2006).
 - This means there are a set of n data points and a set of adversarially chosen cluster centers for which the algorithm requires $2^{\Omega(\sqrt{n})}$ iterations.



k-means: Convergence

k-means may converge to a **locally-optimal** solution, which could be arbitrarily worse than the optimal clustering:

Theorem

For any $\alpha > 1$, there is an instance of the K-means problem for which k-means might find a solution whose cost is at least $\alpha \cdot \phi^$, where ϕ^* is the minimum K-means cost.*



k-means: Convergence

k-means may converge to a **locally-optimal** solution, which could be arbitrarily worse than the optimal clustering:

Theorem

For any $\alpha > 1$, there is an instance of the K-means problem for which k-means might find a solution whose cost is at least $\alpha \cdot \phi^$, where ϕ^* is the minimum K-means cost.*

The theorem relies on constructing a clustering instance due to Kanungo et al. (2004), on which k-means performs arbitrarily worse than the optimal.



k-means: Construction of a bad example

Example by Kanungo et al. (2004):

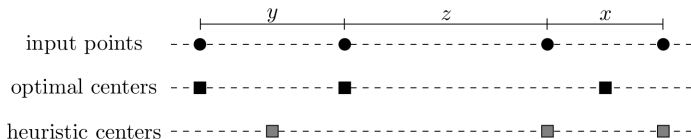
- Consider 4 points on the line such that $z > y > x$.



k-means: Construction of a bad example

Example by Kanungo et al. (2004):

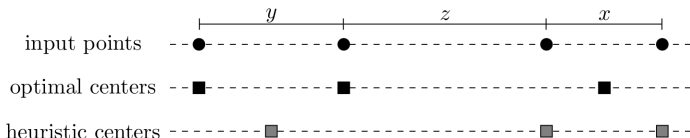
- Consider 4 points on the line such that $z > y > x$.
- For $k = 3$, the optimal cost is $\phi^* = x^2/2$.



k-means: Construction of a bad example

Example by Kanungo et al. (2004):

- Consider 4 points on the line such that $z > y > x$.
- For $k = 3$, the optimal cost is $\phi^* = x^2/2$.



Consider the heuristic solution:

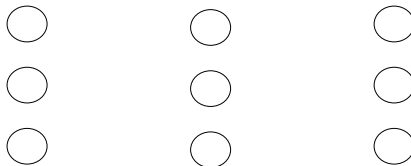
- $\phi = y^2/2$.
- The approximation ratio is y^2/x^2 , which can be made arbitrarily bad (increase y while keeping $y < z$).
- Let initial centers: first, third, fourth points. **k-means** terminates (why?)



Class Exercise

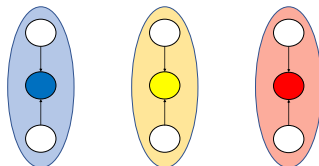
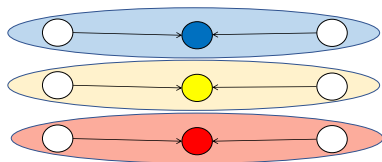
Class Exercise

Consider the clustering instance below. Discuss with your neighbors the worst-case that might be found by 3-means. Also, discuss why 3 clusters.



Class Exercise

The worst choice of centers (left) vs. the optimal clustering (right): The filled circles indicate the initial (and final) centers under each scheme.



k-means: Pros and Cons

Strengths:

- Simple and easy to implement, and relatively efficient
- Fast convergence in practice (often in a few iterations)

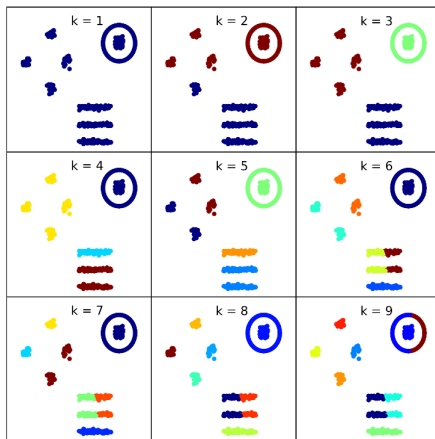
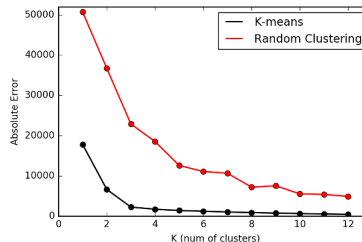
Weaknesses:

- Needs k in advance
- Converges to an arbitrarily bad locally optimal solution
- Sensitive to the choice of initialization
- Sensitive to outliers



How to Choose k ?

How many clusters?



Increasing k always leads to smaller cost, but choosing smaller k is consistent with Occam's razor.



Seeding Methods

The choice of seeds significantly impacts the quality of k -means's output.

In face of this problem:

- Repeat k -means several times with **different initial cluster centers**, and accept the best clustering (i.e., the one with the smallest cost)
- Resort to using wiser seeding methods, that are more adaptive to data.



Seeding Methods

The choice of seeds significantly impacts the quality of k -means's output.

In face of this problem:

- Repeat k -means several times with **different initial cluster centers**, and accept the best clustering (i.e., the one with the smallest cost)
- Resort to using wiser seeding methods, that are more adaptive to data.



Seeding Methods

The choice of seeds significantly impacts the quality of k -means's output.

In face of this problem:

- Repeat k -means several times with **different initial cluster centers**, and accept the best clustering (i.e., the one with the smallest cost)
- Resort to using wiser seeding methods, that are more adaptive to data.

Some popular seeding methods:

- Farthest Traversal
- k -means++



k-means++

k-means++ is a seeding method for k-means based on **adaptive sampling**, and is proposed in (Arthur & Vassilvitskii, 2007).



k-means++

k-means++ is a seeding method for k-means based on **adaptive sampling**, and is proposed in (Arthur & Vassilvitskii, 2007).

Given a set C , define:

$$D(x, C) := \min_{c \in C} \|x - c\|$$

input: \mathcal{X}

initialization: Choose c_1 uniformly at random from \mathcal{X} .

for $i = 2, \dots, k$

- Take c_i , by sampling x with probability

$$\frac{D(x, \{c_1, \dots, c_{i-1}\})^2}{\sum_{y \in \mathcal{X}} D(y, \{c_1, \dots, c_{i-1}\})^2}$$

k-means step: Run k-means using the initial centers c_1, \dots, c_k .



k-means++: Example

Consider the dataset below and assume $k = 3$. How does k-means++ proceeds?



k-means++: Example

Consider the dataset below and assume $k = 3$. How does k-means++ proceeds?



k-means++: Example

- **Step 1.** c_1 is chosen uniformly at random. So any point can be chosen as c_1 w.p. $\frac{1}{4}$.
- **Step 2.** Assume that $c_1 = (5, 0)$ is chosen. Now,

$$D((0, 1), c_1) = \sqrt{26}, \quad D((0, -1), c_1) = \sqrt{26}, \quad D((-5, 0), c_1) = 10, \quad D((5, 0), c_1) = 0,$$

$$\text{so that: } c_2 = \begin{cases} (0, 1) & \text{w.p. } \frac{26}{152} \\ (0, -1) & \text{w.p. } \frac{26}{152} \\ (-5, 0) & \text{w.p. } \frac{100}{152} \end{cases}$$

- **Step 3.** Assume that $c_2 = (0, 1)$ is chosen. Now,

$$D((0, -1), \{c_1, c_2\}) = \min(2, \sqrt{26}) = 2, \quad D((-5, 0), \{c_1, c_2\}) = \min(10, \sqrt{26}) = \sqrt{26}$$

$$D((5, 0), \{c_1, c_2\}) = D((0, 1), \{c_1, c_2\}) = 0,$$

$$\text{so that: } c_3 = \begin{cases} (-5, 0) & \text{w.p. } \frac{26}{30} \\ (0, -1) & \text{w.p. } \frac{4}{30} \end{cases}$$

- After c_3 is chosen, k-means will be applied with the chosen c_1, c_2 , and c_3 .

Assuming $c_3 = (-5, 0)$ is chosen, the final (output) clustering will be

$$C_1 = \{(5, 0)\}, \quad C_2 = \{(0, 1), (0, -1)\}, \quad C_3 = \{(-5, 0)\}$$

with the associated cluster centers at $(5, 0), (0, 0), (-5, 0)$ (why?). And the associated cost is $\phi(\mathcal{X}, \{C_1, C_2, C_3\}) = 2$.



k-means++: Clustering Quality

- k-means++ improves over k-means empirically.



k-means++: Clustering Quality

- k-means++ improves over k-means empirically.
- k-means++ is $O(\log(k))$ -competitive in expectation.

Theorem

For the output constructed using k-means++, the corresponding cost function satisfies:

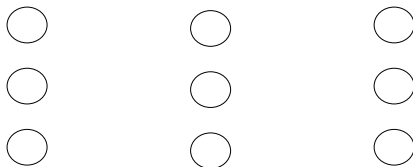
$$\mathbb{E}[\phi] \leq 8(\log(k) + 2) \cdot \phi^*$$



Class Exercise

Class Exercise

Consider the clustering instance below. Discuss with your neighbor how k -means++ could avoid outputting a bad clustering.



Useful References

- Stuart Lloyd. “Least squares quantization in PCM.” *IEEE Transactions on Information Theory* 28.2 (1982): 129–137.
- David Arthur and Sergei Vassilvitskii. “k-means++: the advantages of careful seeding.” *Proceedings of the Eighteenth Annual ACM-SIAM Symposium on Discrete Algorithms*. 2007.
- Johannes Blömer, et al., “Theoretical analysis of the k-means algorithm – a survey.” *Algorithm Engineering*. Springer, Cham, 2016. 81–116.
- Tapas Kanungo, et al. “A local search approximation algorithm for k-means clustering.” *Computational Geometry* 28.2-3 (2004): 89–112.

