



Faculty of Science



Domain-Specific Programming Languages

Torben Mogensen

Programming Language Design



GPL Versus DSL

General-purpose languages (GPLs) are programming languages intended for a broad range of problem domains, and typically for use by trained programmers.

Domain-specific languages (DSLs) are programming languages intended for either a narrow range of problem domains or for use by a certain type of non-programmers (and sometimes both).

Mostly a matter of degree: Few languages are completely general purpose, nor are they normally useless outside a single, narrow domain.



Why Should You Design a DSL?

- A DSL can potentially increase productivity within the domain.
- A DSL can use notation and concepts known to domain experts, which can help domain experts to understand programs and allow them to write simple programs.
- A DSL can be written to either guarantee certain properties or make it possible to analyse programs to determine these.
- An implementation of a DSL can employ domain-specific optimisations not usually found in implementations of GPLs.
- A program written in a DSL can have multiple different well-defined semantics.
- Training non-programmers to use a DSL may be easier than training them to use a GPL.
- If some element of the problem domain changes, it may be sufficient to handle this in the implementation of the DSL rather than having to rewrite all programs. Example: Dates.



Why Should You *Not* Design a DSL?

- The problem domain may not be specific enough or may not have features that lend themselves to domain-specific notation or optimisation.
- Designing, implementing and maintaining a DSL is a significant investment.
- Interfacing with other software may be difficult.
- If the intended users already know a GPL, training them to use a DSL is an extra investment.
- A program with command-line options or menu choices may suffice.



How do You Design a DSL?

- ➊ Identify the problem domain.
- ➋ Collect domain knowledge.
- ➌ Domain analysis and theory building.
- ➍ Concrete language design.
- ➎ Implementation.

May need iteration. Should always consider the needs and knowledge of the user over that of the designer / implementer. Nevertheless, good knowledge of language design and implementation needed to avoid pitfalls.



Before You Continue

Exercise 11.1:

- Discuss to what extent JSON and XML can be considered DSLs.
- Discuss to what extent a notation for regular expressions is a DSL.



How do You Implement a DSL?

- Embedded language:** The language is represented within an existing language, either as data structure or program structure. Type system of host language can help ensure syntactic correctness of embedded programs.
- Preprocessor:** Domain-specific features are expanded by a preprocessor into existing features in an existing language, leaving all other parts unchanged.
- Compiler/interpreter modification:** A compiler or interpreter for an existing language is modified to add, modify or remove features from original language as needed.
- Stand-alone language:** Implemented using a traditional interpreter or compiler with its own parser, type system etc.



Example DSL: SQL SELECT Statement

Grammar for subset of SQL SELECT statement:

SelectStatement → SELECT *Columns* FROM **TableName** *WhereClause*

Columns → **ColumnName**

Columns → *Columns* , *Columns*

WhereClause → WHERE *Condition*

WhereClause →

Condition → *Value* **RelOp** *Value*

Condition → *Condition* AND *Condition*

Value → **ColumnName**

Value → **Constant**

An example of a SELECT statement is

SELECT name FROM students WHERE enrollmentYear < 2015



Implementation as Embedded Language

Implement inside existing language.

- + It is usually easier to make a suitable data structure or program structure within an existing language than it is to make or modify an interpreter or compiler.
- + A programmer using the DSL can use abstraction and structuring features of the GPL around the DSL parts.
- + It is easy to interface programs written in the DSL with other programs.
- You are limited to use the syntax of the implementation language.
- Error messages are often expressed in terms of the underlying implementation language and not in terms of concepts from the DSL.
- It may be impossible to statically verify domain-specific properties, so these may have to be checked at runtime.



Representing Syntax in Embedded Language

- As a string, e.g.,
`"SELECT name FROM students WHERE enrollmentYear < 2015"`
 - Readable syntax, but requires parser
 - Syntax not checked at compile time
 - Possible to build syntax at run time
 - SQL injection!
- As a datatype (as in ML, F#, or Haskell), e.g.,
`SELECTFROM (["name"], "students",
 WHERE (Column "enrollmentYear" << Number 2015))`
 - Fairly readable syntax, no parser required
 - Syntax checked at compile time (type system of host language)
 - Possible to build syntax at run time
- Using method chaining, e.g.,
`Database("students").
 WHERElt(Column("enrollmentYear"), Number(2015)).SELECT("name")`
 - Less readable syntax
 - Syntax checked at compile time (type system of host language)
 - No parser or match/case-statement required
 - Not possible to build syntax at run time



Implementation by Preprocessor

- + Domain-specific notation can (to some extent) be used.
- + Simple domain-specific properties can be verified statically.
- + Compile-time errors can be expressed in terms of domain-specific concepts – as long as the user only uses domain-specific features.
- + A programmer using the DSL can use the abstraction and structuring features of the GPL around the DSL parts.
- + It is easy to interface programs written in the DSL with other programs.
- Run-time errors are in terms of expanded code, not original code.
- Unless full parser is written, syntax errors reported by preprocessor may be primitive.



Implementation by Compiler/Interpreter Modification

- + Domain-specific notation can fairly freely be used.
- + More complex domain-specific properties can be verified statically.
- + Both compile-time and run-time errors can be expressed in terms of domain-specific concepts.
- + A programmer using the DSL can use the abstraction and structuring features of the GPL, but restrictions can be imposed.
- + It is easy to interface programs written in the DSL with other programs.
- Requires access to source code of compiler/interpreter.
- Tie-in to specific (eventually obsolete?) compiler/interpreter.



Implementation as Stand-Alone Language

Complete, traditional implementation.

- + Full domain-specific notation can be used, including diagrams and complex formula notation.
- + Domain-specific properties can be guaranteed or verified statically.
- + Errors can be expressed in terms of domain-specific concepts.
- + More advanced domain-specific optimisations can be used.
- + It is possible to implement multiple different semantics for the same language.
- Requires implementation and maintenance of a full compiler or interpreter.
- Interfacing with other programs may be difficult. It is usually not difficult to make other programs call programs written in the DSL, but the converse may compromise domain-specific properties that are otherwise guaranteed or verifiable.



Before You Continue

Exercise 11.2:

The DSL that has arguably been most successful in being used by non-programmers is Excel, the spreadsheet program in Microsoft's Office package.

Discuss which properties of Excel makes it suitable for non-programmers, and which properties that might hinder this.



Examples of DSLs

Scratch: A language for teaching children to program.

T_EX and L^AT_EX: Languages for type-setting.

Graphviz: A language for writing graph-like diagrams.

Troll: A language for rolling dice and calculating probabilities thereof.

CRL: A language for cashflow reengineering.

Futhark: A language for programming graphics processors.



Scratch



Intended for teaching children how to program.

- Visual LEGO-inspired syntax.
- Problem domain is making interactive 2D animations and games.
- Nevertheless, has advanced features such as concurrency and recursion.



T_EX and L^AT_EX

T_EX was designed by Donald Knuth to typeset “The Art of Computer Programming”. Includes Turing-complete macro language – a BASIC interpreter has been programmed in T_EX.

L^AT_EX built on top of T_EX by Leslie Lamport. Adds features for structuring a document. It is more widely used than raw T_EX.


L^AT_EX is the de-facto standard for writing articles, reports, and textbooks in the natural sciences, but has also been used to typeset music notation, circuit diagrams, role-playing games, and much more.



Graphviz

Graphviz (<https://graphviz.org/>) is a domain-specific language for visualising trees and graphs: You describe a graph using textual notations, and Graphviz will convert this to an image.

```
digraph G {  
    main -> parse -> execute  
    main -> {init, cleanup}  
    execute -> {make_string, printf, compare}  
    init -> make_string  
    main -> printf  
}
```



Troll

A language for rolling dice and calculating probabilities thereof (two different semantics for one language).

Dice can be combined in a myriad of ways. Examples:

- Count the number of different results obtained when rolling five ten-sided dice:

```
count different 5d10
```

- Generate random attributes for D&D:

```
"Str |>Dex|>Con|>Int|>Wis|>Chr"  
|| 6'sum largest 3 4d6
```

See more at <https://topps.diku.dk/~torbenm/troll.msp>.



CRL

A language for creating financial papers (loans, etc.) from others.

Example:

```
declarations
```

```
    fraction = 0.7
```

```
input
```

```
    cashflow a, b
```

```
structure
```

```
    c = a + b
```

```
    (d,e) = Divide(c, fraction)
```

```
output
```

```
    cashflow d,e
```

Uses *linear types* to prevent multiple use of the same money: Every variable is used *exactly* once.



Futhark

Functional language for doing parallel computations on graphics processors. Not intended for complete programs, but for compute-intensive cores that are called from other programs.

Uses `map`, `filter`, `reduce`, and `scan` or equivalent functions expressed as loops. Uses *uniqueness types* to allow arrays to be destructively overwritten.

Matrix multiplication:

```
fun main(x: [n][m]i32, y: [m][p]i32): [n][p]i32 =  
  map (\xr -> map (\yc ->  
                    reduce (+) 0 (zipWith (*) xr yc))  
      (transpose y))
```

x



Before You Continue

Which DSLs not mentioned in the slides do you know, and what makes them domain-specific?

